


Programming Language Basics

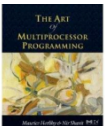


Companion slides for
The Art of Multiprocessor Programming
by Maurice Herlihy & Nir Shavit

With some very minor changes by APS

Languages for Multiprocessor Programming

- Java
- PThreads
 - C and C++
- C#
- MPI
- Etc...




Threads

- Execution of a sequential program
- You can tell a thread
 - What to do
 - When to start
- You can
 - Wait for it to finish
- Other stuff:
 - Interrupt it, give it priority, etc.




Threads in Java

- Class **java.lang.Thread**
- Each thread has a method
 - **Void run()**
- Executes when it starts
- Thread vanishes when it returns
- You must provide this method




Creating a Thread

- Create a **Runnable** object
 - **Runnable** is an interface
 - Provides **run()** method
- Pass **Runnable** object to thread constructor



A Runnable Class

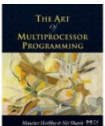
```
public class Hello implements Runnable {  
    String message;  
    public Hello(m) {  
        message = m;  
    }  
    public void run() {  
        System.out.println(message);  
    }  
}
```



A Runnable Class


```
public class Hello implements Runnable {  
    String message;  
    public Hello(m) {  
        message = m;  
    }  
    public void run() {  
        System.out.println(message);  
    }  
}
```

Runnable interface



Creating a Thread


```
String m = "Hello from " + i;  
Runnable h = new Hello(m);  
Thread t = new Thread(h);
```



Creating a Thread

```
String m = "Hello from " + i;  
Runnable h = new Hello(m);  
Thread t = new Thread(h);
```

Create a Runnable object




Creating a Thread

```
String m = "Hello from " + i;
```

```
Runnable h = new Hello(m);
```

```
Thread t = new Thread(h);
```

Create the thread




Syntactic Help

- Defining a single-use class like Hello can be a nuisance
- Java provides special syntax
- Anonymous inner classes
 - May be more trouble than it's worth
 - You should recognize it



Anonymous Inner Class

```
t = new Thread(  
    new Runnable() {  
        public void run() {  
            System.out.println(m);  
        }  
    }  
);
```



Anonymous Inner Class

```
t = new Thread(  
    new Runnable() {  
        public void run() {  
            System.out.println(m);  
        }  
    }  
);
```


Creates object of
anonymous Runnable
class



Anonymous Inner Class

```
t = new Thread(  
    new Runnable() {  
        public void run() {  
            System.out.println(m);  
        }  
    }  
);
```


**Calls Thread constructor
with anonymous object**



Starting a Thread

```
t.start();
```

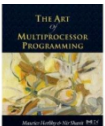
- Starts the new thread
- Caller returns immediately
- Caller & thread run in parallel



Joining a Thread


```
t.join();
```

- Blocks the caller
- Waits for the thread to finish
- Returns when the thread is done



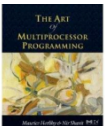
Monitors

- Each object has an implicit lock
- Managed by **synchronized** modifier
 - Methods
 - Code blocks
- OK for easy cases
- Not always for hard cases



Call Center Scenario

- Calls arrive faster than they can be answered
 - Play recorded message
 - “your call is very important to us ...”
 - Put call in queue
 - Play insipid music ...
 - Operators dequeue call when ready



Bad Queue Implementation

```
class Queue<T> {  
    int head = 0, tail = 0;  
    T[] items = new T[QSIZE];  
    public enq(T x) {  
        items[(tail++) % QSIZE] = x;  
    }  
    public T deq() {  
        return items[(head++) % QSIZE]  
    }  
}
```



Bad Queue Implementation

```
class Queue<T> {
```

```
    int head = 0, tail = 0;
```

```
    T[] items = new T[QSIZE];
```

```
    public enq(T x) {
```

```
        items[(tail++) % QSIZE] = x;
```


```
    }
```

```
    public T deq() {
```

```
        return items[(head++) % QSIZE]
```

```
    } }
```

Works for generic
type T



Bad Queue Implementation

```
class Queue<T> {  
    int head = 0, tail = 0;  
    T[] items = new T[QSIZE];  
    public enq(T x) {  
        items[(tail++) % QSIZE] = x;  
    }  
    public T deq() {  
        return items[(head++) % QSIZE];  
    }  
}
```


Array of T items



Bad Queue Implementation

```
class Queue<T> {  
    int head = 0, tail = 0;  
    T[] items = new T[QSIZE];  
    public enq(T x) {  
        items[(tail++) % QSIZE] = x;  
    }  
    public T deq() {  
        return items[(head++) % QSIZE];  
    }  
}
```


next slot to dequeue, 1st empty slot



Bad Queue Implementation




```
class Queue<T> {  
    int head = 0, tail = 0;  
    T[] items = new T[QSIZE];  
    public enq(T x) {  
        items[(tail++) % QSIZE] = x;  
    }  
    public T deq() {  
        return items[(head++) % QSIZE];  
    }  
}
```

Put in empty slot,
advance head




Of course, this doesn't work

head






Mutual Exclusion

- Only one thread modifying queue fields at a time
- Use **synchronized** methods
 - Locks object on call
 - Releases lock on return



Mutual Exclusion

head



Synchronized Method

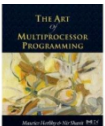
```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        items[(tail++) % QSIZE];  
    }  
    ...  
}
```



Synchronized Method


```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        items[(head++) % QSIZE];  
    }  
    ...  
}
```

Lock acquired on entry,
released on exit



Syntactic Sugar

```
class Queue<T> {  
    ...  
    public enq(T x) {  
        synchronized (this) {  
            items[(tail++) % QSIZE];  
        }  
    }  
    ...  
}
```



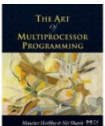
Syntactic Sugar

```
class Queue<T> {  
    ...  
    public enq(T x) {  
        synchronized (this) {  
            items[(tail++) % QSIZE];  
        }  
    }  
    ...  
}
```

Same meaning, more
verbose


Vocabulary

- A **synchronized** method locks the object
- No other thread can call another **synchronized** method for that same object
- Code in middle is critical section




Re-entrant Locks

- What happens if you lock the same object twice?
 - In Java, no deadlock
 - Keeps track of number of times locked and unlocked
 - Unlock occurs when they balance out



Still Doesn't Work


```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        items[(tail++) % QSIZE] = x;  
    }  
    ...  
}
```



Still Doesn't Work

```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        items[(tail++) % QSIZE] = x;  
    }  
    ...  
}
```

What if the array is full?




Waiting

- What if
 - Enqueuer finds a full array?
 - Dequeuer finds an empty array?
- Throw an exception?
 - What can caller do?
 - Repeated retries wasteful
- Wait for something to happen



Waiting Synchronized Method

```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        while (head == tail) {};  
        items[(tail++) % QSIZE] = x;  
    }  
    ...  
}
```



Waiting Synchronized Method


```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        while (head == tail) {};  
        items[(tail++) % QSIZE] = x;  
    }  
    ...  
}
```

Spin while the array is full



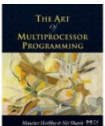
Deadlock

- Enqueuer is
 - Waiting for a dequeuer
 - While holding the lock
- Dequeuer
 - Waiting for enqueuer to release lock
- Nothing will ever happen




Waiting Thread

- Release lock while waiting
- When “something happens”
 - Re-acquire lock
 - Either
 - Re-release lock & resume waiting
 - Finish up and return




Styles of Waiting

- Spinning
 - Repeatedly retest condition
- Blocking
 - Ask OS to run someone else



Styles of Waiting


- Spinning
 - Good for very short intervals
 - Expensive to call OS
 - Works **only** on multiprocessors!
- Blocking
 - Good for longer intervals
 - Processor can do work
- Clever libraries sometimes mix



The wait() Method

```
q.wait();
```

- Releases lock on q
- Sleeps (gives up processor)
- Awakens (resumes running)
- Reacquires lock & returns



The wait() Method

```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        while (head == tail) {  
            wait();  
        };  
        items[(head++) % QSIZE] = x;  
    }  
    ...  
}
```




Waiting Synchronized Method

```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        while (head == tail) {  
            wait();  
        };  
        items[(head++) % QSIZE] = x;  
    }  
    ...  
}
```

Keep retesting condition

}




Waiting Synchronized Method

```
class Queue<T> {  
    ...  
    public synchronized enq(T x) {  
        while (head == tail) {  
            wait();  
        };  
        items[(head++) % QSIZE] = x;  
    }  
    ...  
}
```


Keep retesting condition

Release lock & sleep



Wake up and Smell the Coffee


- When does a waiting thread awaken?
 - Must be notified by another thread
 - when something has happened
- Failure to notify in a timely way is called a “lost wakeup”



The wait() Method

```
q.notify();
```

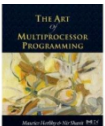
- Awakens one waiting thread
- Which will reacquire lock & returns



The wait() Method


```
q.notifyAll();
```

- Awakens all waiting threads
- Which will reacquire lock & return



The wait() Method


```
public synchronized enq(T x) {  
    while (head == tail) {  
        wait();  
    };  
    items[(head++) % QSIZE] = x;  
    if ((head-tail+QSIZE) % QSIZE == 1) {  
        notify();  
    }  
}
```



The wait() Method

```
public synchronized enq(T x) {  
    while (head == tail) {  
        wait();  
    };  
    items[(head++) % QSIZE] = x;  
    if ((head - tail + QSIZE) % QSIZE == 1) {  
        notify();  
    }  
}
```

Wait for empty slot



The wait() Method

```
public synchronized enq(T x) {  
    while (head == tail) {  
        wait();  
    };  
    items[(head++) % QSIZE] = x;  
    if ((head - tail + QSIZE) % QSIZE == 1) {  
        notify();  
    }  
}
```


Stuff item into array



The wait() Method


```
public synchronized enq(T x) {  
    while (head == tail) {  
        wait();  
    };  
    items[(head++) % QSIZE] = x;  
    if ((head-tail+QSIZE) % QSIZE == 1) {  
        notify();  
    }  
}
```

If the queue was empty,
wake up a dequeuer




Lost Wakeup


- This code has a lost wakeup bug
- Possible to have
 - Waiting dequeuer
 - Non-empty queue
- Because not enough threads awakened





Empty queue, waiting dequeuers





Enqueuer puts item in queue






Since queue was empty, wakes dequeuer



1st Dequeueer slow, overtaken by enqueueers




1st Dequeueur finishes




Solutions

- Don't write buggy code d'oh!
- Always call **notifyAll()**
- Can also use timed waits
 - Wake up after specified time




Thread-Local Data

- In many of our examples we assume
 - Threads have unique ids
 - In range $0, \dots, n-1$
- Where do they come from?
 - Long-lived data
 - Unique to a thread



Thread-Local Data in Java


- **ThreadLocal<T> class**
- No built-in language support
- Library classes
 - Syntax is awkward
 - Very useful anyway



ThreadLocal methods

```
ThreadLocal<T> local;  
T x = ...;  
local.set(x);
```


- Changes calling thread's version of object
- Other threads' versions unaffected



ThreadLocal methods

```
ThreadLocal<T> local;  
T x = local.get();
```

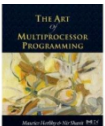
- Returns calling thread's version of object



Initializing ThreadLocals

```
T x = local.initialValue();
```

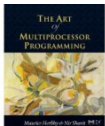
- Called by **get()** method the first time the thread-local variable is accessed.



Example

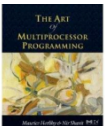
```
int me = ThreadID.get()
```

- Return unique thread id
- Take a number first time called



Thread-Local IDs


```
public class ThreadID {  
    private static volatile int nextID = 0;  
    private static LocalID threadID =  
        new LocalID();  
    public static int get() {  
        return threadID.get();  
    }  
    ... // define LocalID here  
}
```



Thread-Local IDs

```
public class ThreadID {  
    private static volatile int nextID = 0;  
    private static LocalID threadID =  
        new LocalID();  
    public static int get() {  
        return threadID.get();  
    }  
    ... // define LocalID here  
}
```

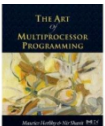
Next ID to assign



Thread-Local IDs

```
public class ThreadID {  
    private static volatile int nextID = 0;  
    private static LocalID threadID =  
        new LocalID();  
    public static int get() {  
        return threadID.get();  
    }  
    ... // define LocalID here  
}
```

Declare & initialize thread-local ID



Thread-Local IDs


```
public class ThreadID {  
    private static volatile int nextID = 0;  
    private static LocalID threadID =  
        new LocalID();  
    public static int get() {  
        return threadID.get();  
    }  
    ... // define LocalID here  
}
```

Return value of thread-local ID



The Inner Class


```
private static class LocalID
    extends ThreadLocal<Integer> {
    protected synchronized Integer
        initialValue() {
            return nextID++;
        }
}
```



The Inner Class

```
private static class LocalID
    extends ThreadLocal<Integer> {
    protected synchronized Integer
        initialValue() {
            return nextID++;
        }
}
```

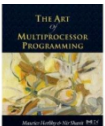
Subclass of ThreadLocal<Integer>



The Inner Class


```
private static class LocalID
    extends ThreadLocal<Integer> {
    protected synchronized Integer
        initialValue() {
            return nextID++;
        }
}
```

Overrides initialValue()



Summary

- Threads
 - And how to control them
- Synchronized methods
 - Wait, notify, and NotifyAll
- Thread-Local objects



This work is licensed under a Creative Commons Attribution-ShareAlike 2.5 License.

- **You are free:**
 - **to Share** — to copy, distribute and transmit the work
 - **to Remix** — to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
 - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.

