Multiprocessor Architecture Basics

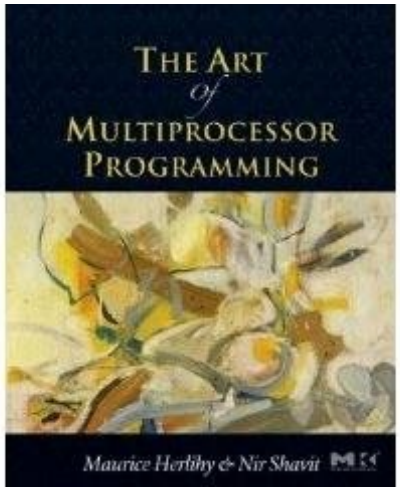

Companion slides for

The Art of Multiprocessor Programming

by Maurice Herlihy & Nir Shavit

With some very minor changes by APS

Multiprocessor Architecture

- Abstract models are (mostly) OK to understand algorithm correctness and progress
- To understand how concurrent algorithms actually perform
- You need to understand something about multiprocessor architectures

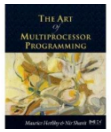

Pieces

- Processors
- Threads
- Interconnect
- Memory
- Caches

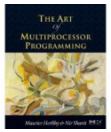

Old-School Multiprocessor

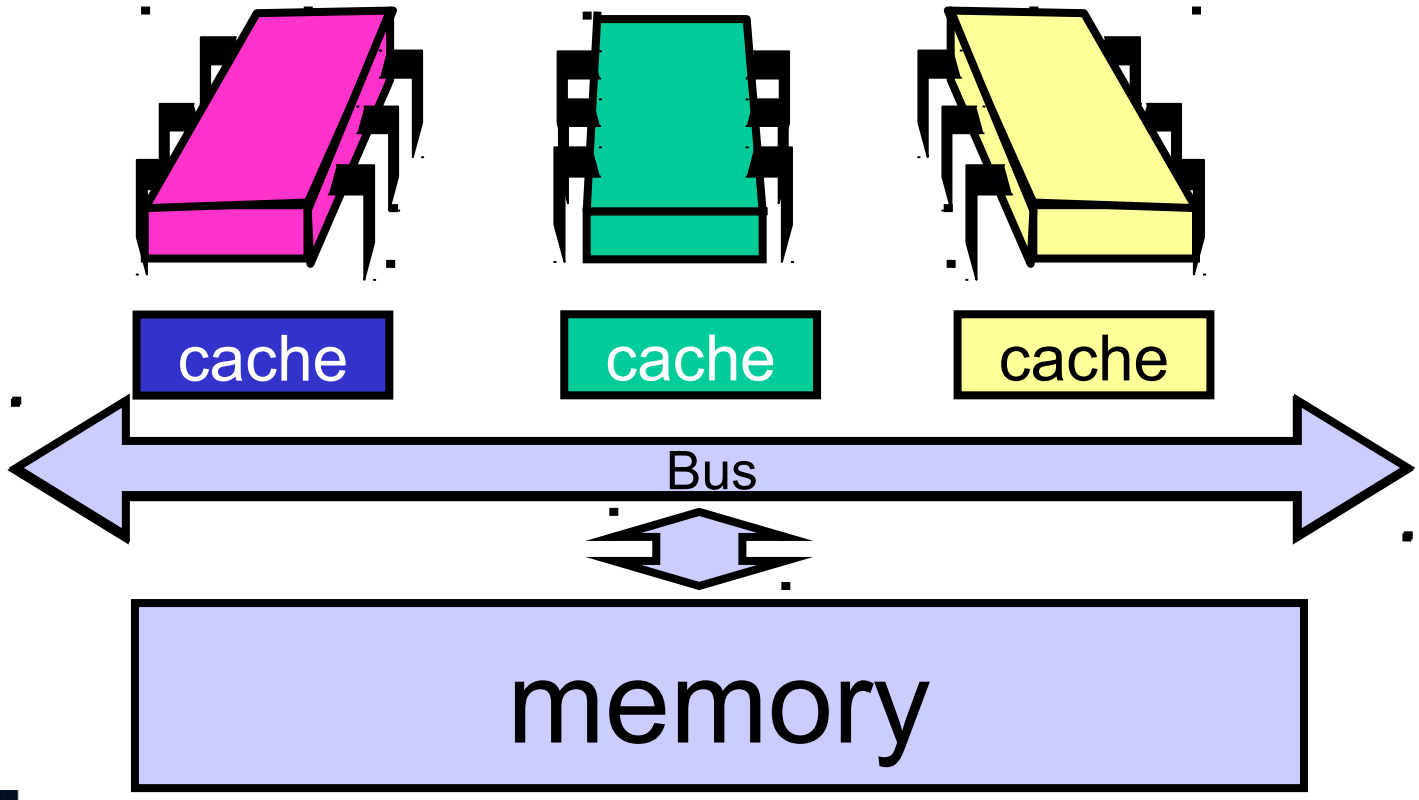

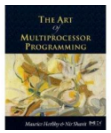

Old School

- Processors on different chips
- Processors share off chip memory resources
- Communication between processors typically slow

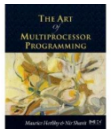

Multicore Architecture

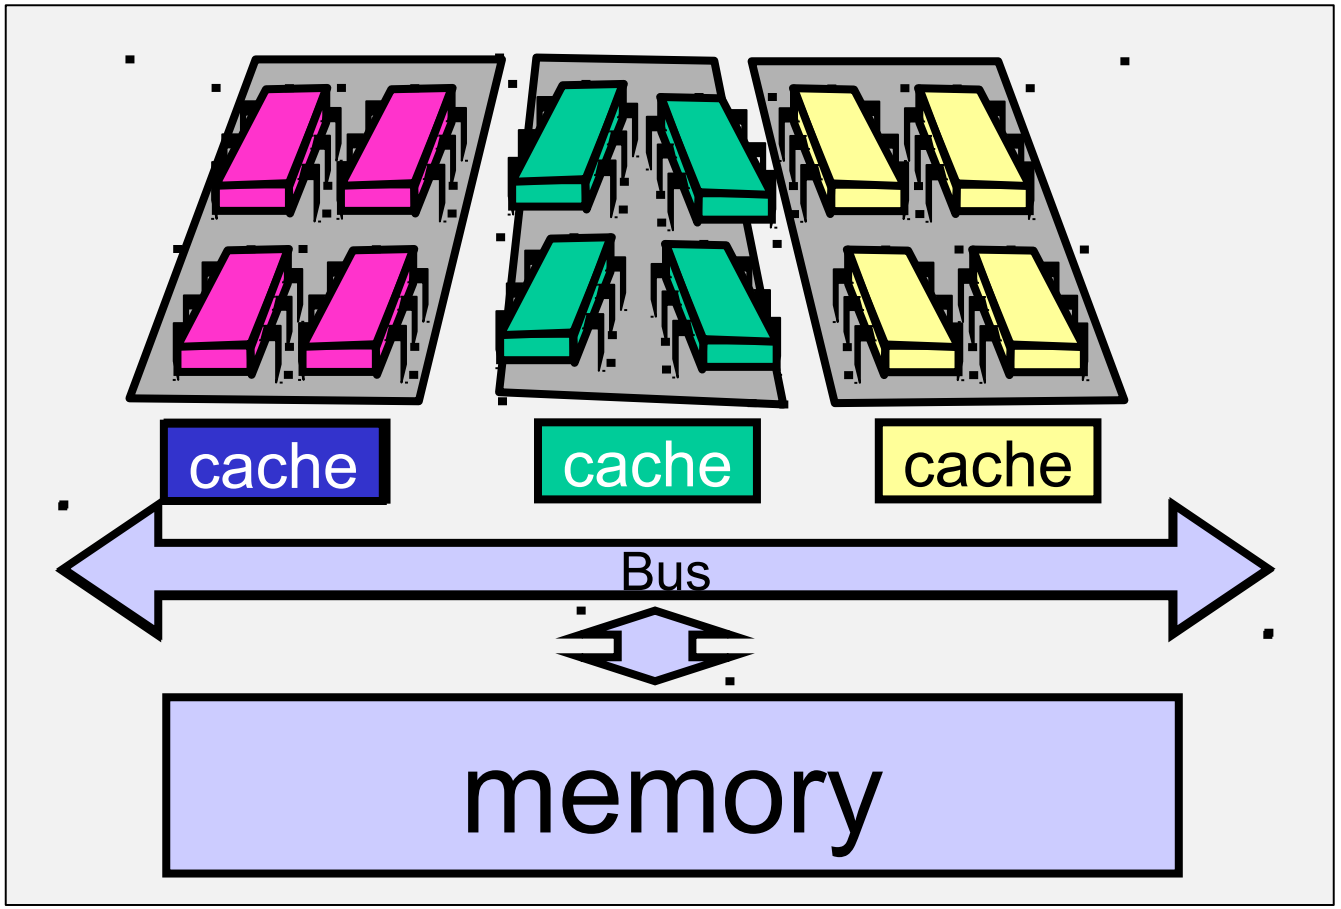

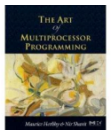

Multicore

- All Processors on same chip
- Processors share on chip memory resources
- Communication between processors now very fast

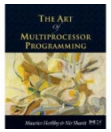

SMP vs NUMA

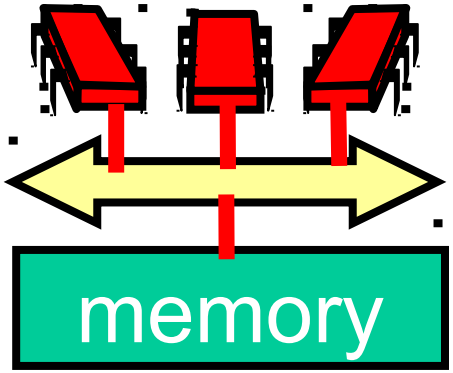

SMP

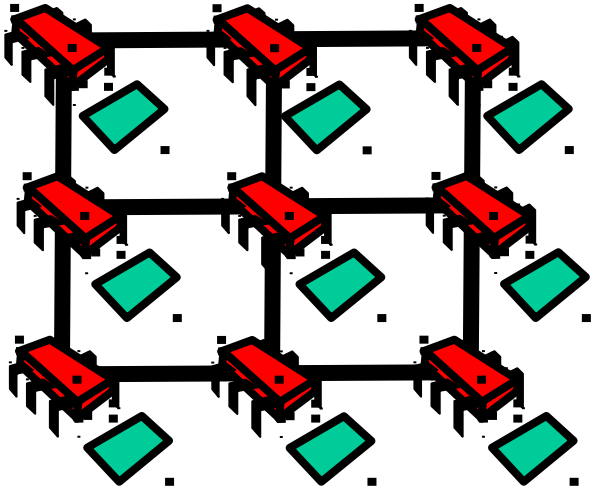

NUMA

- SMP: symmetric multiprocessor
- NUMA: non-uniform memory access
- CC-NUMA: cache-coherent …

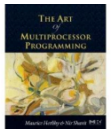

Future Multicores

- Short term: SMP
- Long Term: most likely a combination of SMP and NUMA properties

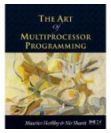

Understanding the Pieces

- Lets try to understand what the pieces that make the multiprocessor machine are
- And how they fit together

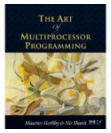

Processors

- Cycle:
  - Fetch and execute one instruction
- Cycle times change
  - 1980: 10 million cycles/sec
  - 2005: 3,000 million cycles/sec

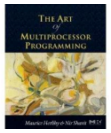

Computer Architecture

- Measure time in cycles
  - Absolute cycle times change
- Memory access: ~100s of cycles
  - Changes slowly
  - Mostly gets worse

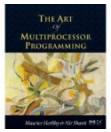

Threads

- Execution of a sequential program
- Software, not hardware
- A processor can run a thread
- Put it aside
  - Thread does I/O
  - Thread runs out of time
- Run another thread

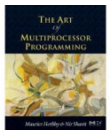

Analogy

- You work in an office
- When you leave for lunch, someone else takes over your office.
- If you don't take a break, a security guard shows up and escorts you to the cafeteria.
- When you return, you may get a different office

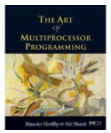

Interconnect

- Bus
  - Like a tiny Ethernet
  - Broadcast medium
  - Connects
    - Processors to memory
    - Processors to processors
- Network
  - Tiny LAN
  - Mostly used on large machines

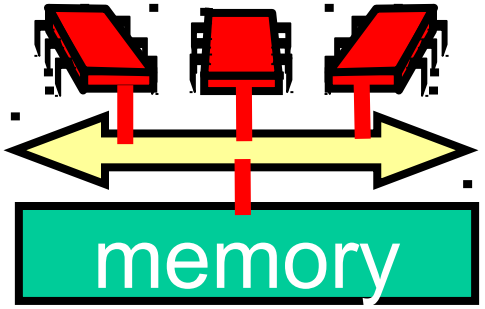

SMP

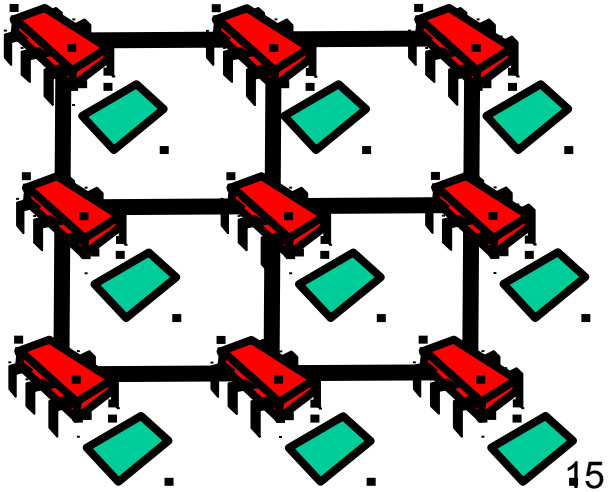

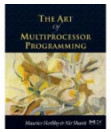

Interconnect

- Interconnect is a finite resource
- Processors can be delayed if others are consuming too much
- Avoid algorithms that use too much bandwidth

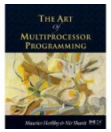

Processor and Memory are Far Apart

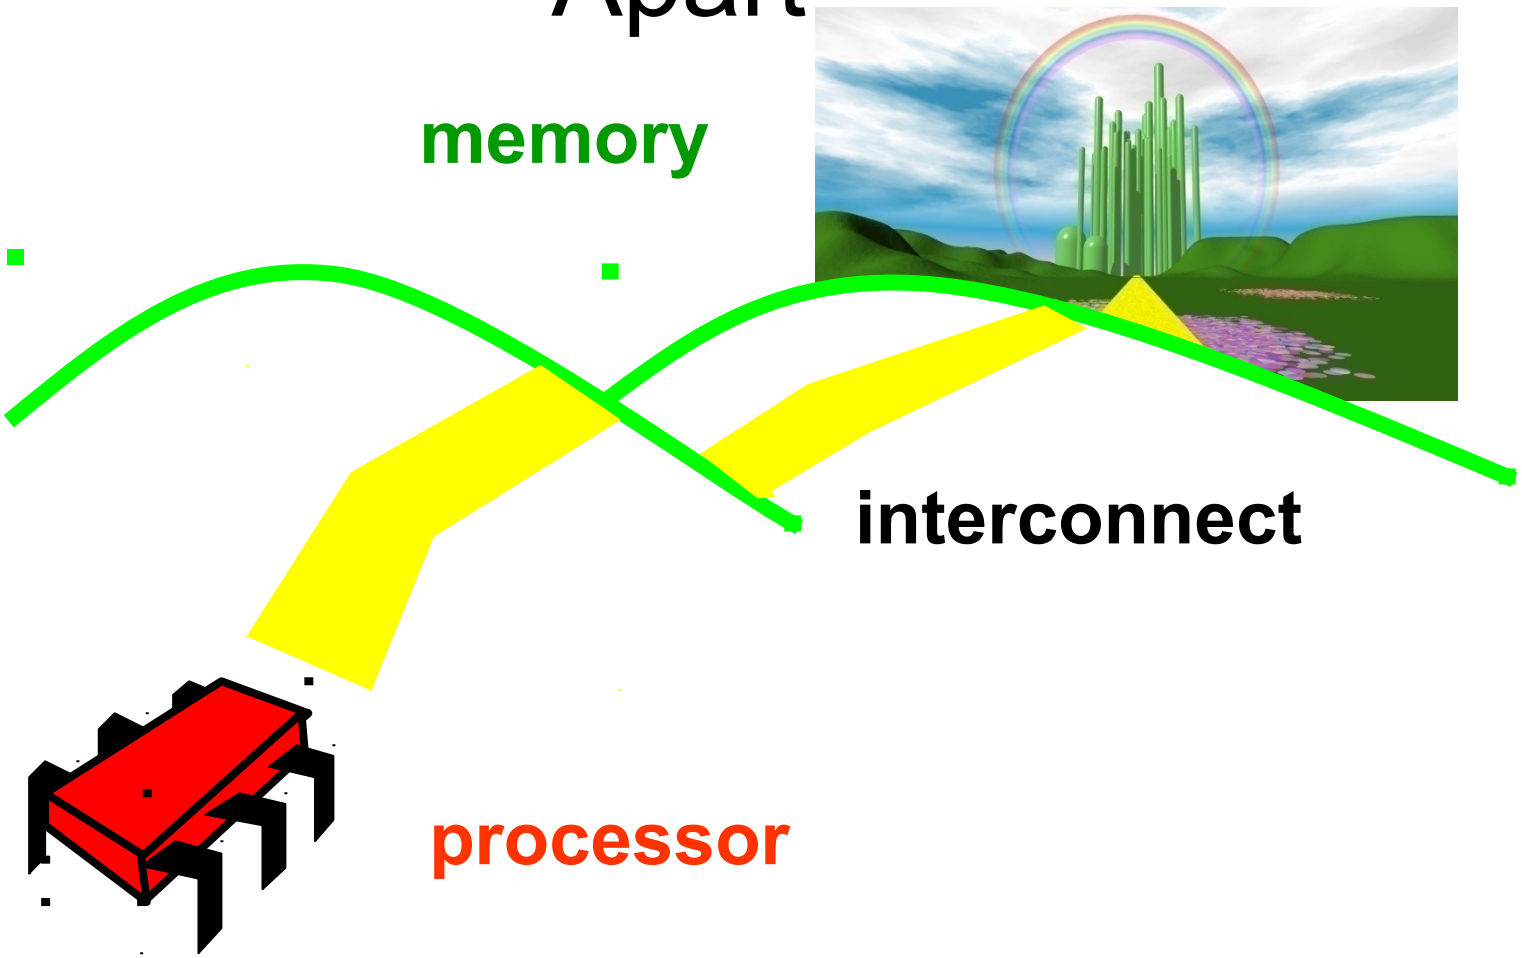

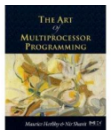

Reading from Memory

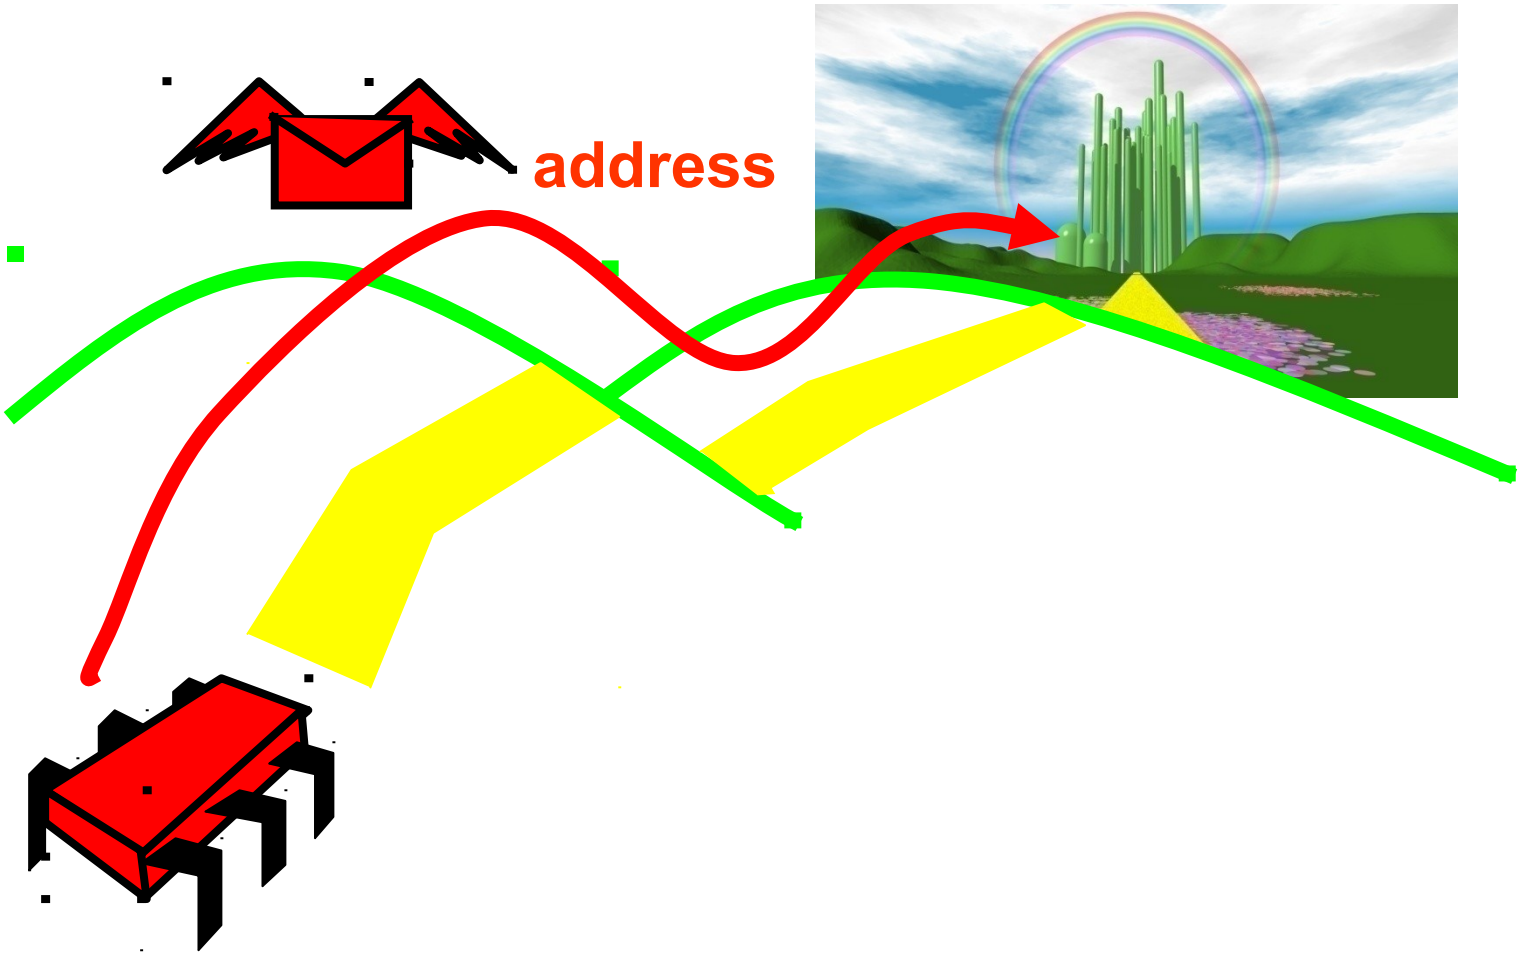

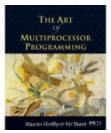

Reading from Memory

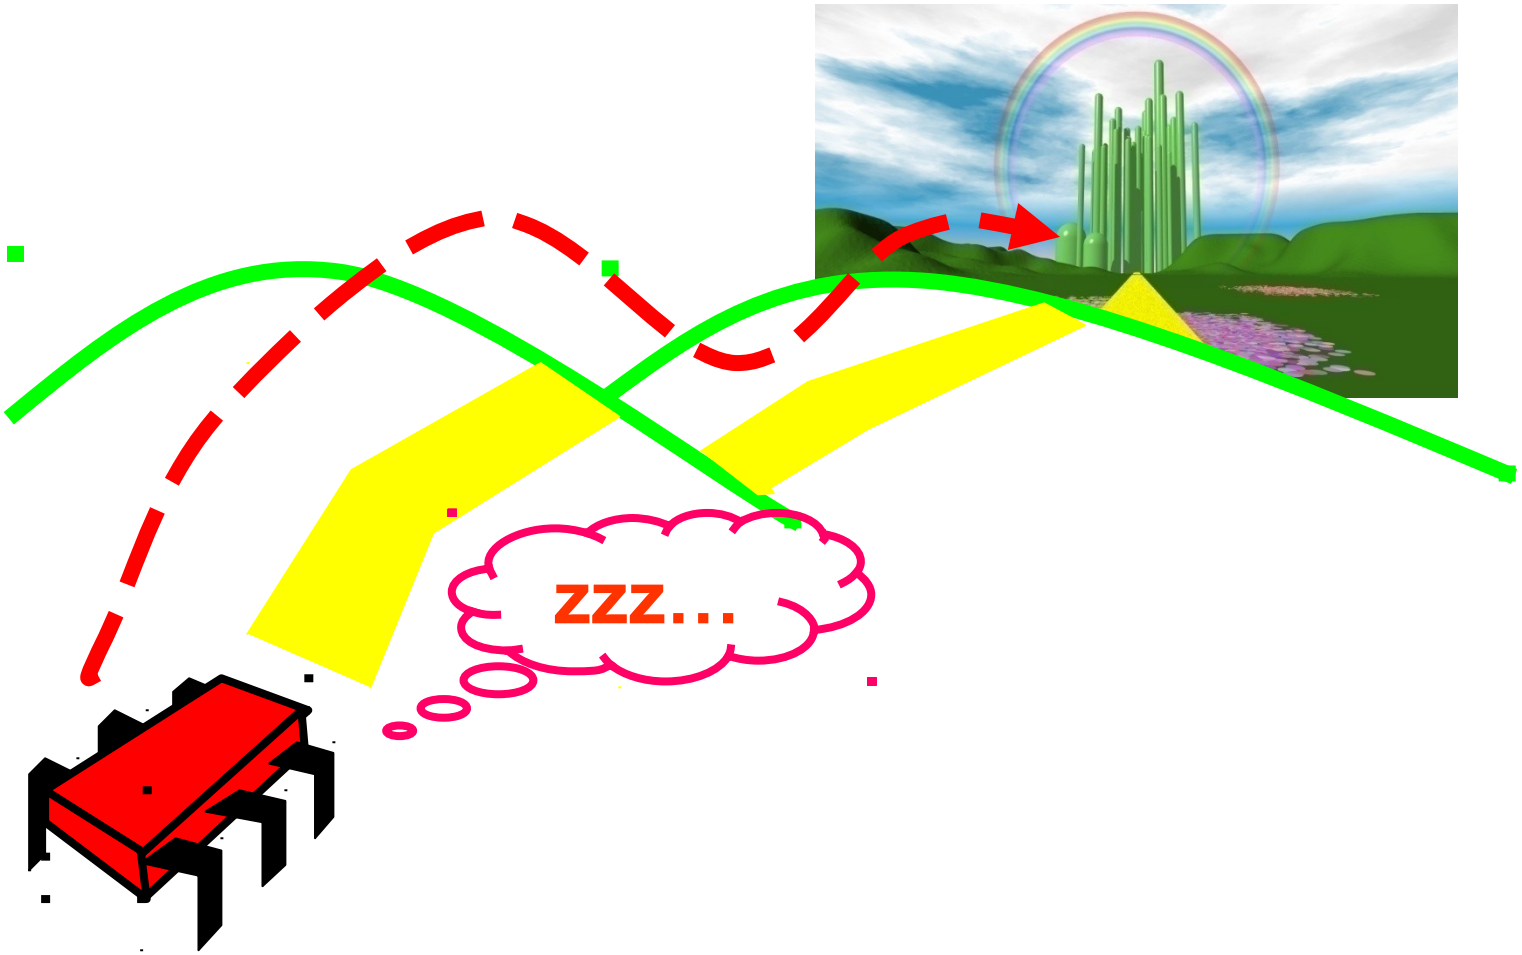

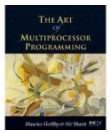

Reading from Memory

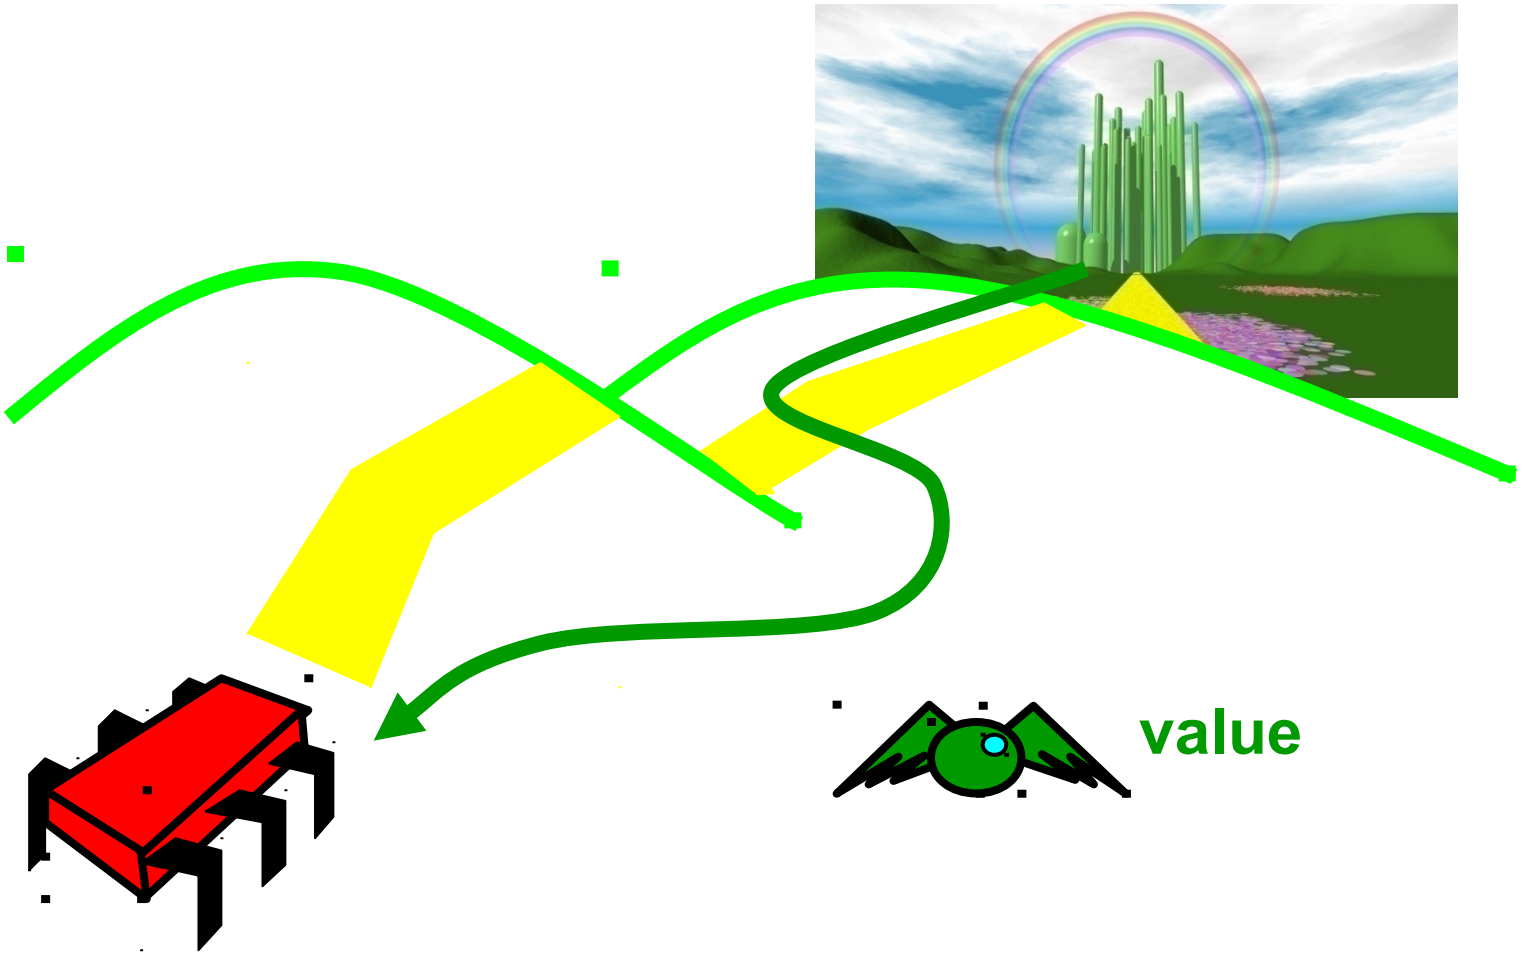

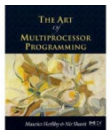

Writing to Memory

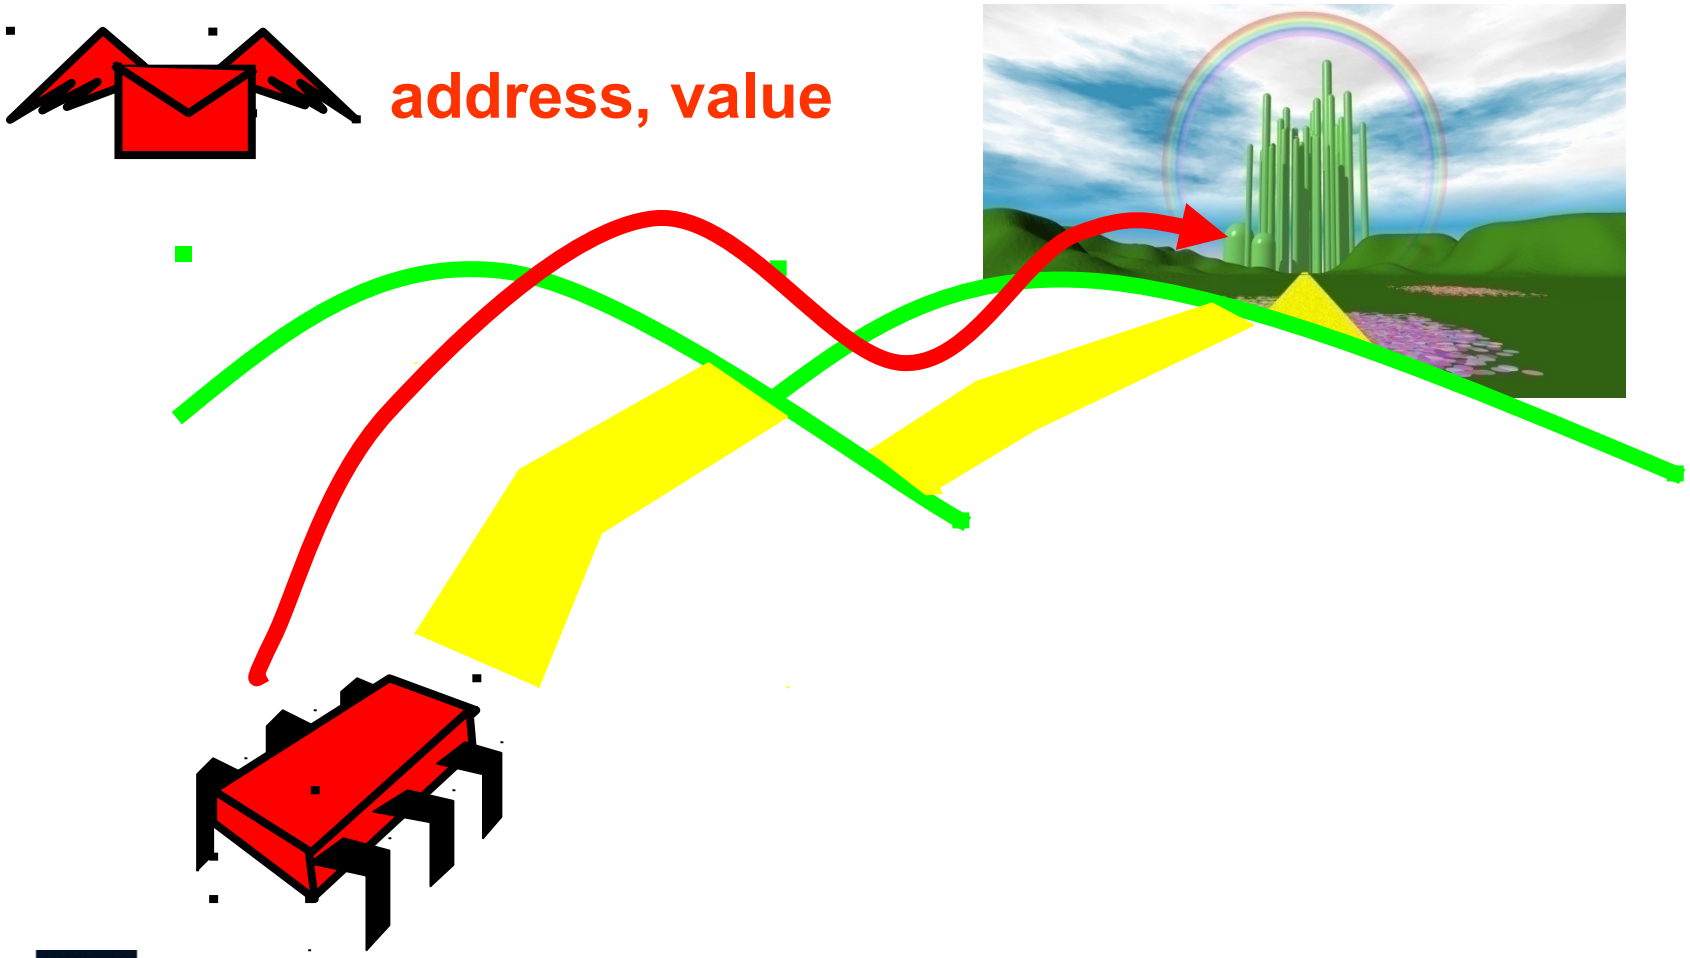

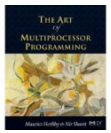

Writing to Memory

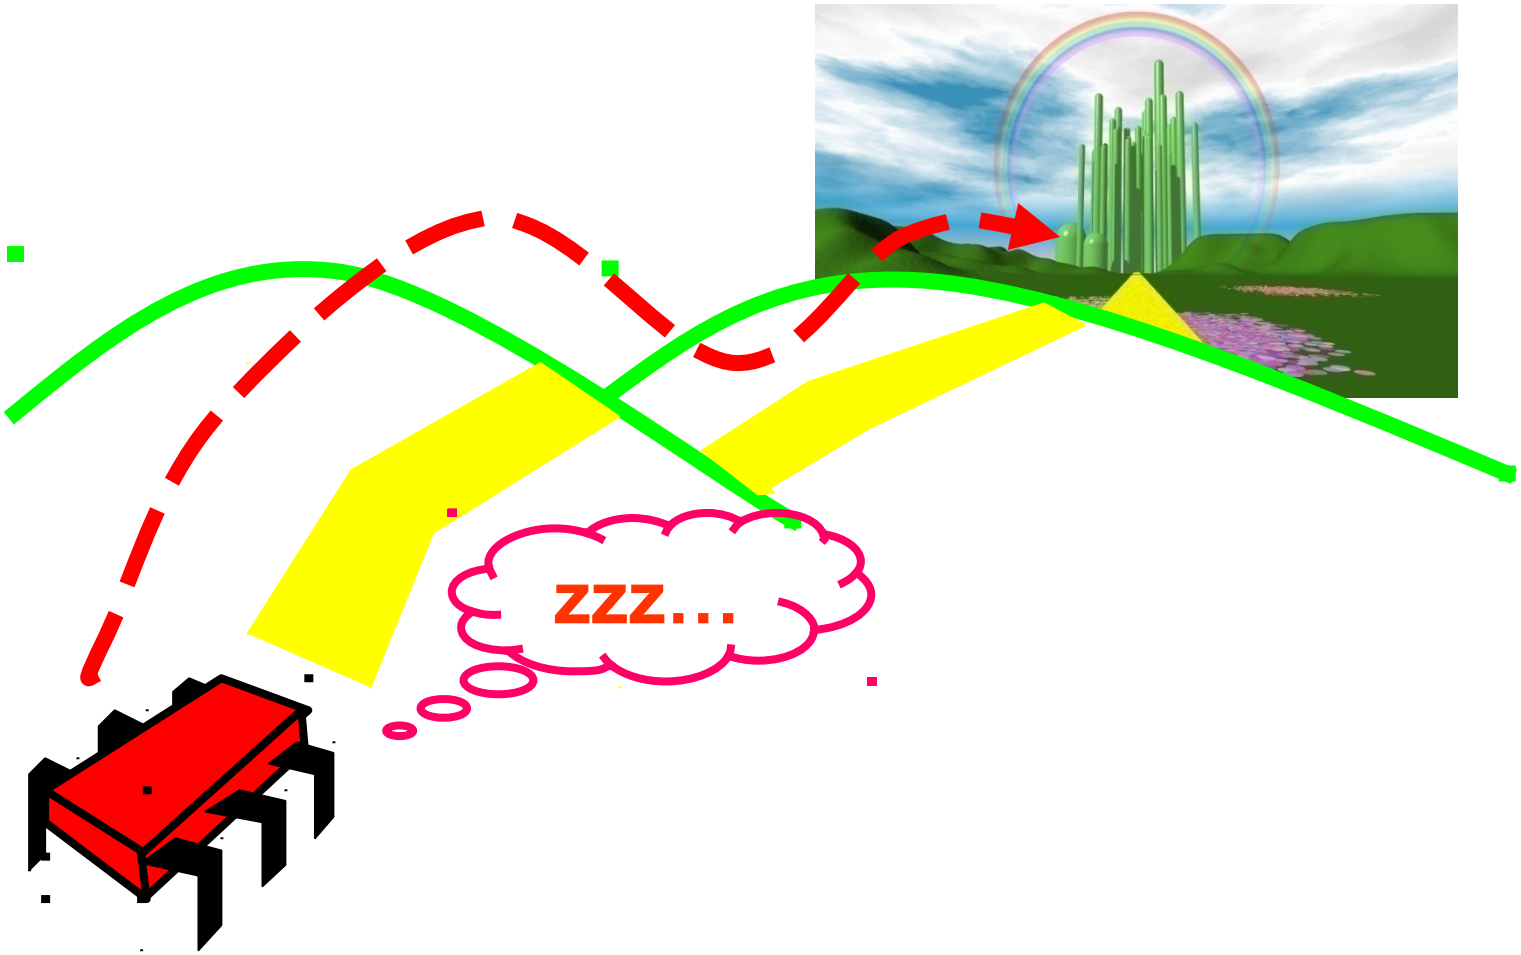

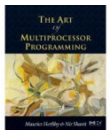

Writing to Memory

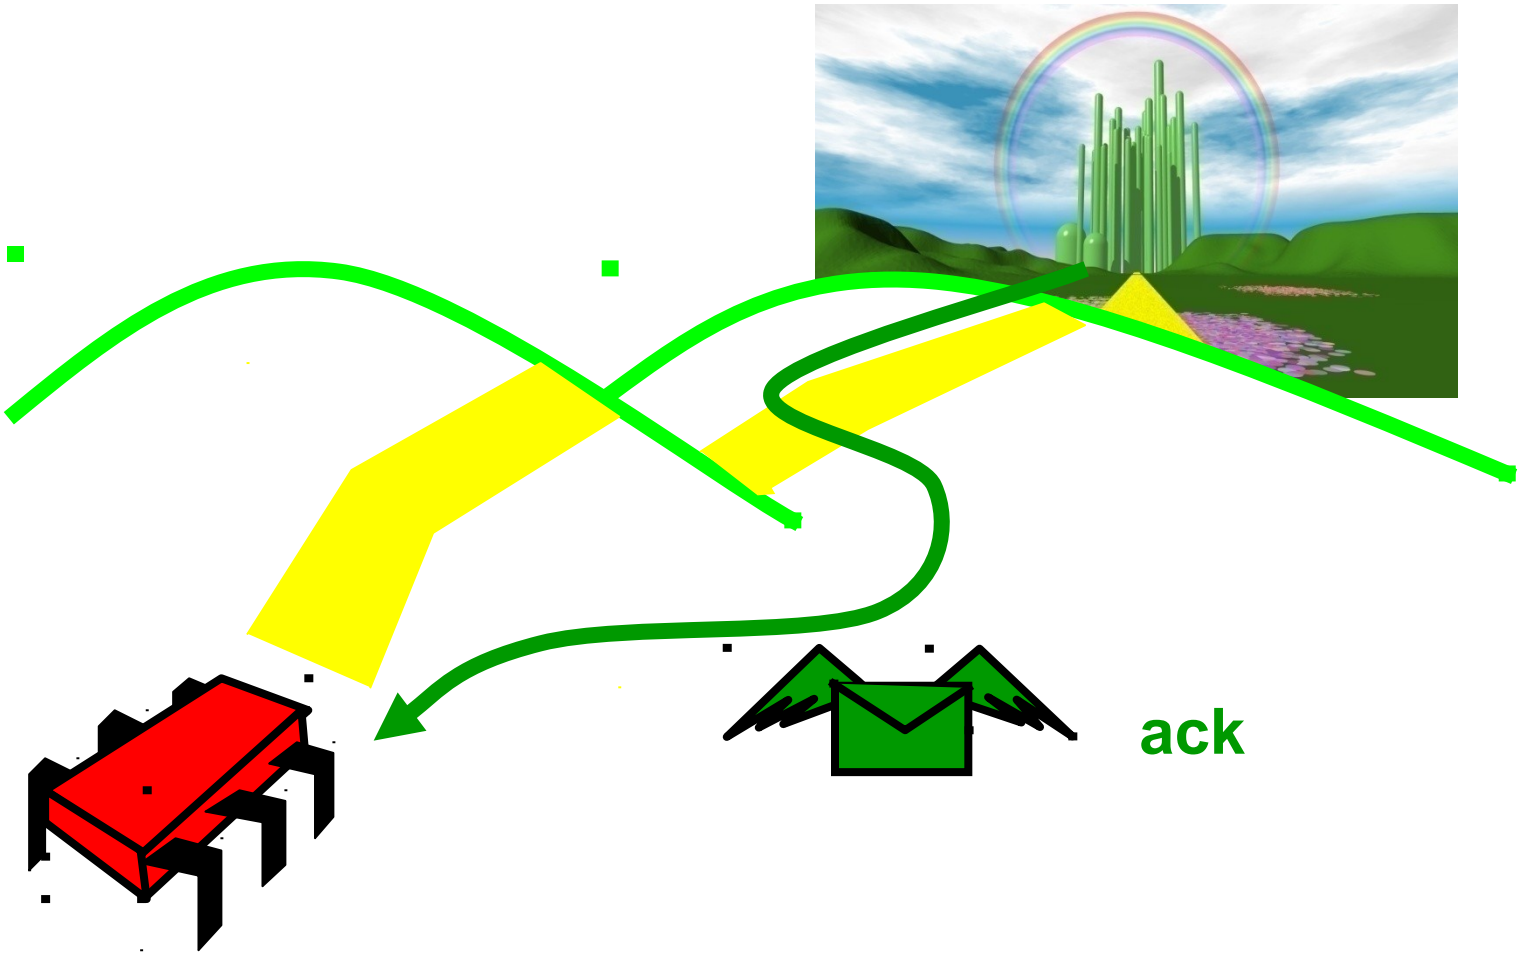

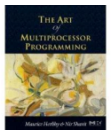

Cache: Reading from Memory

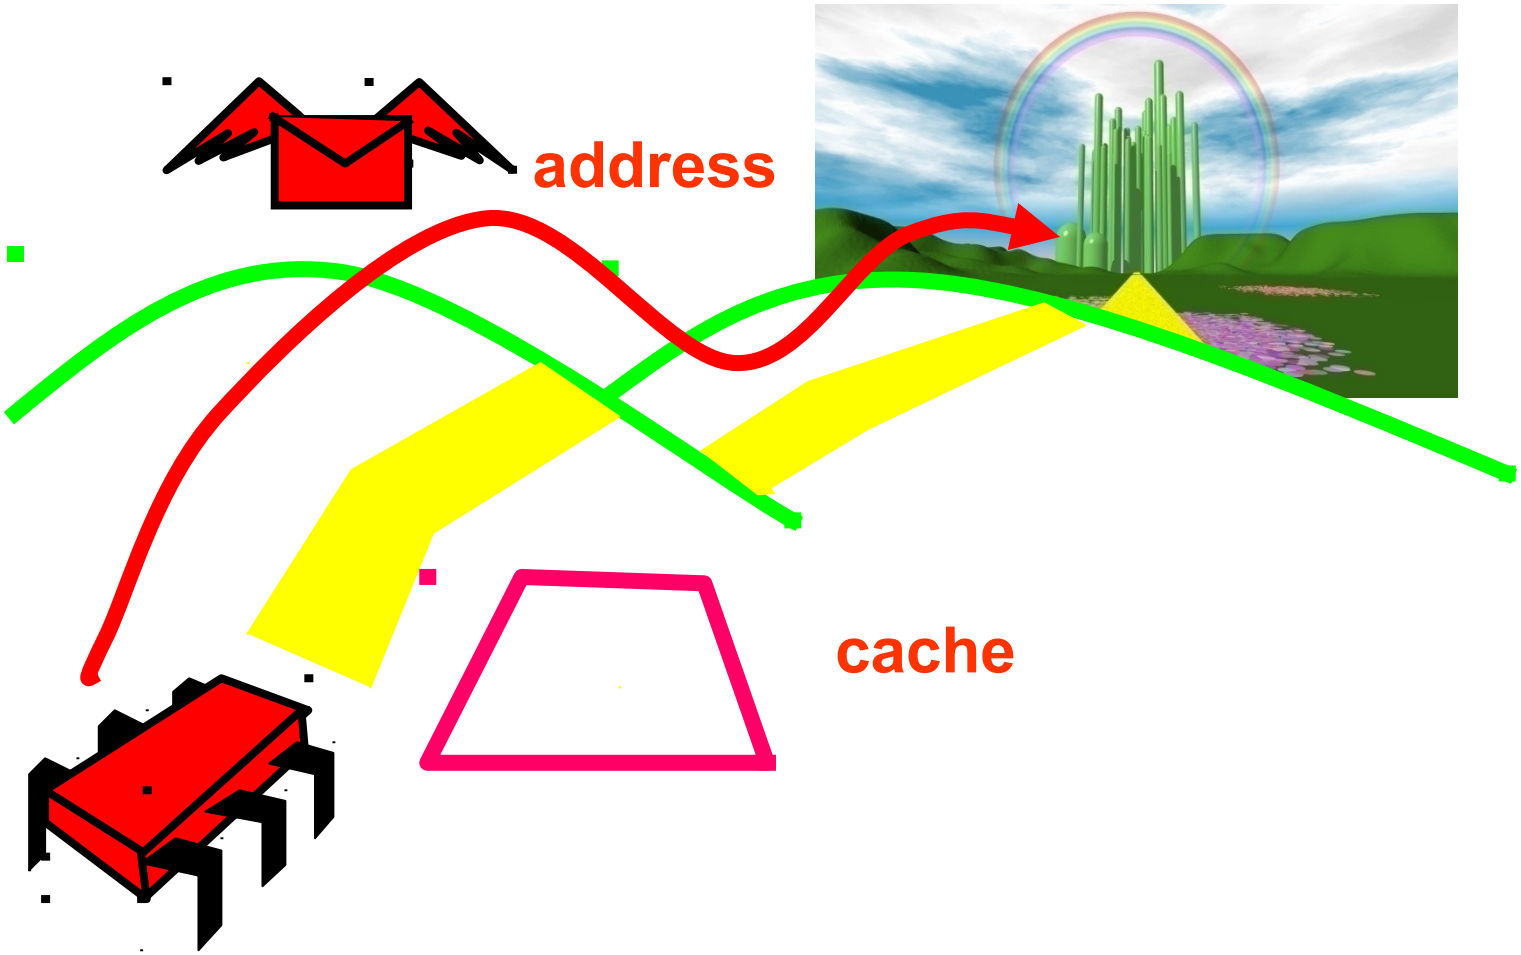

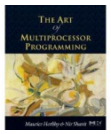

Cache: Reading from Memory

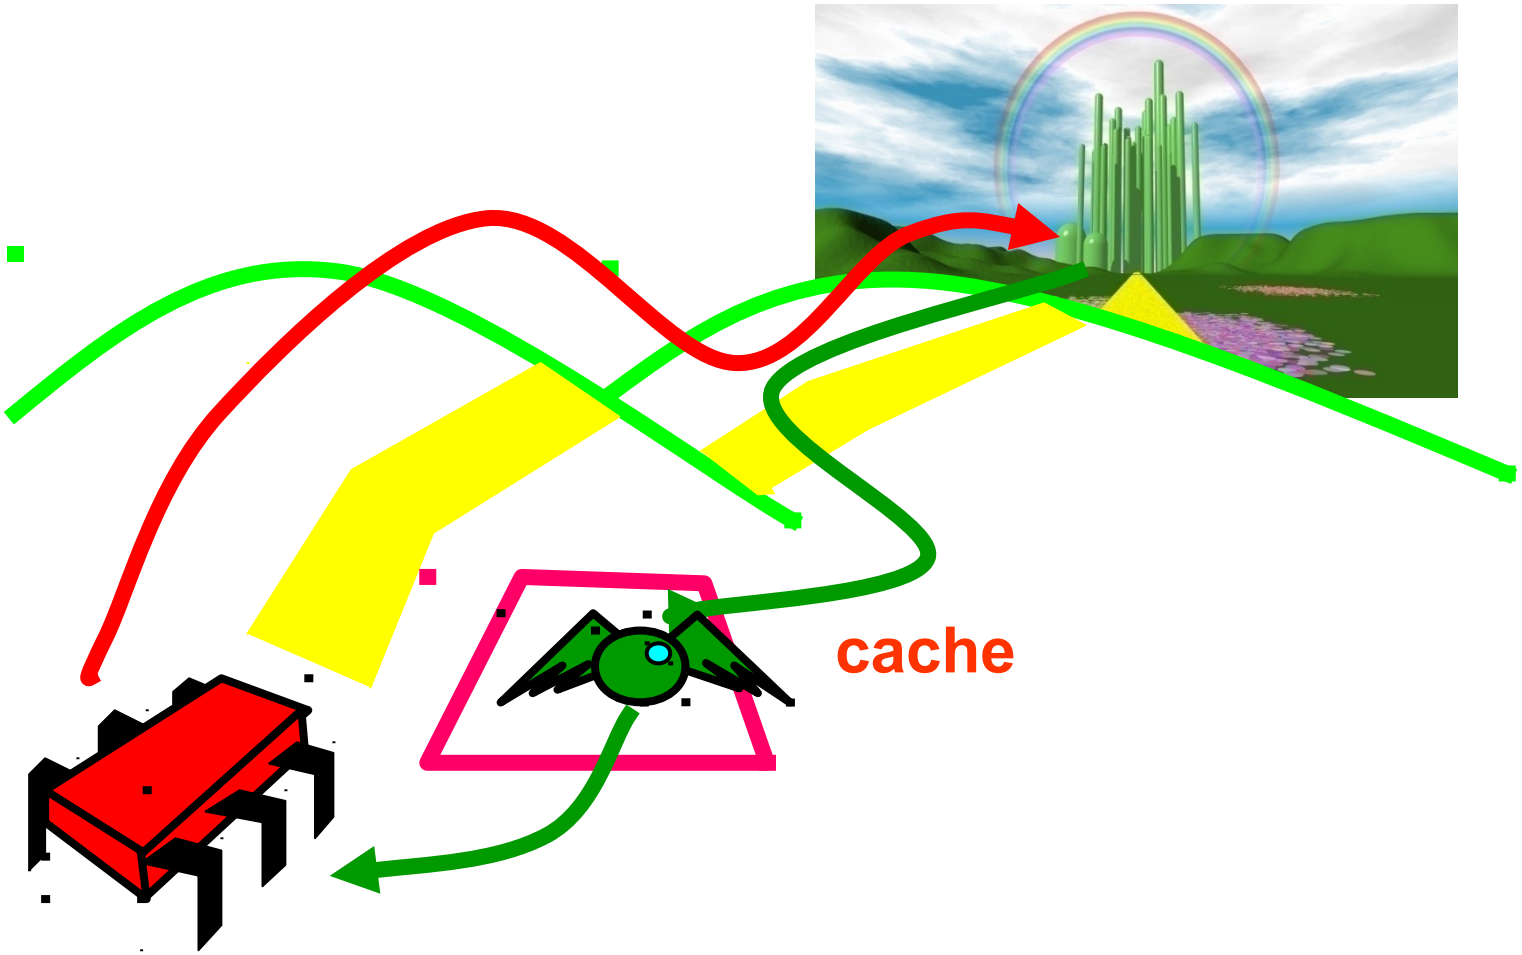

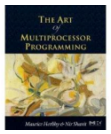

Cache: Reading from Memory

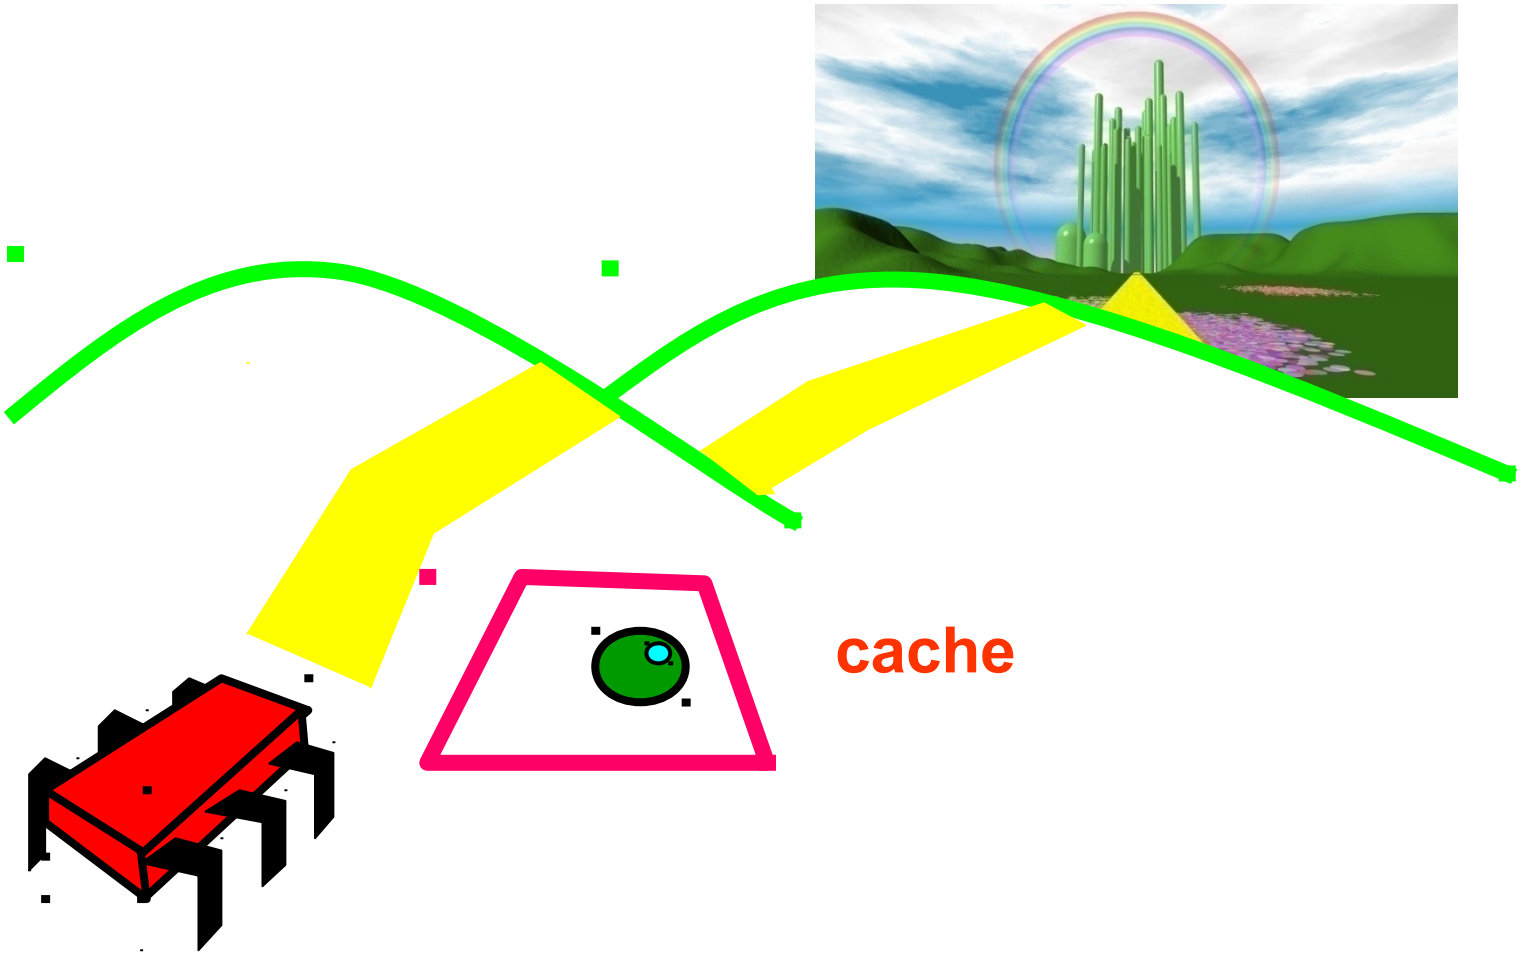

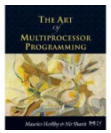

Cache Hit

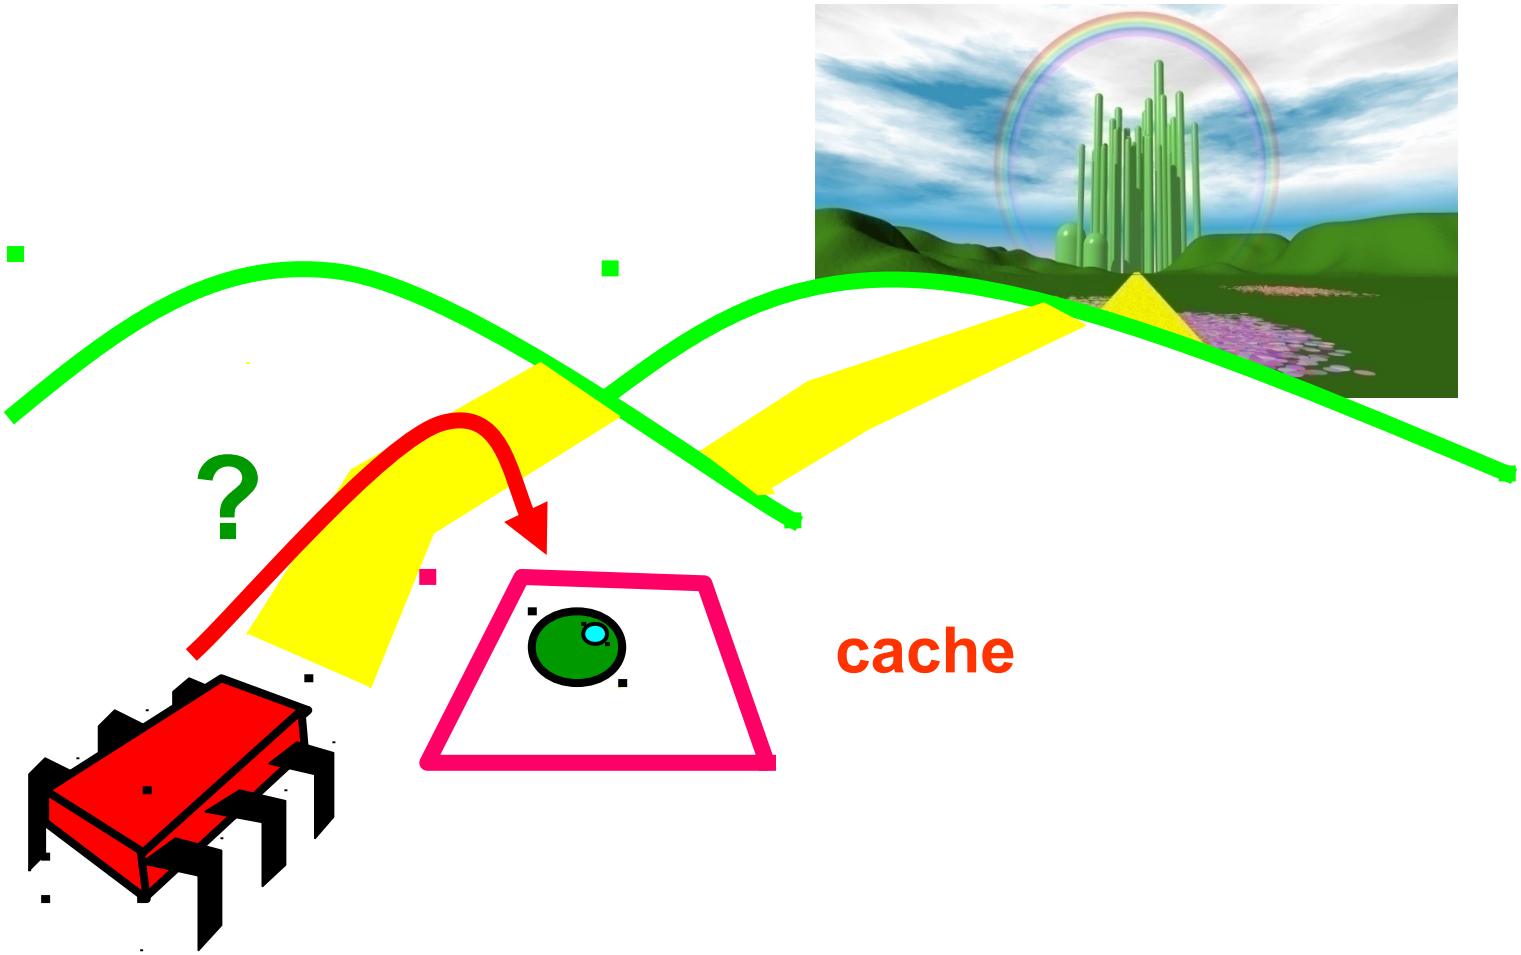

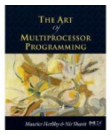

Cache Hit

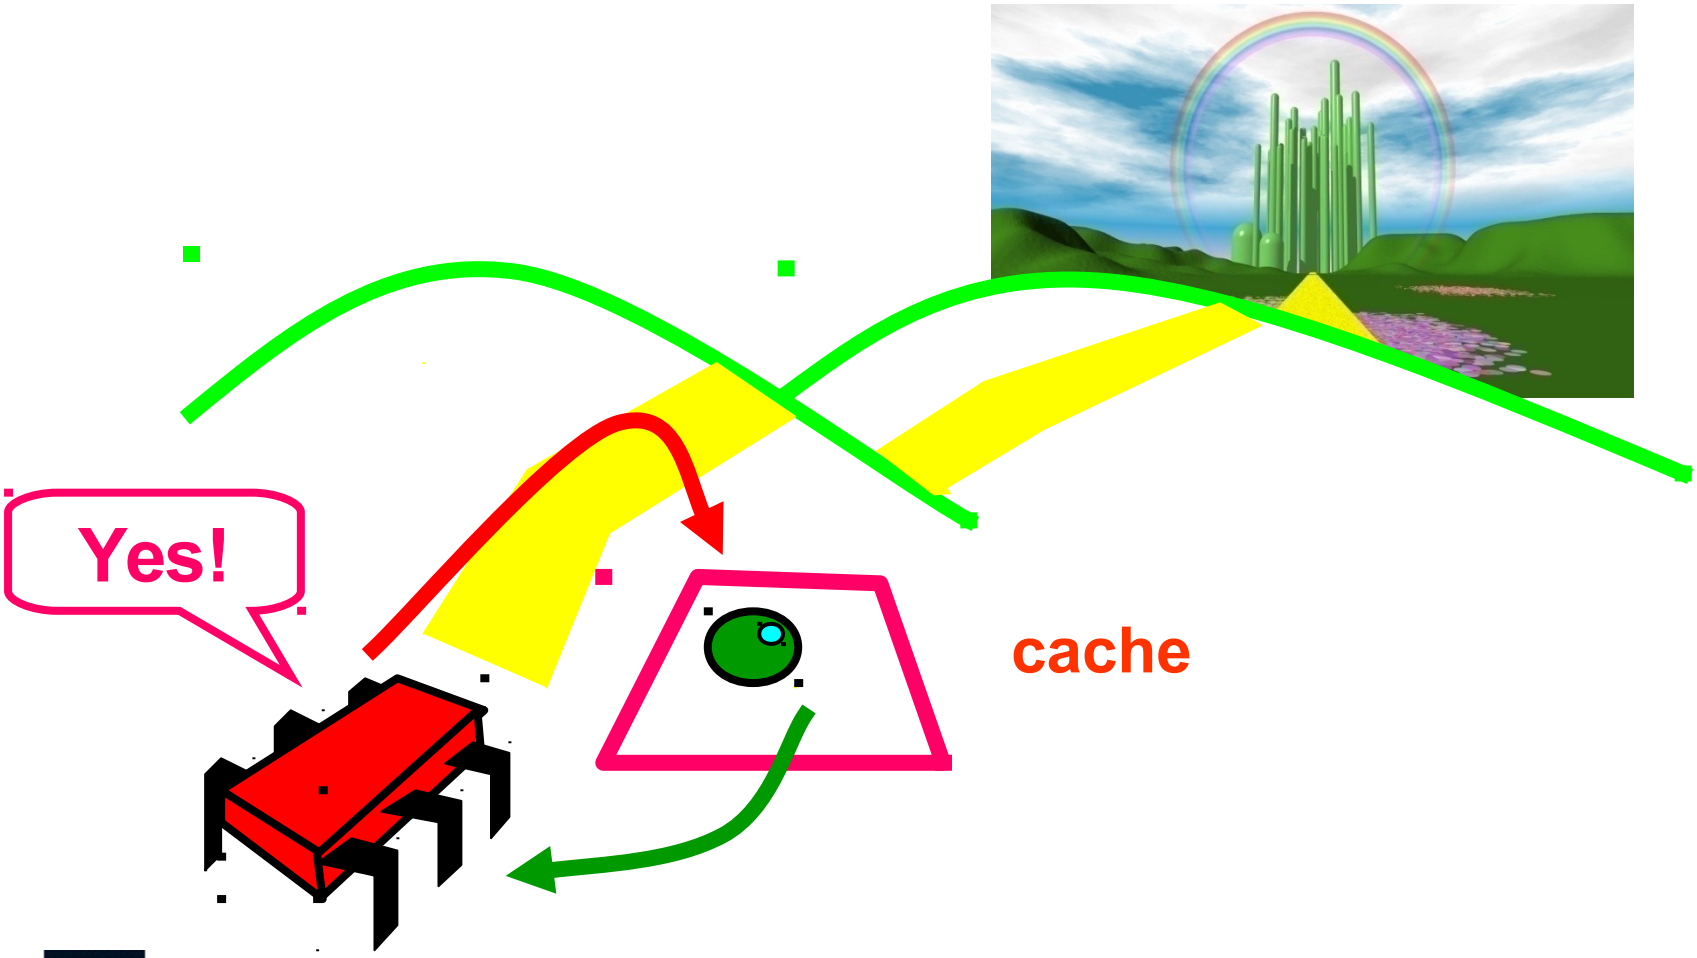

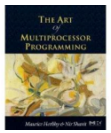

Cache Miss

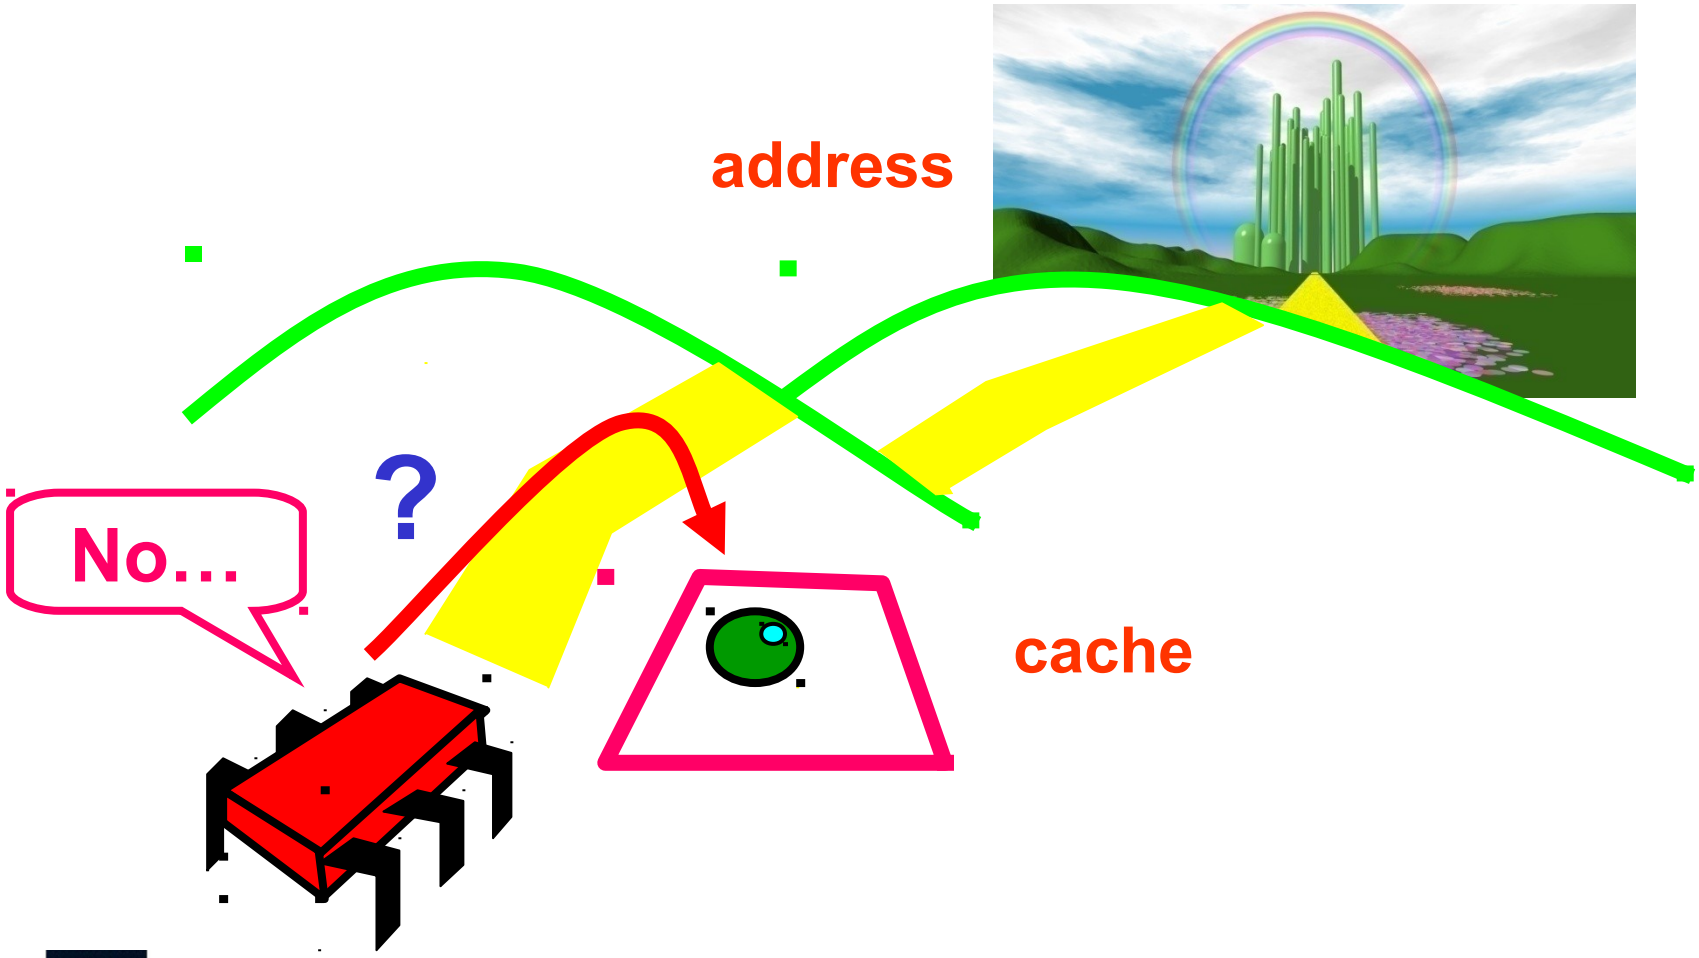

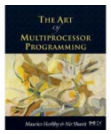

Cache Miss

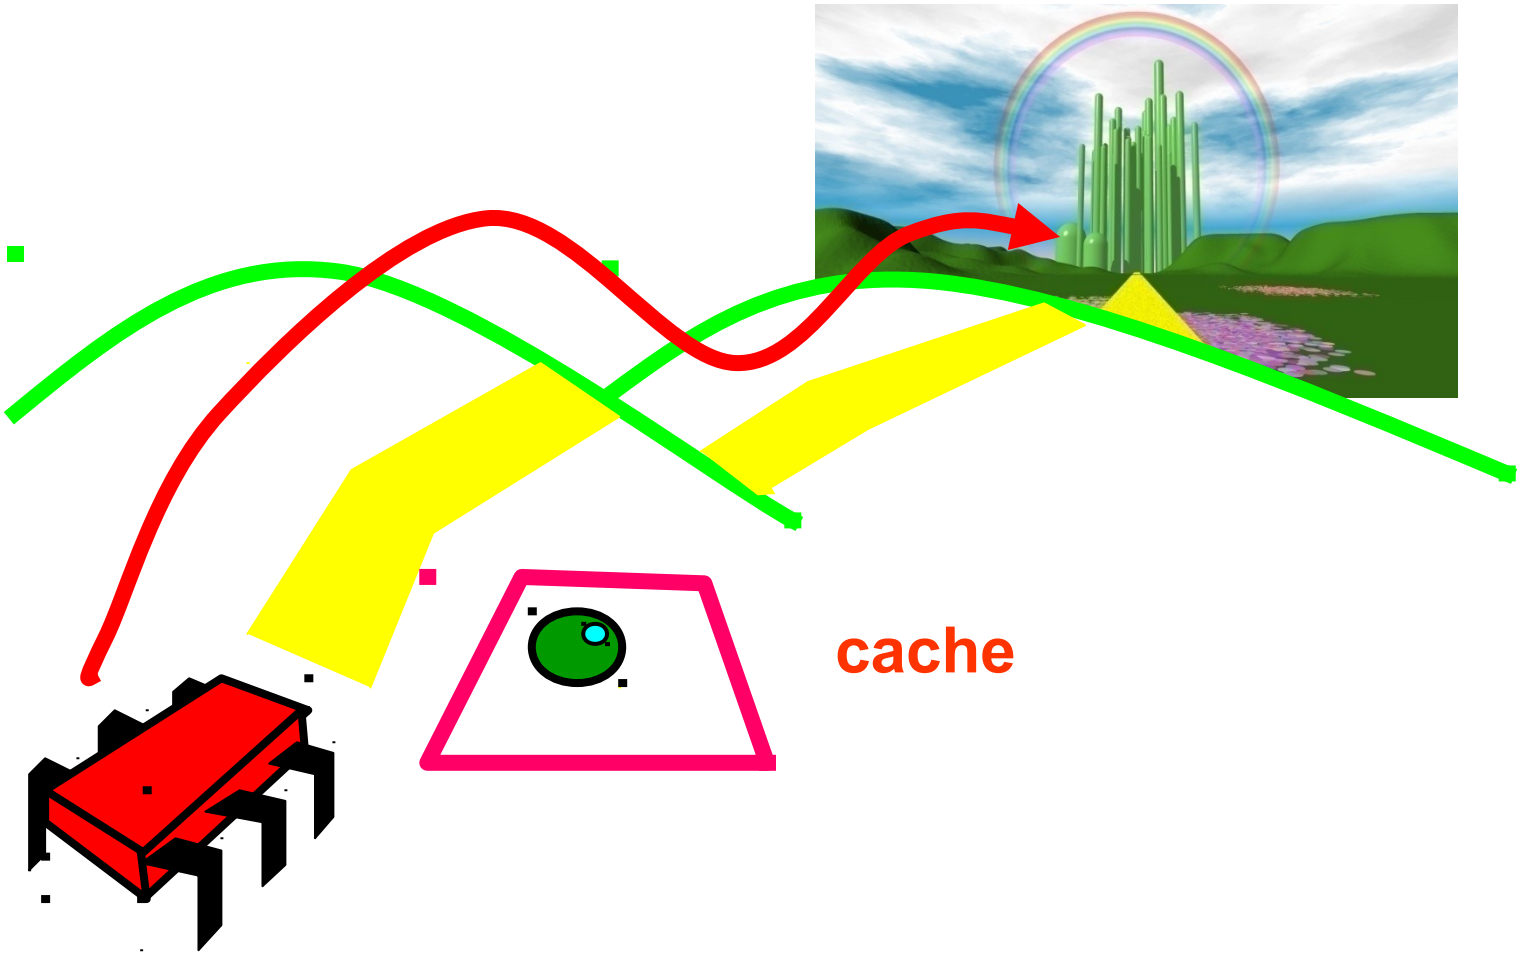

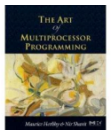

Cache Miss

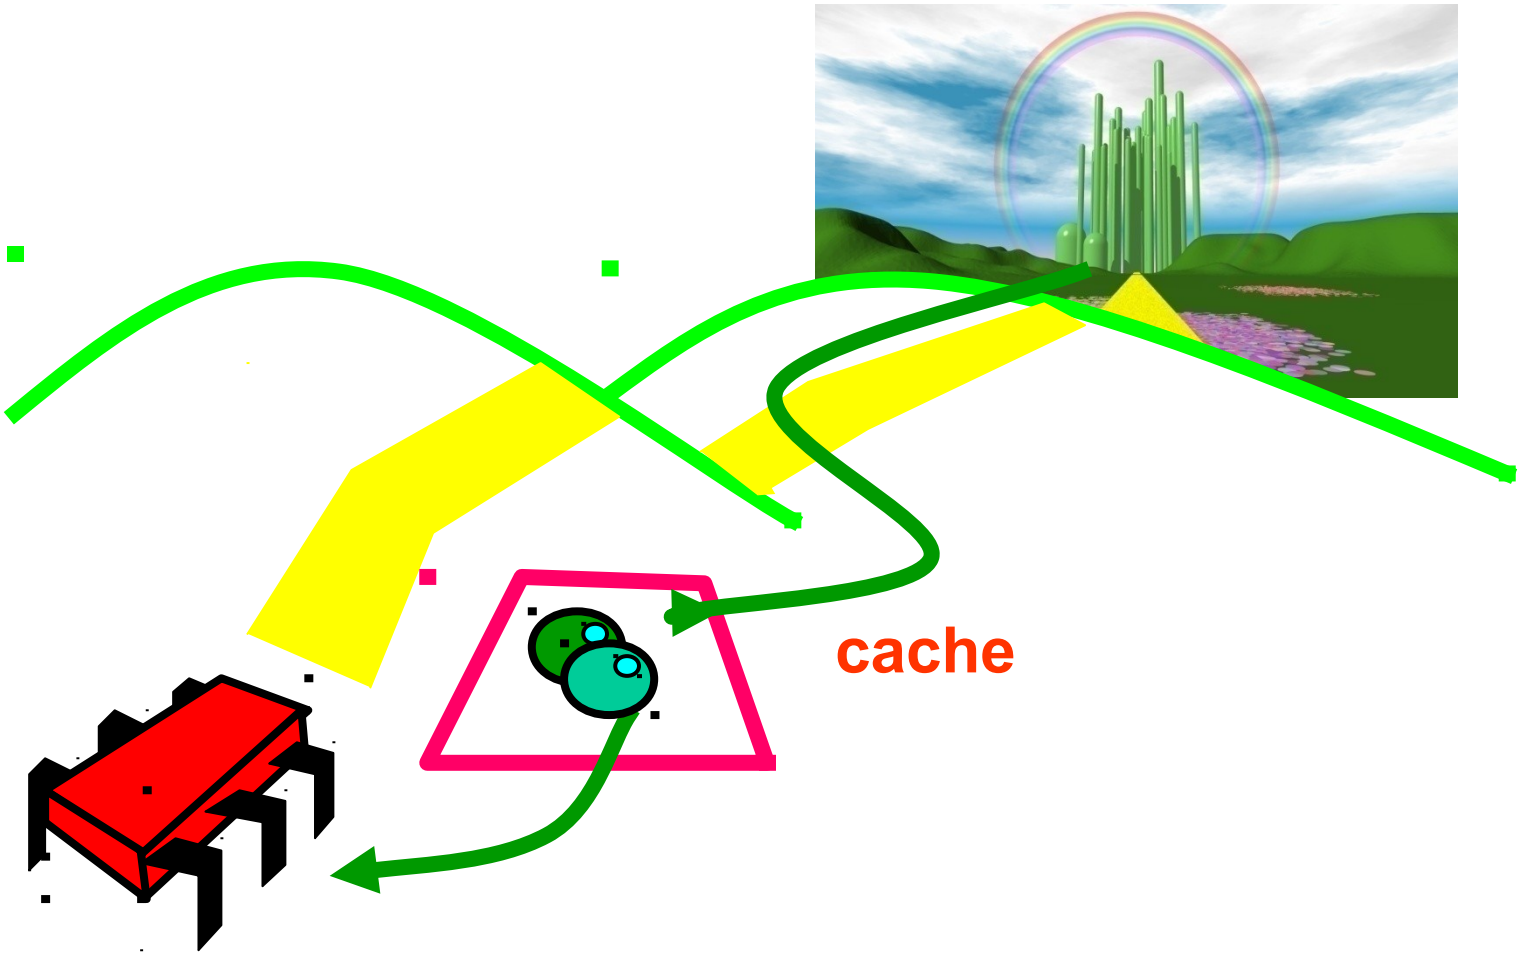

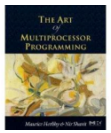

Local Spinning

- With caches, spinning becomes practical
- First time
  - Load flag bit into cache
- As long as it doesn't change
  - Hit in cache (no interconnect used)
- When it changes
  - One-time cost
  - See cache coherence below

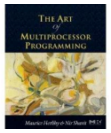

Granularity

- Caches operate at a larger granularity than a word
- Cache line: fixed-size block containing the address (today 64 or 128 bytes)

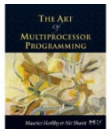

Locality

- If you use an address now, you will probably use it again soon
  - Fetch from cache, not memory
- If you use an address now, you will probably use a nearby address soon
  - In the same cache line

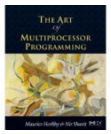

Hit Ratio

- Proportion of requests that hit in the cache
- Measure of effectiveness of caching mechanism
- Depends on locality of application

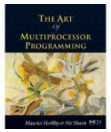

L1 and L2 Caches

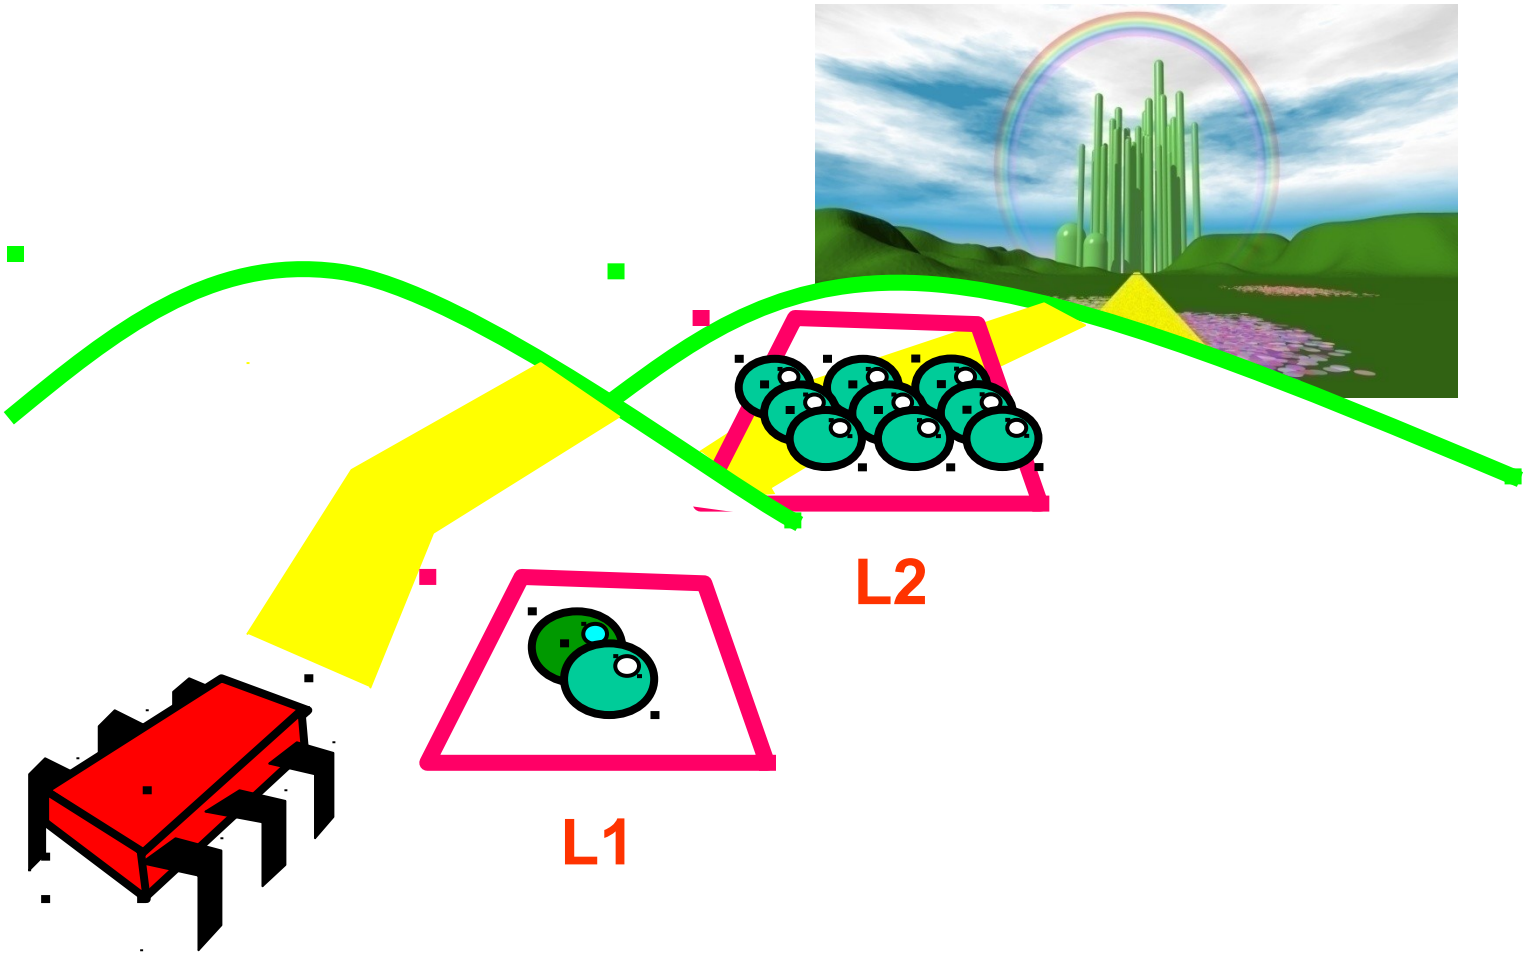

L1 and L2 Caches

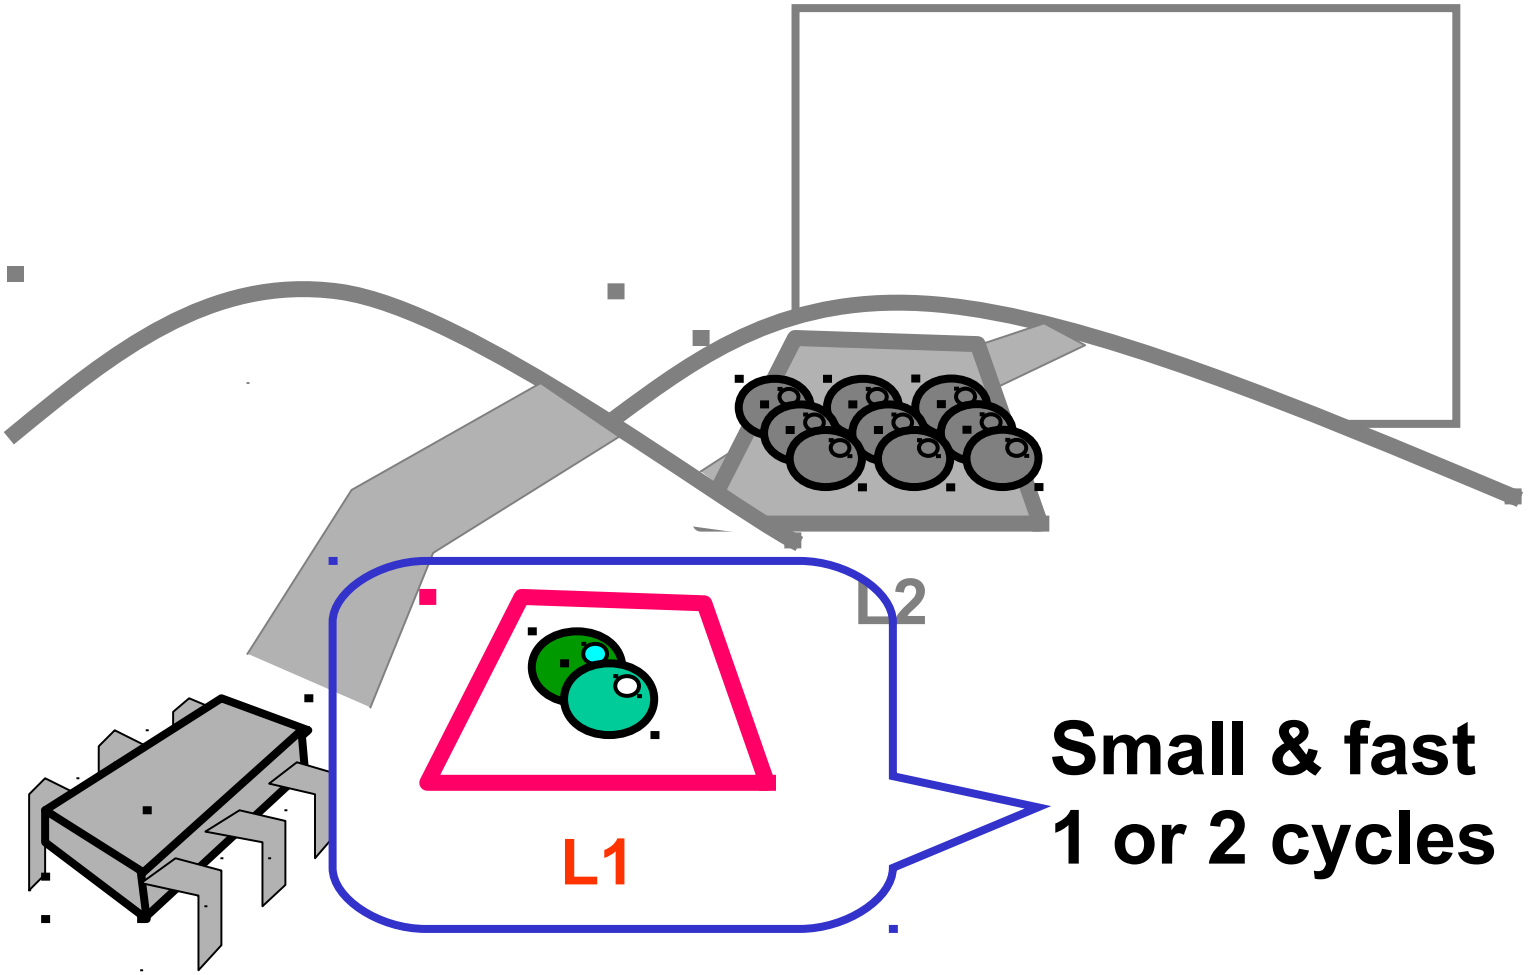

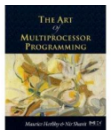

L1 and L2 Caches

Larger and slower
10s of cycles
~128 byte line

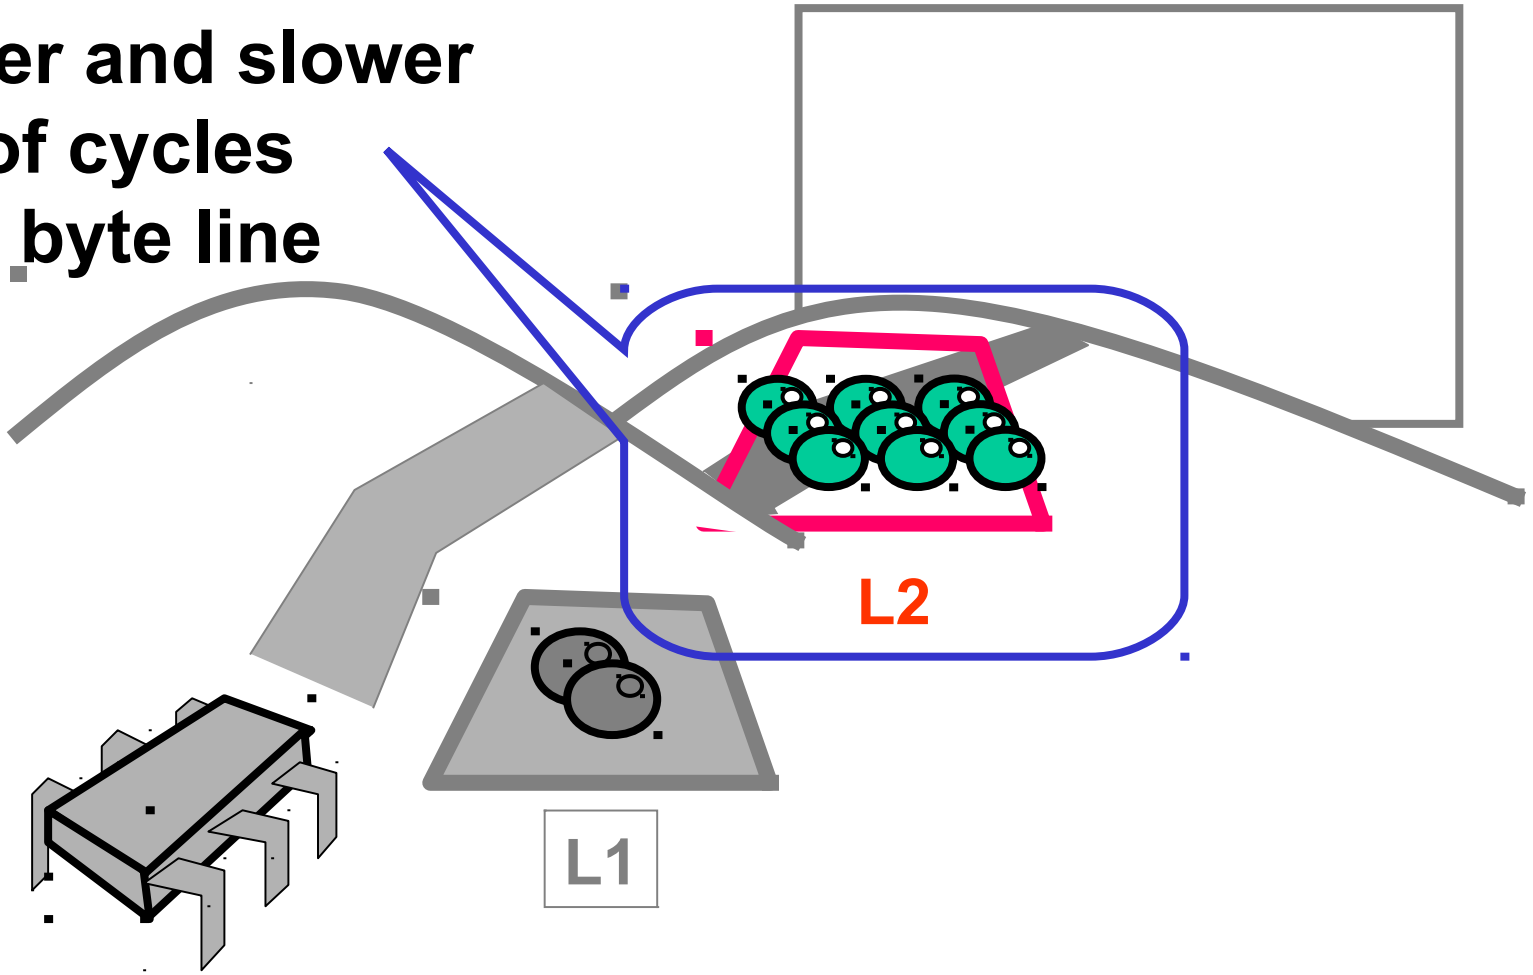

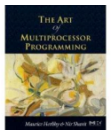

When a Cache Becomes Full…

- Need to make room for new entry
- By evicting an existing entry
- Need a replacement policy
  - Usually some kind of least recently used heuristic

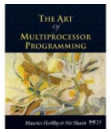

Fully Associative Cache

- Any line can be anywhere in the cache
  - Advantage: can replace any line
  - Disadvantage: hard to find lines

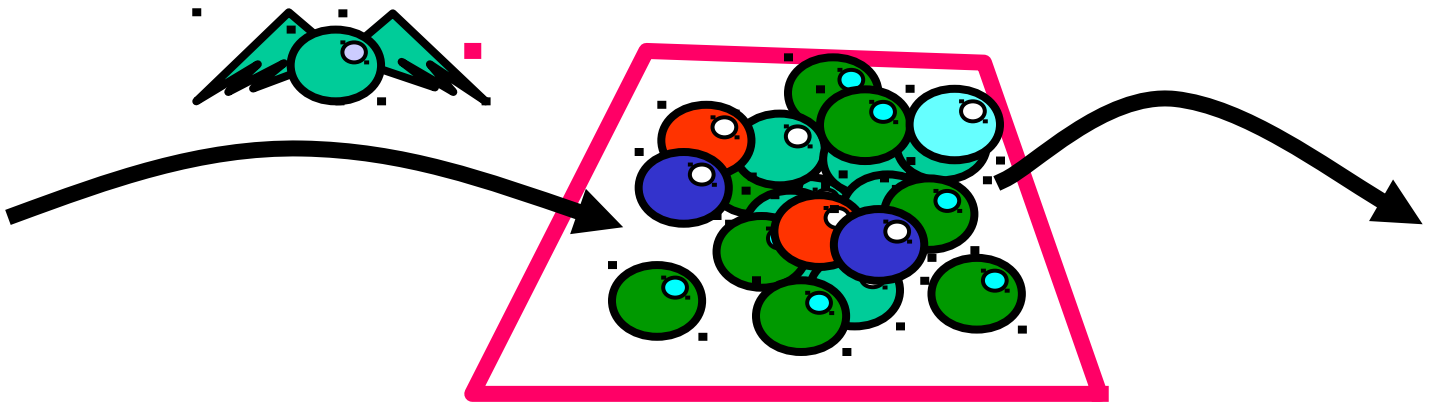

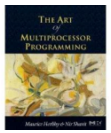

Direct Mapped Cache

- Every address has exactly 1 slot
  - Advantage: easy to find a line
  - Disadvantage: must replace fixed line

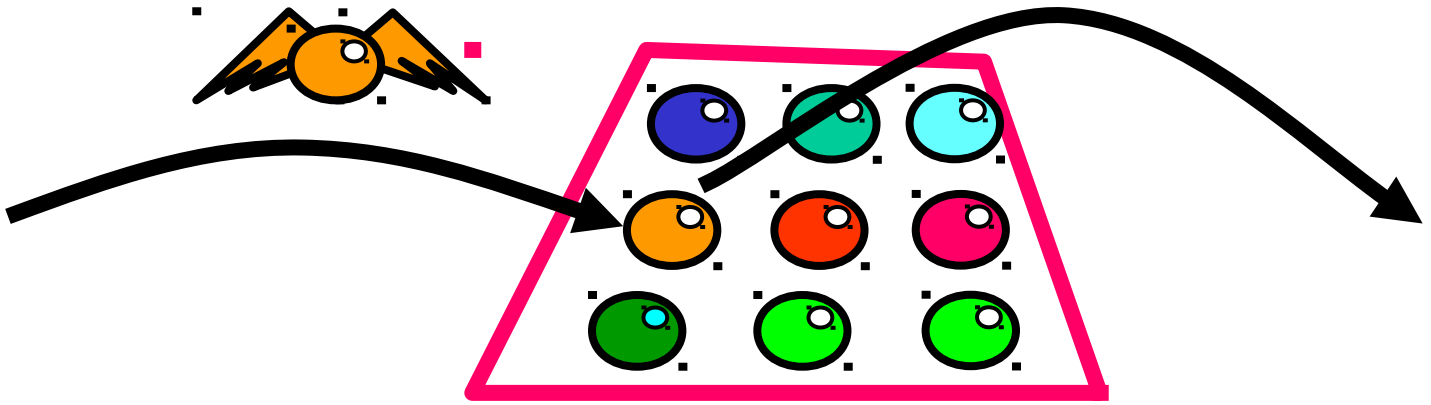

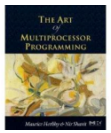

K-way Set Associative Cache

- Each slot holds k lines
  - Advantage: pretty easy to find a line
  - Advantage: some choice in replacing line

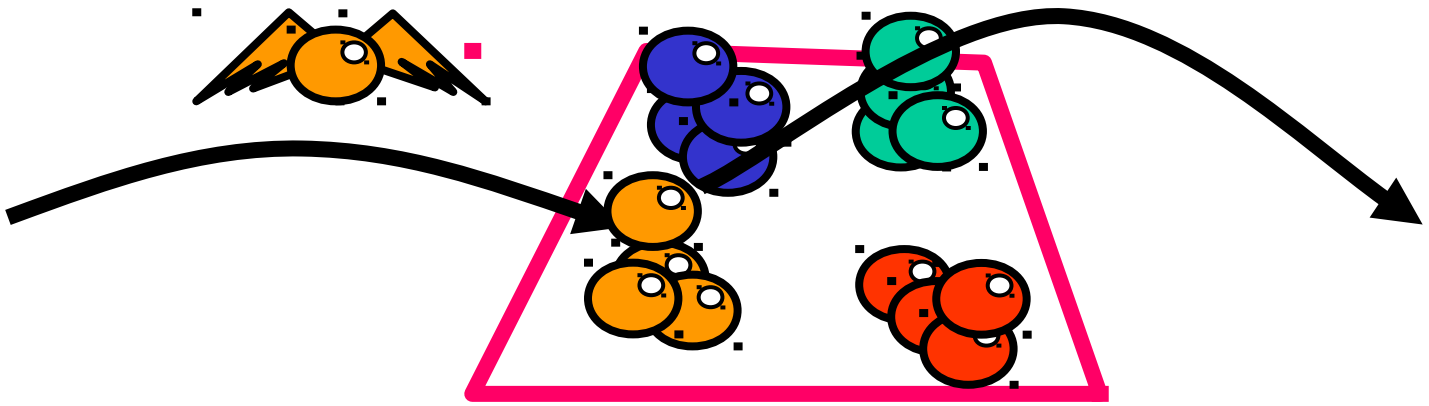

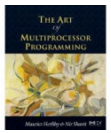

Multicore Set Associativity

- k is 8 or even 16 and growing…
  - Why? Because cores share sets
  - Threads cut effective size if accessing different data

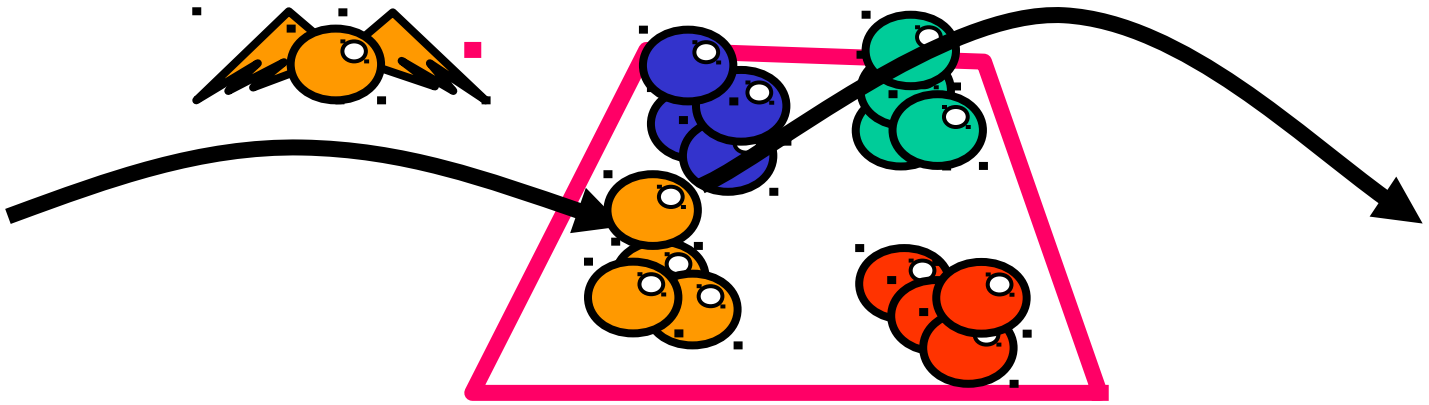

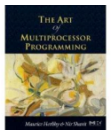

Cache Coherence

- A and B both cache address x
- A writes to x
  - Updates cache
- How does B find out?
- Many cache coherence protocols in literature

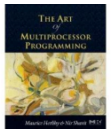

MESI

- Modified
  - Have modified cached data, must write back to memory

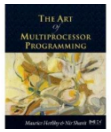

MESI

- Modified
  - Have modified cached data, must write back to memory
- Exclusive
  - Not modified, I have only copy

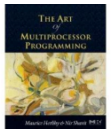

MESI

- Modified
  - Have modified cached data, must write back to memory
- Exclusive
  - Not modified, I have only copy
- Shared
  - Not modified, may be cached elsewhere

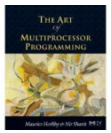

MESI

- Modified
  - Have modified cached data, must write back to memory
- Exclusive
  - Not modified, I have only copy
- Shared
  - Not modified, may be cached elsewhere
- Invalid
  - Cache contents not meaningful

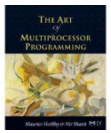

Processor Issues Load Request

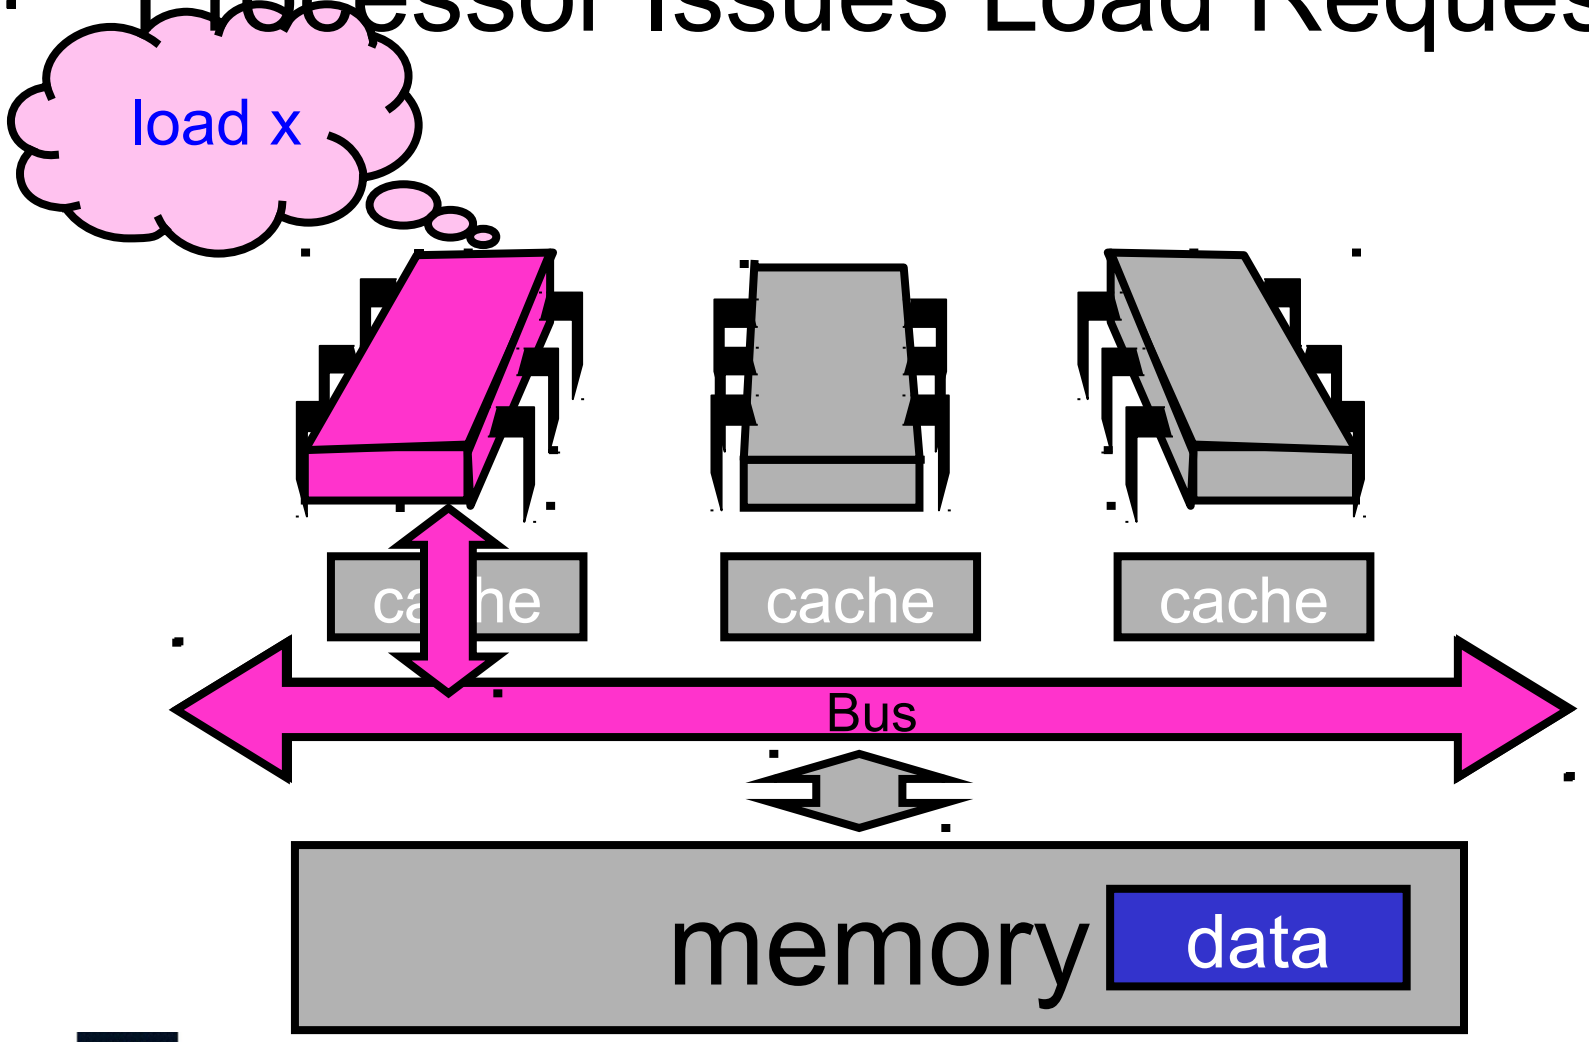

Memory Responds

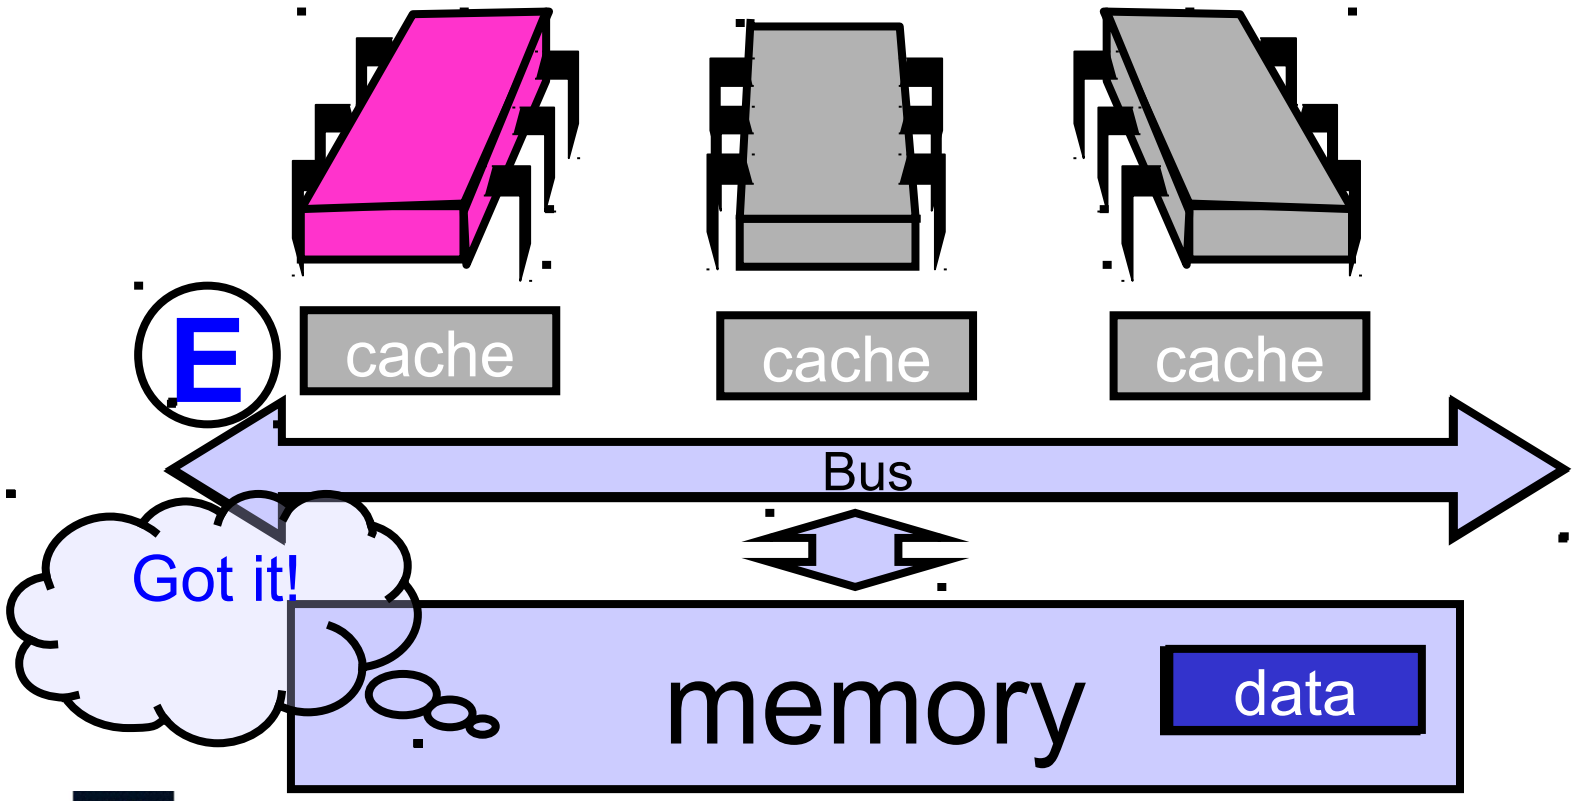

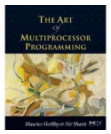

Processor Issues Load Request

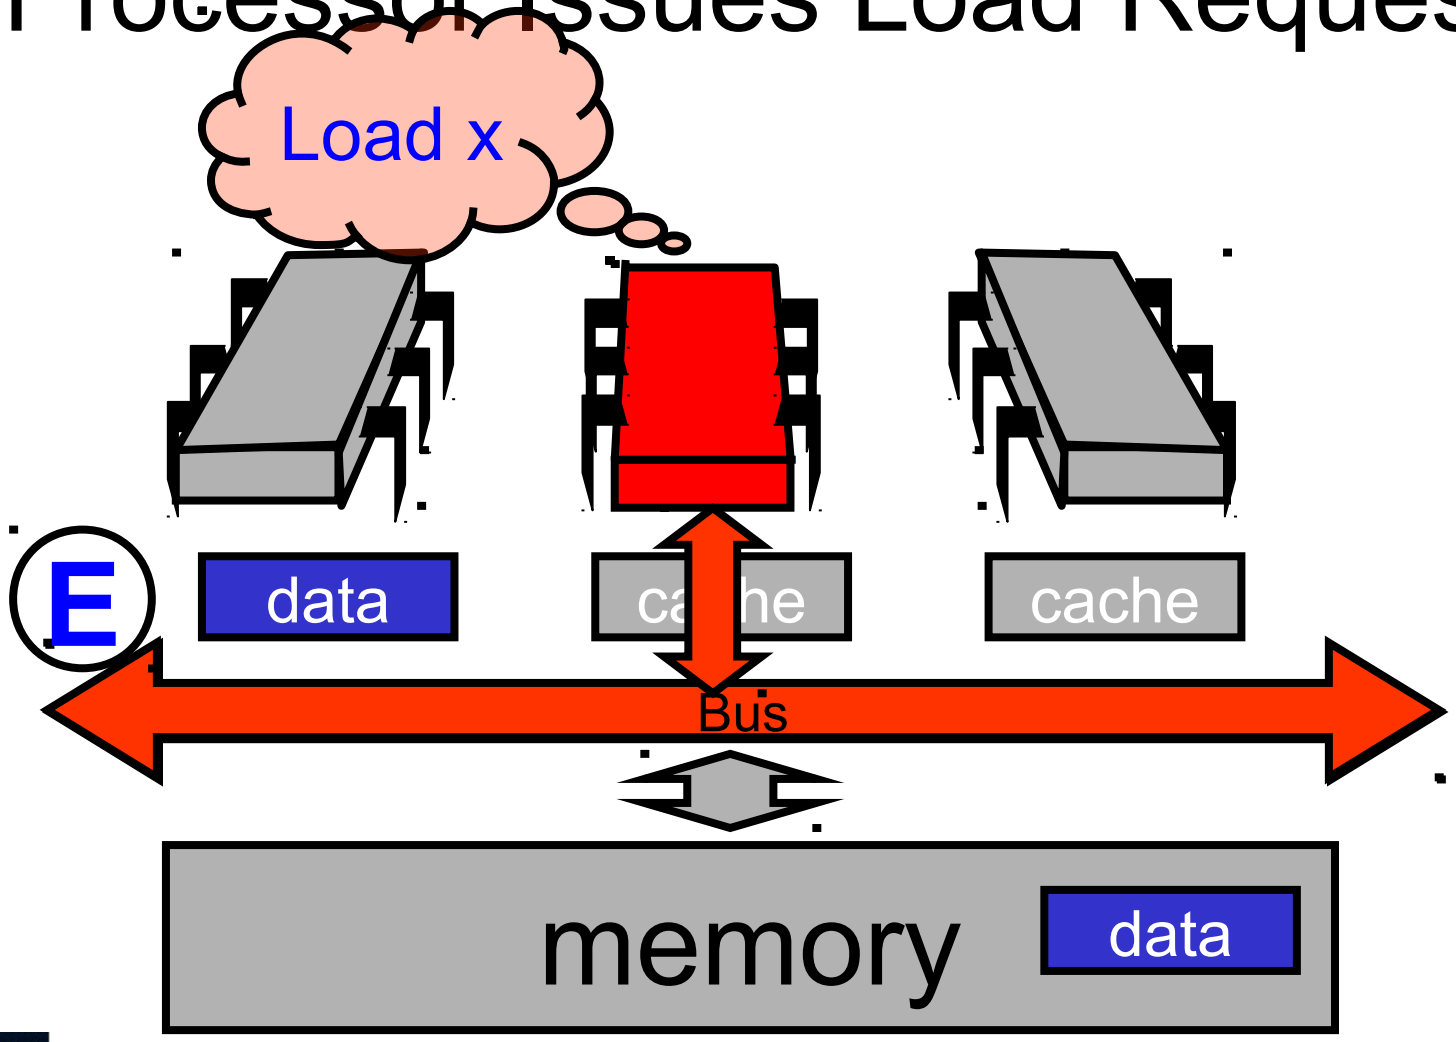

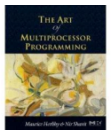

Other Processor Responds

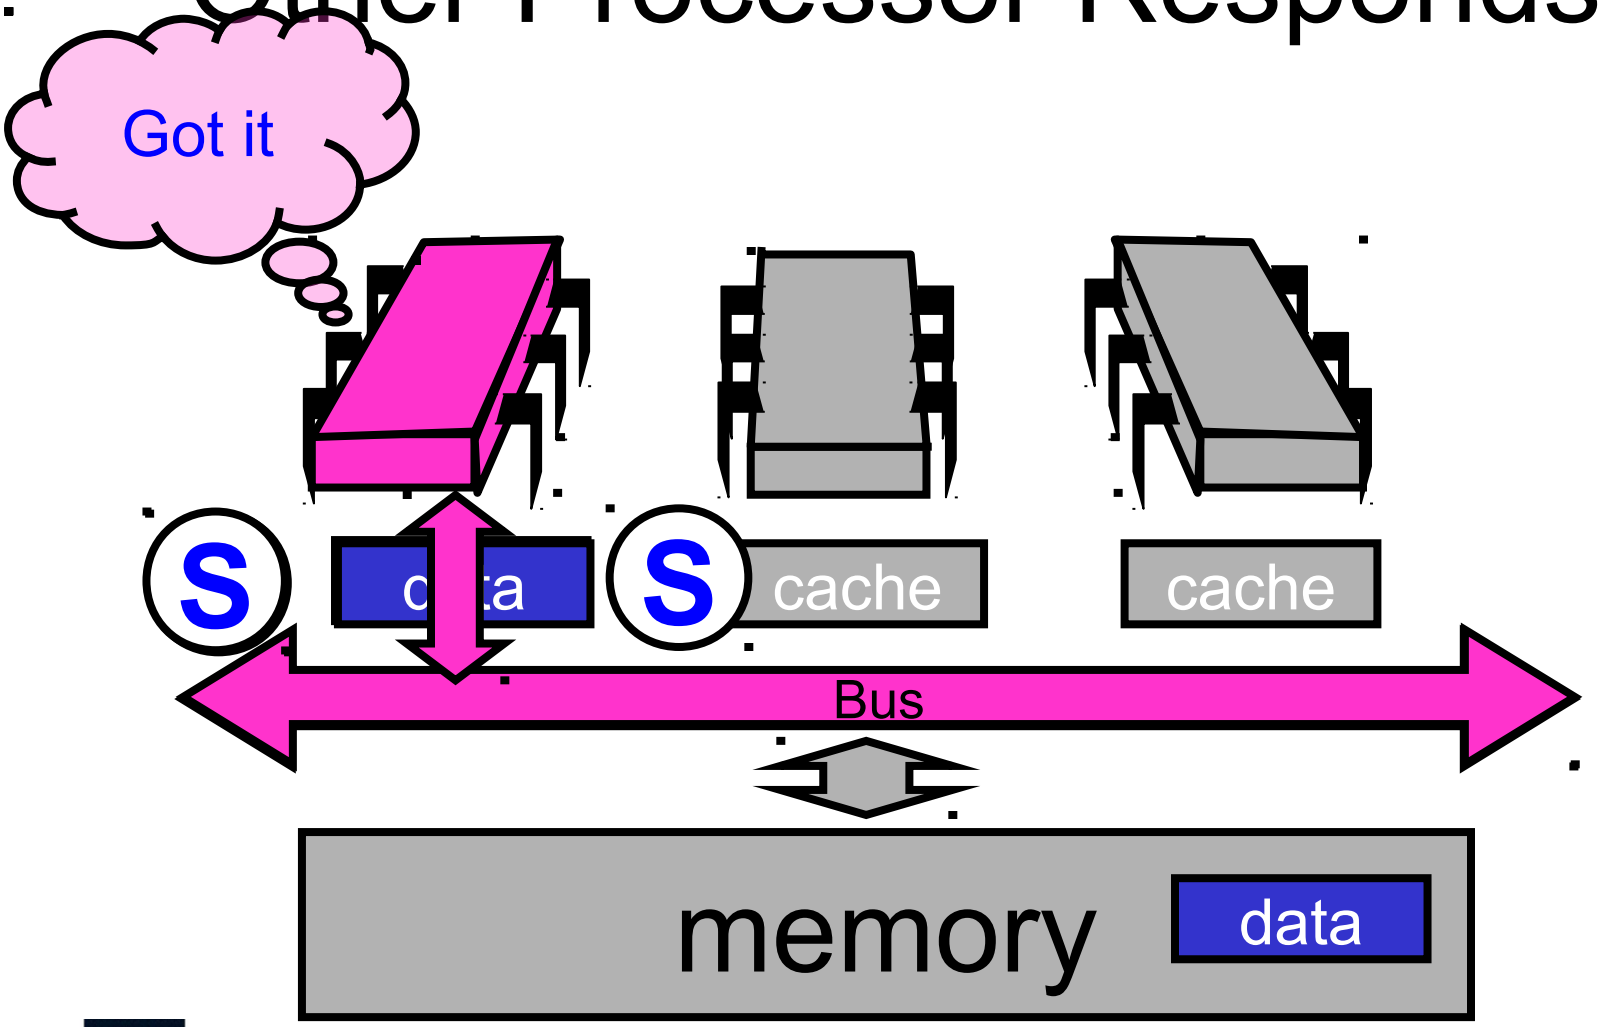

Modify Cached Data

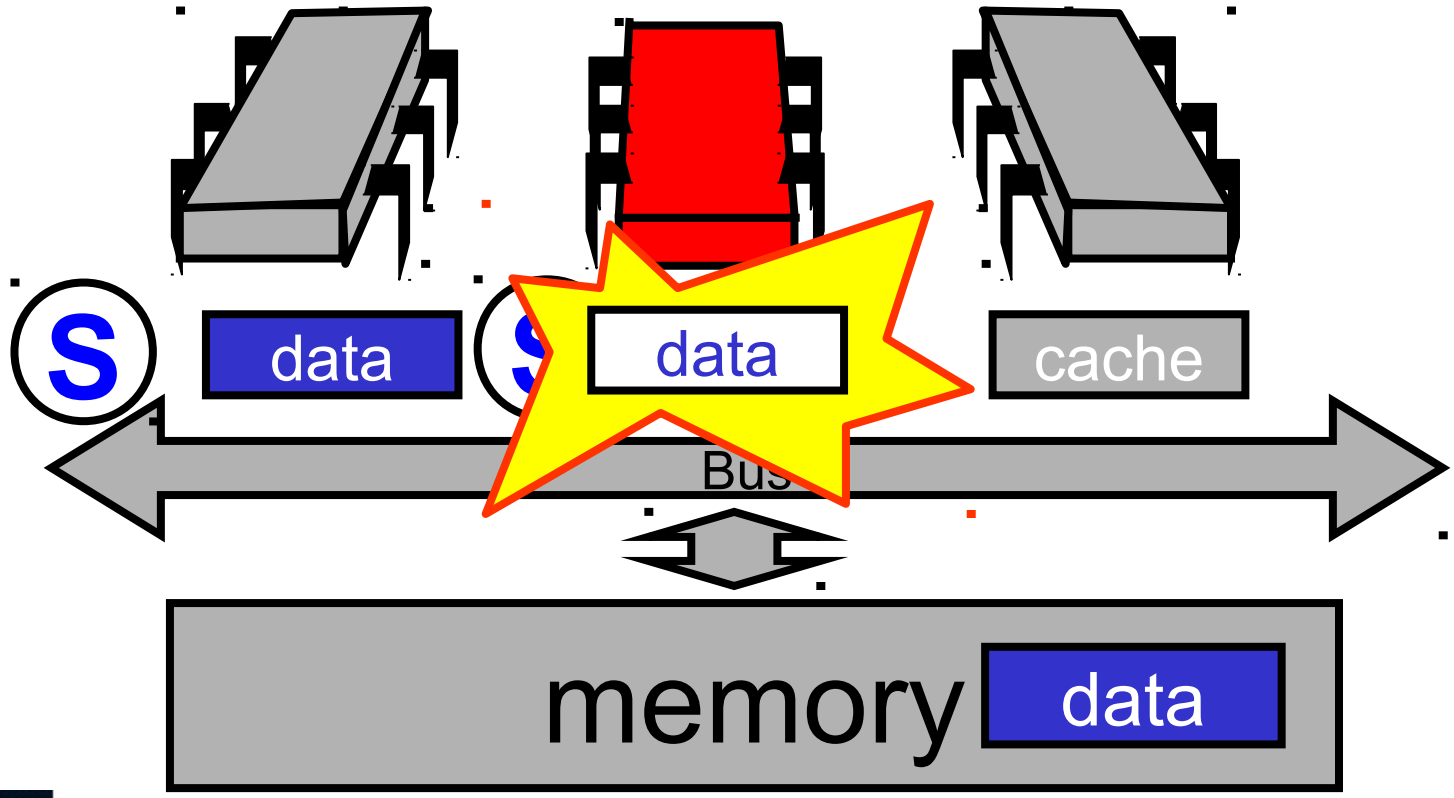

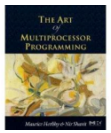

Write-Through Cache

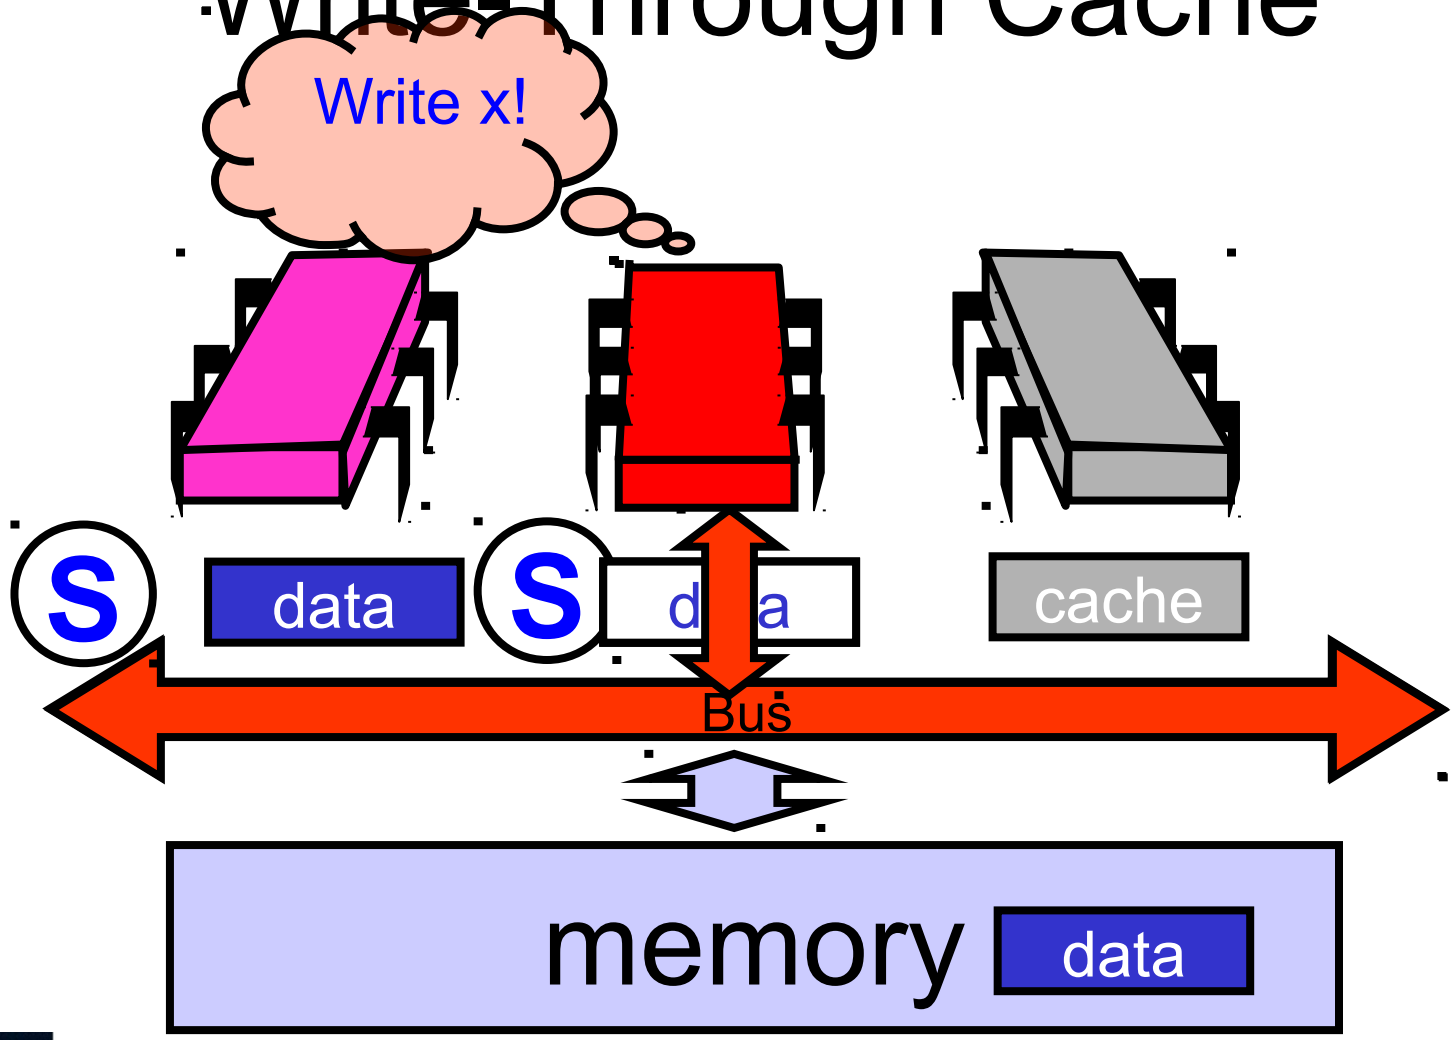

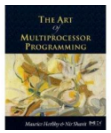

Write-Through Caches

- Immediately broadcast changes
- Good
  - Memory, caches always agree
  - More read hits, maybe
- Bad
  - Bus traffic on all writes
  - Most writes to unshared data
  - For example, loop indexes …

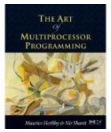

Write-Through Caches

- Immediately broadcast changes
- Good
  - Memory, caches always agree
  - More read hits, maybe
- Bad
  - Bus traffic on all writes
  - Most writes to unshared data
  - For example, loop indexes …

"show stoppers"

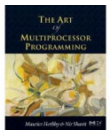

Write-Back Caches

- Accumulate changes in cache
- Write back when line evicted
  - Need the cache for something else
  - Another processor wants it

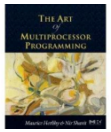

Invalidate

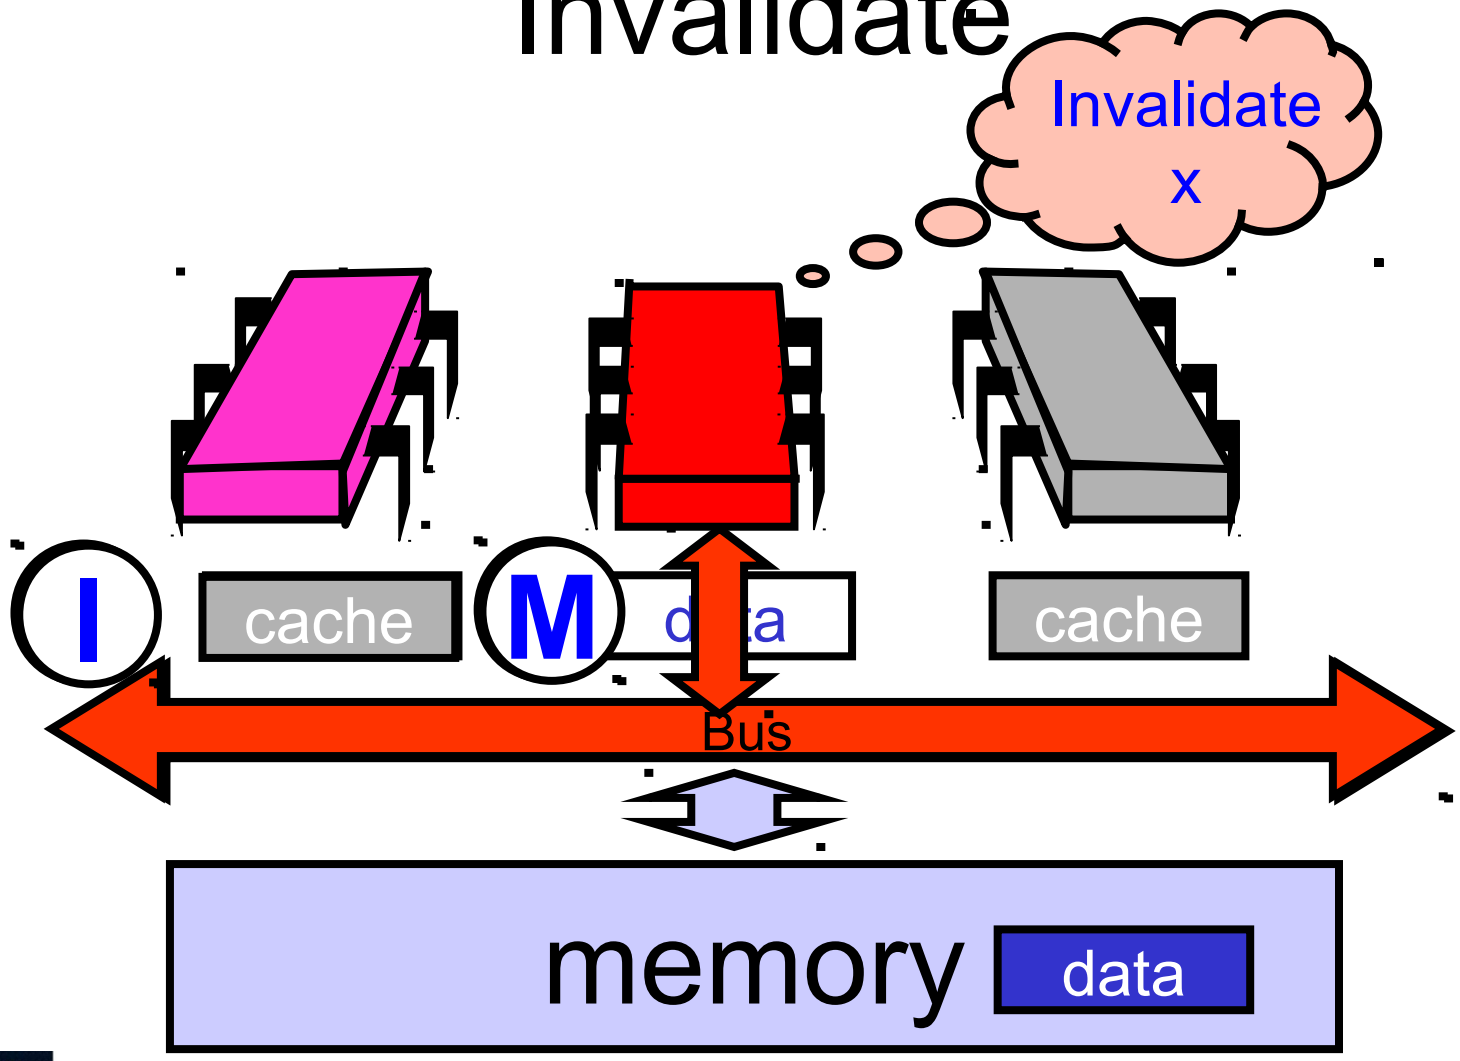

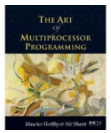

Problem

- Sometimes the compiler reorders memory operations
- Can improve
  - cache performance
  - interconnect use
- But unexpected concurrent interactions

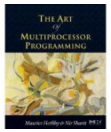

Write Buffers

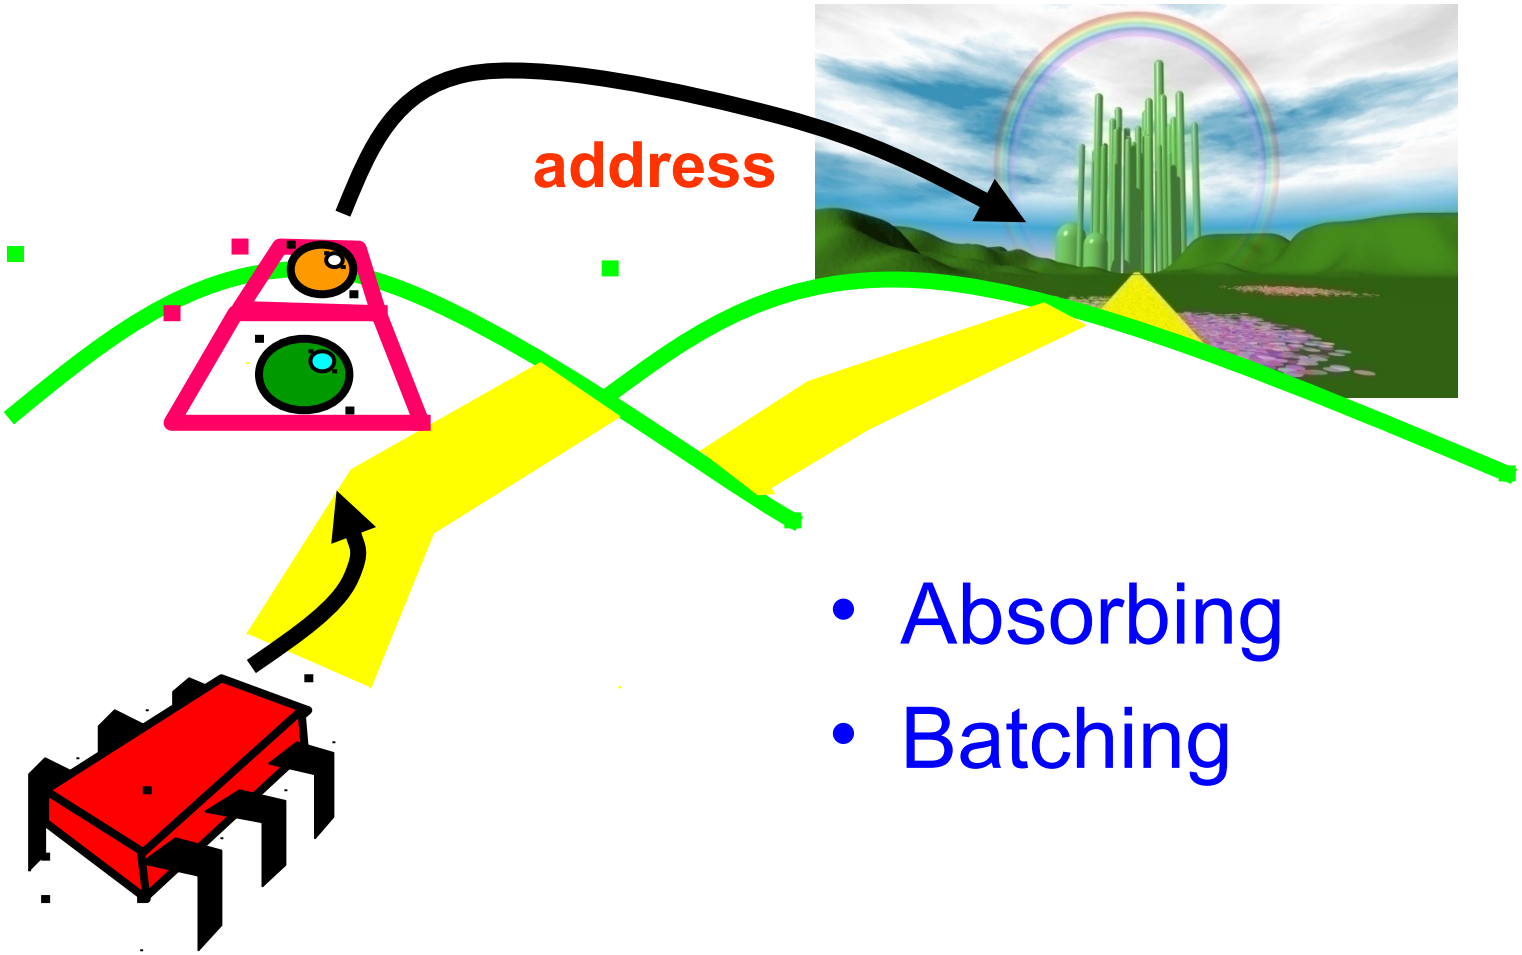

- Absorbing
- Batching

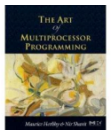

Volatile

- In Java, if a variable is declared volatile, operations won't be reordered
- Write buffer always spilled to memory before thread is allowed to continue a write
- Expensive, so use it only when needed

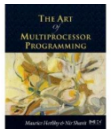

SOME RIGHTS RESERVED

This work is licensed under a
Creative Commons Attribution-ShareAlike 2.5 License.
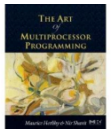
- **You are free**:
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions**:
  - **Attribution**. You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike**. If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - http://creativecommons.org/licenses/by-sa/3.0/.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.