


Mutual Exclusion




Companion slides for
The Art of Multiprocessor
Programming
by Maurice Herlihy & Nir Shavit
With some very minor changes by APS

Mutual Exclusion




- We will clarify our understanding of mutual exclusion
- We will also show you how to reason about various properties in an asynchronous concurrent setting



Mutual Exclusion

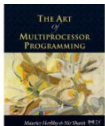


- Formal problem definitions
- Solutions for **2** threads
- Solutions for ***n*** threads
- Fair solutions
- Inherent costs




Warning

- You will never use these protocols
 - Get over it
- You are advised to understand them
 - The same issues show up everywhere
 - Except hidden and more complex





Why is Concurrent Programming so Hard?

- Try preparing a seven-course banquet
 - By yourself
 - With one friend
 - With twenty-seven friends ...
- Before we can talk about programs
 - Need a language
 - Describing time and concurrency





Events

- An **event** a_0 of thread **A** is
 - Instantaneous
 - No simultaneous events (break ties)




Threads

- A *thread* A is (formally) a sequence a_0, a_1, \dots of events
 - “Trace” model
 - Notation: $a_0 \rightarrow a_1$ indicates order





Example Thread Events

- Assign to shared variable
- Assign to local variable
- Invoke method
- Return from method
- Lots of other things ...



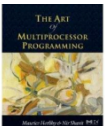
Threads are State Machines

Events are transitions



States

- **Thread State**
 - Program counter
 - Local variables
- **System state**
 - Object fields (shared variables)
 - Union of thread states



Concurrency

- **Thread A**




Concurrency

- Thread A



- Thread B




Interleavings

- **Events of two or more threads**
 - Interleaved
 - Not necessarily independent (why?)




Intervals


- An *interval* $A_0 = (a_0, a_1)$ is
 - Time between events a_0 and a_1



Intervals may Overlap





Intervals may be Disjoint




Precedence

Interval A_0 precedes interval B_0




Precedence




- **Notation:** $A_0 \rightarrow B_0$
- **Formally,**
 - End event of A_0 before start event of B_0
 - Also called “happens before” or “precedes”


Precedence Ordering



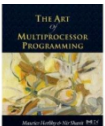
- Remark: $A_0 \rightarrow B_0$ is just like saying
 - 1066 AD \rightarrow 1492 AD,
 - Middle Ages \rightarrow Renaissance,
- Oh wait,
 - what about this week vs this month?



Precedence Ordering



- Never true that $A \rightarrow A$
- If $A \rightarrow B$ then not true that $B \rightarrow A$
- If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$
- Funny thing: $A \rightarrow B$ & $B \rightarrow A$ might both be false!



Partial Orders

(review)


- **Irreflexive:**
 - Never true that $A \rightarrow A$
- **Antisymmetric:**
 - If $A \rightarrow B$ then not true that $B \rightarrow A$
- **Transitive:**
 - If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$



Total Orders

(review)

- Also
 - Irreflexive
 - Antisymmetric
 - Transitive
- Except that for every distinct A, B ,
Either $A \rightarrow B$ or $B \rightarrow A$ (Trichotomy)



Repeated Events


```
while (mumble) {  
    a0; a1;  
}
```

**k -th occurrence of
event a_0**

a_0^k

**k -th occurrence of
interval $A_0 = (a_0, a_1)$**

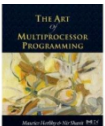
A_0^k



Implementing a Counter

```
public class Counter {  
    private long value;  
  
    public long getAndIncrement() {  
        temp = value;  
        value = temp + 1;  
        return temp;  
    }  
}
```


**Make these steps
indivisible using
locks**



Locks (Mutual Exclusion)

```
public interface Lock {  
  
    public void lock();  
  
    public void unlock();  
}
```

Note: `java.util.concurrent.locks` has more methods



Locks (Mutual Exclusion)


```
public interface Lock {
```

```
    public void lock();
```

acquire lock

```
    public void unlock();
```

```
}
```



Locks (Mutual Exclusion)

```
public interface Lock {
```


```
    public void lock();
```

acquire lock

```
    public void unlock();
```


release lock

```
}
```



Using Locks


```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```



Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```


acquire Lock



Using Locks

```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```


**Release lock
(no matter what)**



Using Locks


```
public class Counter {  
    private long value;  
    private Lock lock;  
    public long getAndIncrement() {  
        lock.lock();  
        try {  
            int temp = value;  
            value = value + 1;  
        } finally {  
            lock.unlock();  
        }  
        return temp;  
    }  
}
```

**critical
section**






Desired 1: Mutual Exclusion

- Let $CS_i^k \leftrightarrow$ be thread i 's k -th critical section execution









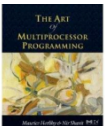
Desired 1: Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be thread j's m-th critical section execution









Desired 1: Mutual Exclusion


- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be j's m-th execution
- Then either
—   or  









Desired 1: Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be j's m-th execution
- Then either
—   or  

$CS_i^k \rightarrow CS_j^m$




Desired 1: Mutual Exclusion

- Let CS_i^k  be thread i's k-th critical section execution
- And CS_j^m  be j's m-th execution
- Then either
—   or  


$CS_i^k \rightarrow CS_j^m$

$CS_j^m \rightarrow CS_i^k$




Desired 2: Deadlock-Free

- If **ANY** thread calls `lock()`, then **SOME** thread will acquire it
- If some thread calls `lock()`
 - And never returns
 - Then other threads must be completing `lock()` and `unlock()` calls infinitely often
- System as a whole makes progress
 - Even if individuals starve



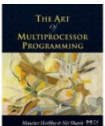
Desired 3: Starvation-Free

- If some thread calls `lock()`
 - It will eventually return
- Individual threads make progress




Two-Thread vs n -Thread Solutions

- **2-thread solutions first**
 - Illustrate most basic ideas
 - Fits on one slide
- **Then n -thread solutions**



Two-Thread Conventions


```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
    ...  
    }  
}
```



Two-Thread Conventions


```
class ... implements Lock {  
    ...  
    // thread-local index, 0 or 1  
    public void lock() {  
        int i = ThreadID.get();  
        int j = 1 - i;  
    }  
}
```

Henceforth: **i** is current thread, **j** is other thread



LockOne

```
class LockOne implements Lock {
private boolean[] flag = new boolean[2];
public void lock() {
    int i = ThreadId.get();
    int j = 1 - i ;
    flag[i] = true;
    while (flag[j]) {}
}
public void unlock(){
    int i = ThreadId.get();
    flag[i] = false;
}
}
```



LockOne

```
class LockOne implements Lock {
```

```
private boolean[] flag = new boolean[2];
```

```
public void lock() {
```

```
    flag[i] = true;
```

```
    while (flag[j]) {}
```

```
}
```

Each thread has flag
(initialised to false)

LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

Set my flag


LockOne

```
class LockOne implements Lock {  
    private boolean[] flag = new boolean[2];  
    public void lock() {  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```

**Wait for other flag to
become false**

LockOne Satisfies Mutual Exclusion


- Assume CS_A^j overlaps CS_B^k
- Consider each thread's last (j -th and k -th) read and write in the `lock()` method before entering
 - Note: only considering 4 events
- Derive a contradiction



From the Code

- $\text{write}_A(\text{flag}[A]=\text{true}) \rightarrow$
 $\text{read}_A(\text{flag}[B]==\text{false}) \rightarrow \text{CS}_A$
- $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$
 $\text{read}_B(\text{flag}[A]==\text{false}) \rightarrow \text{CS}_B$

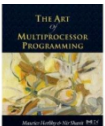
```
class LockOne implements Lock {  
    ...  
    public void lock() {  
        ...  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```



From the Assumption

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow$
 $\text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow$
 $\text{write}_A(\text{flag}[A] = \text{true})$

```
class LockOne implements Lock {  
    ...  
    public void lock() {  
        ...  
        flag[i] = true;  
        while (flag[j]) {}  
    }  
}
```




Combining

- **Assumptions:**

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- **From the code**

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$




Combining

- Assumptions:

- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$
- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$
- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$



Combining

- Assumptions:


- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\rightarrow \text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$



Combining

- Assumptions:


- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$



Combining

- Assumptions:


- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$

- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$



Combining

- Assumptions:


- $\text{read}_A(\text{flag}[B] == \text{false}) \rightarrow \text{write}_B(\text{flag}[B] = \text{true})$

- $\text{read}_B(\text{flag}[A] == \text{false}) \rightarrow \text{write}_A(\text{flag}[A] = \text{true})$



- From the code

- $\text{write}_A(\text{flag}[A] = \text{true}) \rightarrow \text{read}_A(\text{flag}[B] == \text{false})$

- $\text{write}_B(\text{flag}[B] = \text{true}) \rightarrow \text{read}_B(\text{flag}[A] == \text{false})$



Cycle!




Deadlock Freedom

- **LockOne Fails deadlock-freedom**
 - Concurrent execution can deadlock


```
flag[i] = true;    flag[j] = true;  
while (flag[j]){  while (flag[i]){}
```

- Sequential executions OK



LockTwo


```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```



LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

**Let other go
first**



LockTwo

```
public class LockTwo implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {};  
    }  
  
    public void unlock() {}  
}
```

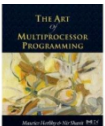
Wait for permission



LockTwo

```
public class Lock2 implements Lock {  
    private int victim;  
    public void lock() {  
        victim = i;  
        while (victim == i) {}  
    }  
    public void unlock() {}  
}
```

Nothing to do



LockTwo Claims

- Satisfies mutual exclusion


- If thread `i` in CS
- Then `victim == j`
- Cannot be both 0 and 1

```
public void LockTwo() {  
    victim = i;  
    while (victim == i) {}  
}
```

- Not deadlock free


- Sequential execution deadlocks
 - if one thread never tries to get the lock
- Concurrent execution does not

Exercise: Prove mutual exclusion



Peterson's Algorithm


```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```



Peterson's Algorithm

Announce I'm
interested

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}  
public void unlock() {  
    flag[i] = false;  
}
```



Peterson's Algorithm

```
public void lock() {
```

```
    flag[i] = true;
```

```
    victim = i;
```

```
    while (flag[j] && victim == i) {};
```

```
}
```


```
public void unlock() {
```

```
    flag[i] = false;
```

```
}
```

**Announce I'm
interested**

Defer to other



Peterson's Algorithm

```
public void lock() {
```

```
    flag[i] = true;
```

```
    victim = i;
```

```
    while (flag[j] && victim == i) {};
```

```
}
```

```
public void unlock() {
```


```
    flag[i] = false;
```

```
}
```

**Announce I'm
interested**

Defer to other

**Wait while other
interested & I'm
the victim**



Peterson's Algorithm

```
public void lock() {
```

```
    flag[i] = true;
```

```
    victim = i;
```

```
    while (flag[j] && victim == i) {};
```

```
}
```

```
public void unlock() {
```

```
    flag[i] = false;
```


```
}
```

Announce I'm
interested

Defer to other

Wait while other
interested & I'm
the victim

No longer
interested




Mutual Exclusion

(1) $\text{write}_B(\text{Flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B)$

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {}  
}
```

From the Code



Also from the Code

**(2) write_A(victim=A) → read_A(flag[B])
→ read_A(victim)**

```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i) {};  
}
```

Assumption


(3) $\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$

W.L.O.G. assume A is the last thread to write victim



Combining Observations

- (1)** $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow \text{write}_B(\text{victim}=B)$
- (3)** $\text{write}_B(\text{victim}=B) \rightarrow \text{write}_A(\text{victim}=A)$
- (2)** $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$
 $\rightarrow \text{read}_A(\text{victim})$




Combining Observations

(1) write_B(flag[B]=true)→

(3) write_B(victim=B)→

**(2) write_A(victim=A)→read_A(flag[B])
→ read_A(victim)**



Combining Observations

(1) $\text{write}_B(\text{flag}[B]=\text{true}) \rightarrow$


```
public void lock() {  
    flag[i] = true;  
    victim = i;  
    while (flag[j] && victim == i){};  
}
```

(3) $\text{write}_B(\text{victim}=B)$ -

(2) $\text{write}_A(\text{victim}=A) \rightarrow \text{read}_A(\text{flag}[B])$

$\rightarrow \text{read}_A(\text{victim})$


A **read** $\text{flag}[B] == \text{true}$ **and** $\text{victim} == A$, so
it could not have entered the CS (QED)



Deadlock Free

```
public void lock() {  
    ...  
    while (flag[j] && victim == i) {};
```

- Thread blocked
 - only at **while** loop
 - only if other's flag is **true**
 - only if it is the **victim**
- **Solo: other's flag is false**
- **Both: one or the other not the victim**



Starvation Free


- Thread *i* blocked only if *j* repeatedly re-enters so that

`flag[j] == true` and
`victim == i`

- When *j* re-enters
 - it sets `victim` to *j*.
 - So *i* gets in

```
public void lock() {  
    flag[i] = true;  
    victim  = i;  
    while (flag[j] && victim == i) {};  
}  
  
public void unlock() {  
    flag[i] = false;  
}
```



Problem? Only two threads



The Filter Algorithm for n Threads

There are $n-1$ “waiting rooms” called levels

- At each level
 - At least one enters level
 - At least one blocked at that level if many try
- Only one thread makes it through
- A thread at level j is also at level $j-1, \dots, 0$




Filter

```
class Filter implements Lock {  
    int[] level; // level[i] for thread i  
    int[] victim; // victim[L] for level L
```

```
    public Filter(int n) {  
        level = new int[n];  
        victim = new int[n];  
        for (int i = 1; i < n; i++) {  
            level[i] = 0;  
        }  
        ...  
    }
```

Thread 2 at level 4



Filter

```
class Filter implements Lock {
```

```
...
```

```
public void lock(){
```

```
    for (int L = 1; L < n; L++) {
```

```
        level[i] = L;
```

```
        victim[L] = i;
```

thread i trying to enter level L

```
        while (( $\exists$  k != i level[k] >= L) &&  
                victim[L] == i ) {};
```


```
    }
```

```
public void unlock() {
```

```
    level[i] = 0;
```

```
}
```


thread i succeeded in
entering level L



Filter

```
class Filter implements Lock {  
    ...  
  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                victim[L] == i) {};  
        }  
    }  
    public void release(int i) {  
        level[i] = 0;  
    }  
}
```

One level at a time



Filter

```
class Filter implements Lock {
```

```
...
```

```
public void lock() {
```

```
    for (int L = 1; L < n; L++) {
```

```
        level[i] = L;
```

```
        victim[L] = i;
```

```
        while (( $\exists$  k != i) level[k] >= L) &&  
            victim[L] == i)
```


```
    }
```

```
public void release(int i)
```

```
    level[i] = 0;
```

```
}
```

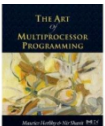
**Announce
intention to enter
level L**



Filter

```
class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists$  k != i) level[k] >= L) &&
                victim[L] == i) {};
        }
    }
    public void release(int i)
        level[i] = 0;
}
```

**Give priority to
anyone but me**



Filter


Wait as long as someone else is at same or higher level, and I'm designated victim

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists$  k != i) level[k] >= L) &&  
            victim[L] == i) {};  
    }  
}  
public void release(int i) {  
    level[i] = 0;  
}
```

Filter


```
class Filter implements Lock {  
    int level[n];  
    int victim[n];  
    public void lock() {  
        for (int L = 1; L < n; L++) {  
            level[i] = L;  
            victim[L] = i;  
            while (( $\exists$  k != i) level[k] >= L) &&  
                victim[L] == i) {};  
        }  
    }  
}
```

Thread *enters* level L when it completes the loop





Claim

- Start at level $L=0$
- At most $n-L$ threads enter level L
- Mutual exclusion at level $L=n-1$




Induction Hypothesis

- No more than $n-(L-1)$ at level $L-1$
- Induction step: by contradiction
 - Assume all at level $L-1$ enter level L
 - A last to write `victim[L]`
 - B is any other thread at level L



Proof Structure



Show that **A must have seen B in level[L] and since victim[L] == A could not have entered**




Just Like Peterson

(1) $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$

```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while (( $\exists$  k != i) level[k] >= L)  
            && victim[L] == i) {};  
    }  
}
```


From the Code



From the Code

**(2) write_A(victim[L]=A) → read_A(level[B])
→ read_A(victim[L])**

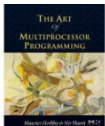
```
public void lock() {  
    for (int L = 1; L < n; L++) {  
        level[i] = L;  
        victim[L] = i;  
        while ((∃ k != i) level[k] >= L)  
            && victim[L] == i) {};  
    }  
}
```



By Assumption


(3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

By assumption, A is the last thread to write victim[L]



Combining Observations

- (1)** $\text{write}_B(\text{level}[B]=L) \rightarrow \text{write}_B(\text{victim}[L]=B)$
- (3)** $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$
- (2)** $\text{write}_A(\text{victim}[L]=A) \rightarrow \text{read}_A(\text{level}[B])$
 $\rightarrow \text{read}_A(\text{victim}[L])$




Combining Observations

(1) $\text{write}_B(\text{level}[B]=L) \rightarrow$

(3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

(2) $\text{read}_A(\text{level}[B])$

$\rightarrow \text{read}_A(\text{victim}[L])$




Combining Observations

(1) $\text{write}_B(\text{level}[B]=L) \rightarrow$

(3) $\text{write}_B(\text{victim}[L]=B) \rightarrow \text{write}_A(\text{victim}[L]=A)$

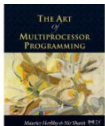
(2) $\dots \rightarrow \text{read}_A(\text{level}[B])$
 $\rightarrow \text{read}_A(\text{victim}[L])$

A read $\text{level}[B] \geq L$, and $\text{victim}[L] = A$, so it could not have entered level L!




No Starvation

- **Filter Lock satisfies properties:**
 - Just like Peterson Alg at any level
 - So no one starves
- **But what about fairness?**
 - Threads can be overtaken by others




Bounded Waiting

- Want stronger fairness guarantees
- Thread not “overtaken” too much
- If **A** starts before **B**, then **A** enters before **B**?
- But what does “start” mean?
- Need to adjust definitions



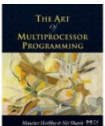
Bounded Waiting

- **Divide `lock()` method into 2 parts:**
 - **Doorway interval:**
 - Written D_A
 - always finishes in finite steps
 - **Waiting interval:**
 - Written W_A
 - may take unbounded steps




First-Come-First-Served

- **For threads A and B:**
 - **If $D_A^k \rightarrow D_B^j$**
 - **A's k-th doorway precedes B's j-th doorway**
 - **Then $CS_A^k \rightarrow CS_B^j$**
 - **A's k-th critical section precedes B's j-th critical section**
 - **B cannot overtake A**




Fairness Again

- **Filter Lock satisfies properties:**
 - No one starves
 - But very weak fairness
 - Can be overtaken arbitrary # of times
 - That's pretty lame...




Bakery Algorithm

- Provides First-Come-First-Served
- How?
 - Take a “number”
 - Wait until lower numbers have been served
- Lexicographic order
 - $(a,i) > (b,j)$
 - If $a > b$, or $a = b$ and $i > j$



Bakery Algorithm

```
class Bakery implements Lock {  
    boolean[] flag;  
    Label[] label;  
    public Bakery (int n) {  
        flag = new boolean[n];  
        label = new Label[n];  
        for (int i = 0; i < n; i++) {  
            flag[i] = false; label[i] = 0;  
        }  
    }  
    ...  
}
```



Bakery Algorithm

```
class Bakery implements Lock {
```

```
    boolean[] flag;
```

```
    Label[] label;
```

```
    public Bakery (int n) {
```

```
        flag = new boolean[n];
```

```
        label = new Label[n];
```



```
        for (int i = 0; i < n; i++) {
```

```
            flag[i] = false; label[i] = 0;
```

```
        }
```

```
    }
```


```
    ...
```



CS

Bakery Algorithm


```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i],i) > (label[k],k));  
    }  
}
```



Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) > (label[k], k));  
    }  
}
```


Doorway



Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) > (label[k], k));  
    }  
}
```


I'm interested



Bakery Algorithm

Take increasing
label (read labels
in some arbitrary
order)

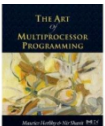
```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) > (label[k], k));  
    }  
}
```



Bakery Algorithm

```
class Bakery implements Lock {  
    ...  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) > (label[k], k));  
    }
```

Someone is
interested




Bakery Algorithm

```
class Bakery implements Lock {  
    boolean flag[n];  
    int label[n];  
  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0], ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
            && (label[i], i) > (label[k], k));  
    }
```


Someone is
interested ...

... whose (label,i) in
lexicographic order is lower



Bakery Algorithm

```
class Bakery implements Lock {  
  
    ...  
  
    public void unlock() {  
        flag[i] = false;  
    }  
}
```



Bakery Algorithm

```
class Bakery implements Lock {
```

```
    ...
```

```
    public void unlock() {
```


```
        flag[i] = false;
```

```
    }
```

```
}
```


**No longer
interested**

labels are always increasing



No Deadlock

- There is always one thread with earliest label
- Ties are impossible (why?)



First-Come-First-Served

- If $D_A \rightarrow D_B$ then
 - A's label is smaller
- And:
 - $\text{write}_A(\text{label}[A]) \rightarrow$
 - $\text{read}_B(\text{label}[A]) \rightarrow$
 - $\text{write}_B(\text{label}[B]) \rightarrow \text{read}_B(\text{flag}[A])$
- So B sees
 - smaller label for A
 - locked out while $\text{flag}[A]$ is true


```
class Bakery implements Lock {  
  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
                        ..., label[n-1])+1;  
  
        while (∃k flag[k]  
                && (label[i], i) >  
                (label[k], k));  
    }  
}
```



Mutual Exclusion

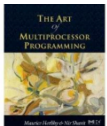
- Suppose A and B in CS together
- Suppose A has earlier label
- When B entered, it must have seen
 - flag[A] is false, or
 - label[A] > label[B]

```
class Bakery implements Lock {  
    public void lock() {  
        flag[i] = true;  
        label[i] = max(label[0],  
                        ..., label[n-1])+1;  
        while ( $\exists k$  flag[k]  
                && (label[i], i) >  
                (label[k], k));  
    }  
}
```




Mutual Exclusion

- Labels are strictly increasing so
- B must have seen `flag[A] == false`




Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow$
 $\text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$



Mutual Exclusion

- Labels are strictly increasing so
- B must have seen $\text{flag}[A] == \text{false}$
- $\text{Labeling}_B \rightarrow \text{read}_B(\text{flag}[A]) \rightarrow \text{write}_A(\text{flag}[A]) \rightarrow \text{Labeling}_A$
- Which contradicts the assumption that A has an earlier label



This work is licensed under a Creative Commons Attribution-ShareAlike 2.5 License.

- **You are free:**
 - to Share — to copy, distribute and transmit the work
 - to Remix — to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- **For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to**
 - <http://creativecommons.org/licenses/by-sa/3.0/>.
- **Any of the above conditions can be waived if you get permission from the copyright holder.**
- **Nothing in this license impairs or restricts the author's moral rights.**

