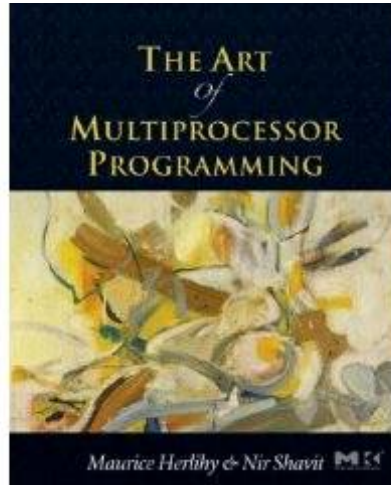


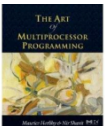
Mutual Exclusion



Companion slides for
The Art of Multiprocessor
Programming
by Maurice Herlihy & Nir Shavit
With some very minor changes by APS

Deep Philosophical Question

- The Bakery Algorithm is
 - Succinct, Elegant,
 - Mutual Exclusion, Deadlock-Free, Starvation-Free
 - Fair.
- Q: So why isn't it practical?
- A: Well, you have to read **N** distinct variables



Shared Memory

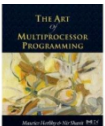
- Shared read/write memory locations called **Registers** (historical reasons)
 - Assume Reads and Writes are atomic
- Come in different flavors
 - Multi-Reader-Single-Writer (**Flag[]**)
 - Multi-Reader-Multi-Writer (**Victim[]**)
 - Not that interesting: SRMW and SRSW



Theorem

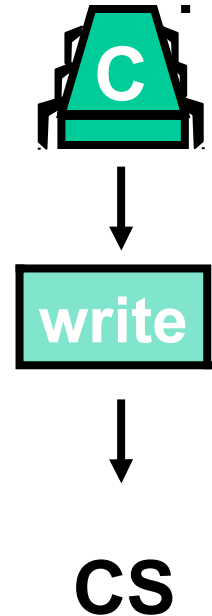
At least **N** MRSW (multi-reader/single-writer) registers are needed to solve **N**-thread deadlock-free mutual exclusion.

N registers like **Flag[]**...



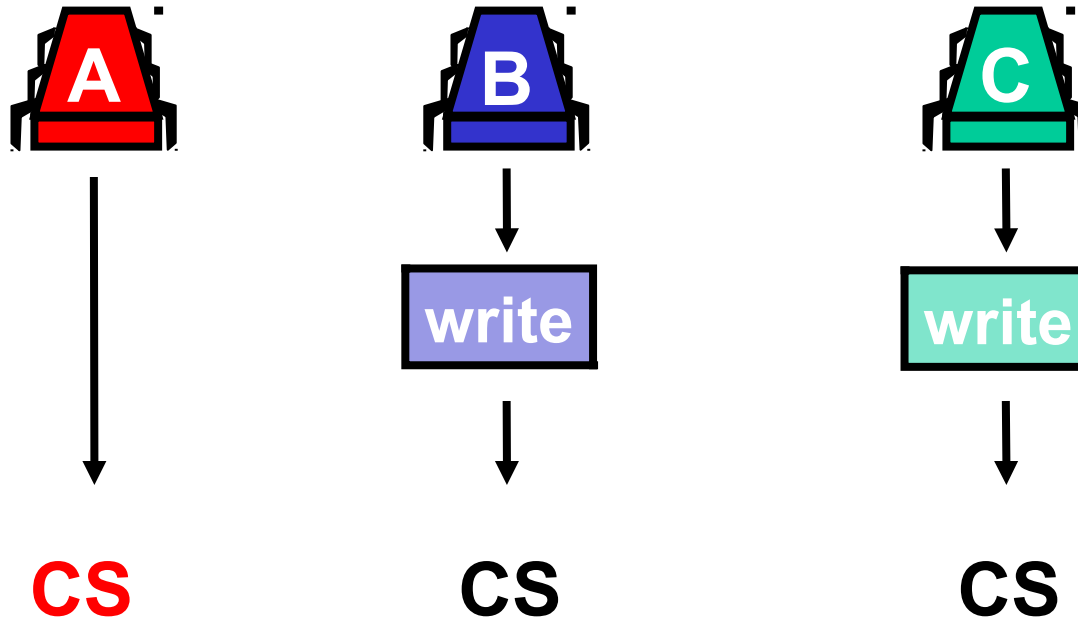
Proving Algorithmic Impossibility

- To show no algorithm exists:
 - assume by way of contradiction one does,
 - show a **bad execution** that violates properties:
 - in our case assume an alg for deadlock free mutual exclusion using $< N$ registers

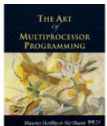


Proof: Need N-MRSW Registers

Each thread must write to some register

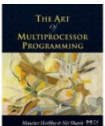


...can't tell whether **A** is in critical section



Upper Bound

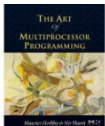
- Bakery algorithm
 - Uses **2N MRSW** registers
- So the bound is (pretty) tight
- But what if we use **MRMW** registers?
 - Like **victim[]** ?



Bad News Theorem

At least **N** MRMW multi-reader/**multi-writer** registers are needed to solve deadlock-free mutual exclusion.

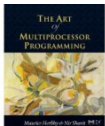
(So multiple writers don't help)



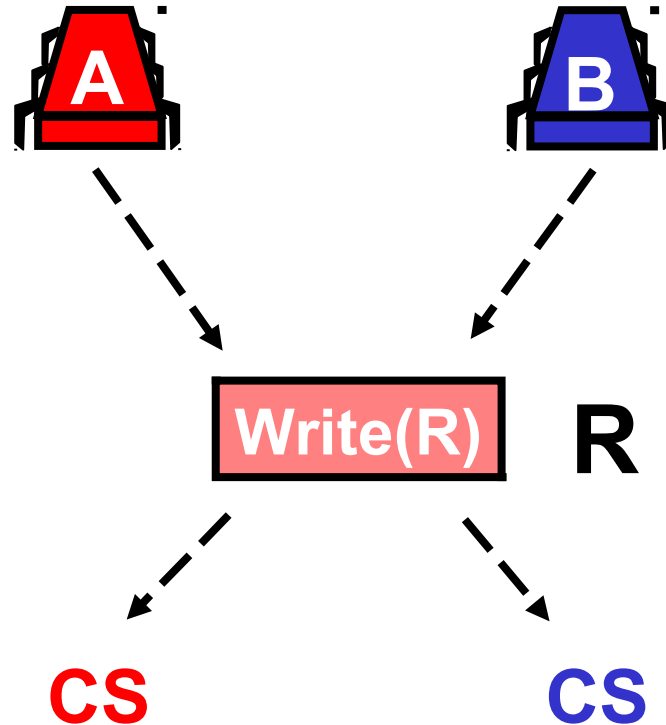
Theorem (For 2 Threads)

Theorem: Deadlock-free mutual exclusion for 2 threads requires at least 2 multi-reader multi-writer registers

Proof: assume one register suffices and derive a contradiction




Two Thread Execution




- Threads run, reading and writing R
- Deadlock free so at least one gets in

Covering State

- One thread each about to write to each shared location on next instruction
- The shared locations “look” like there is no thread currently in, or waiting to enter, the Critical Section

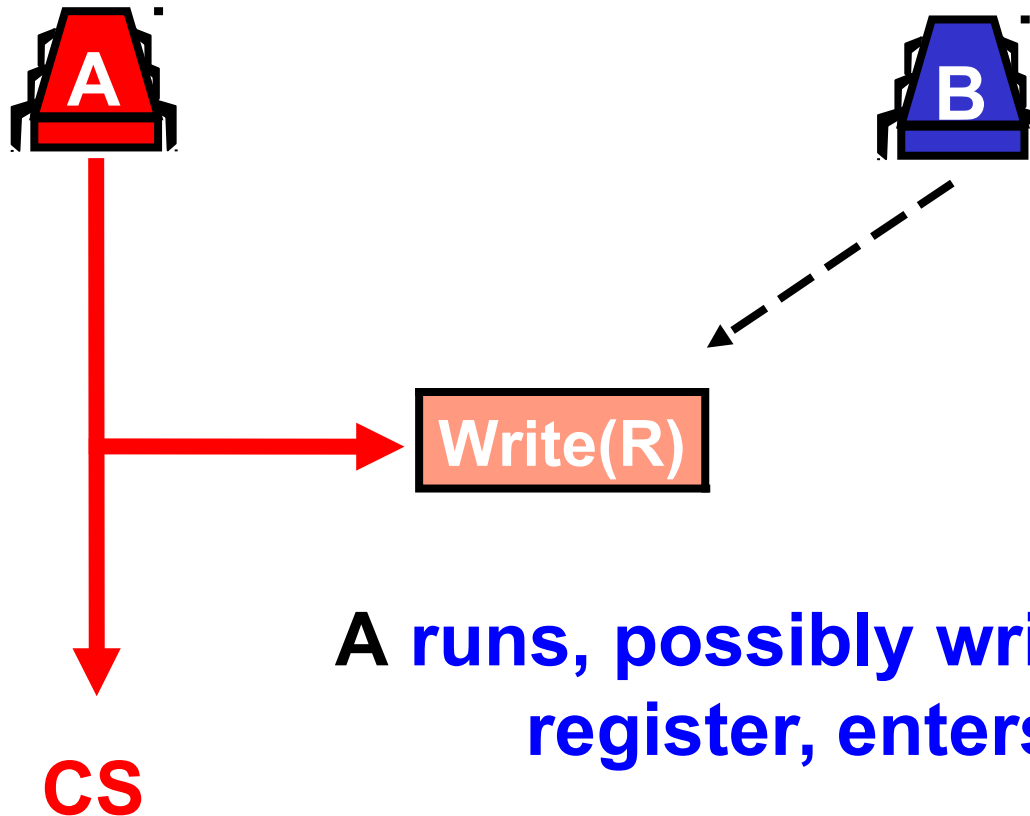


Covering State for One Register Always Exists

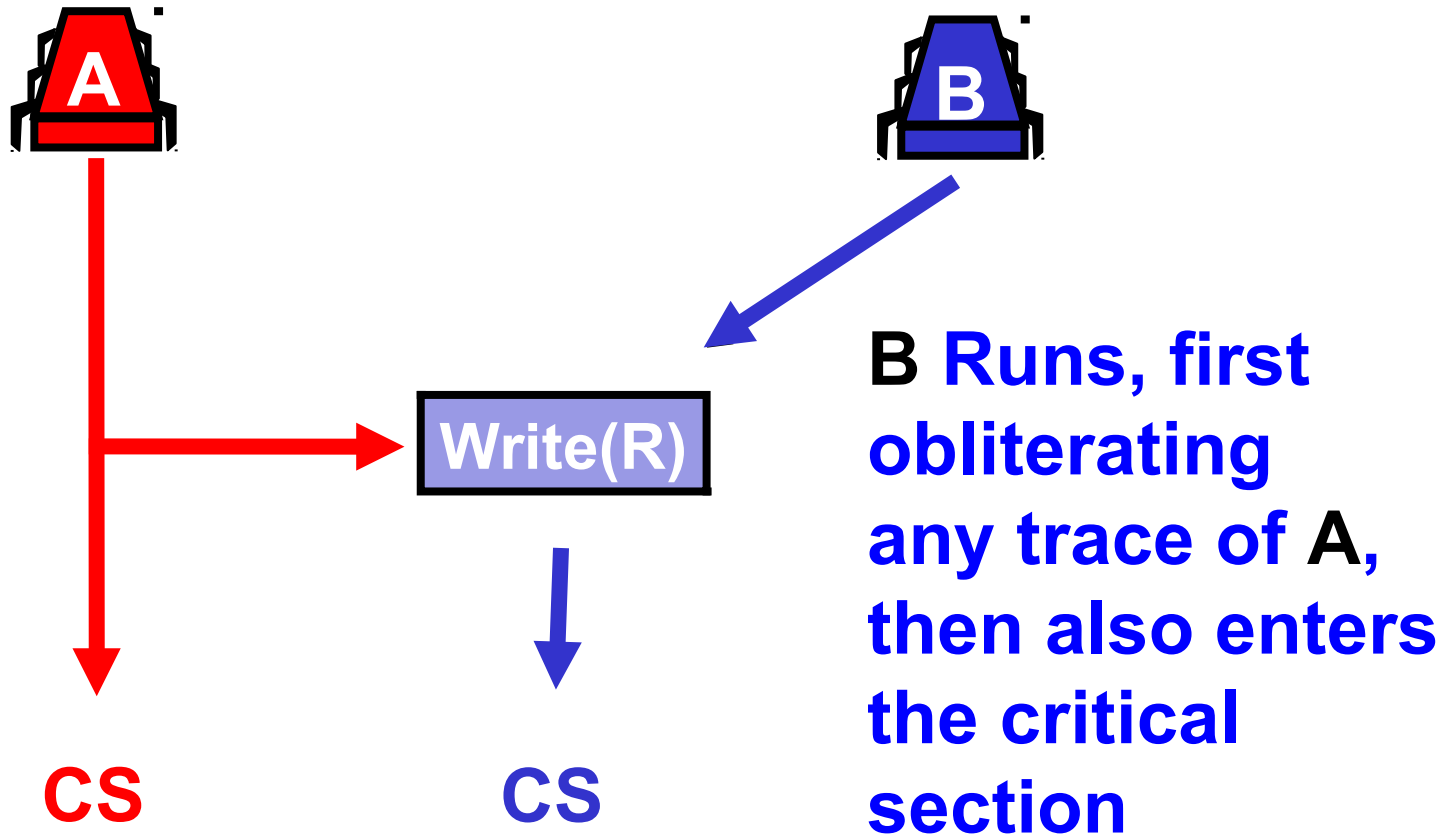


In any protocol **B has to write to the register before entering CS, so stop it just before**

Proof: Assume Cover of 1

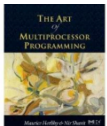


Proof: Assume Cover of 1

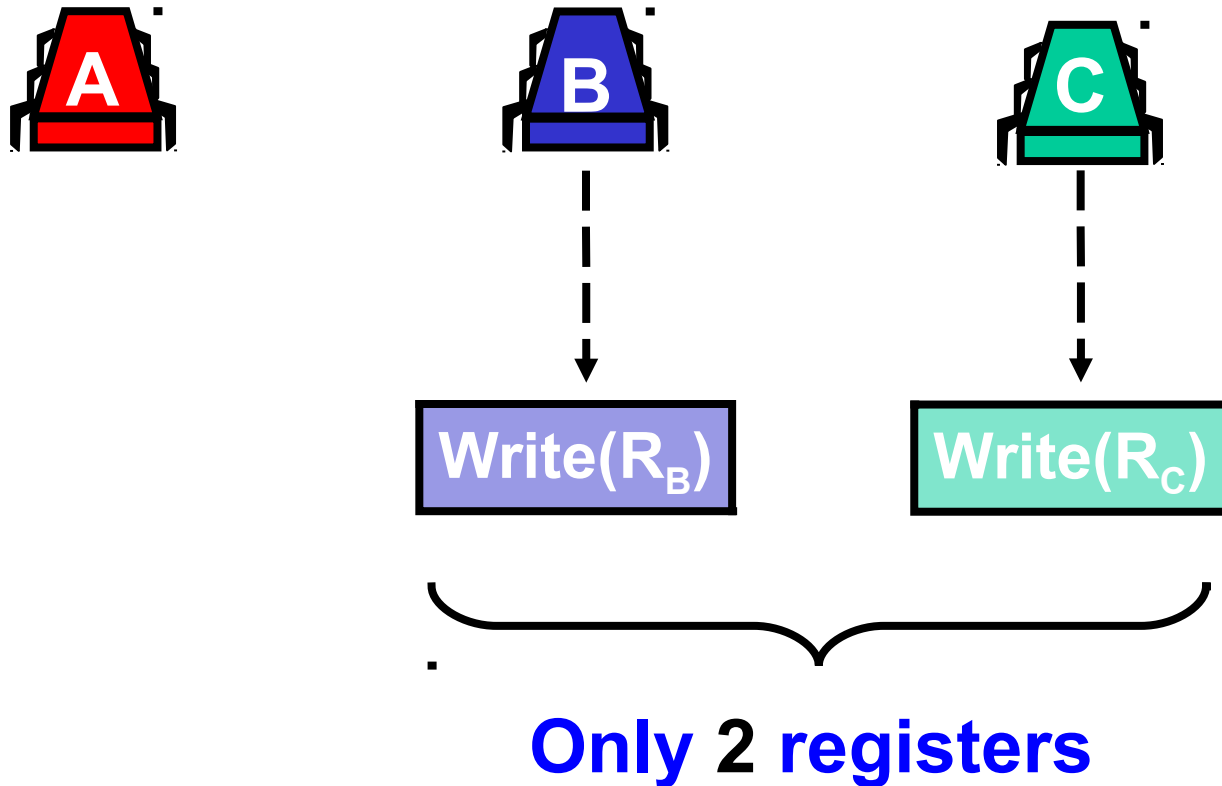


Theorem

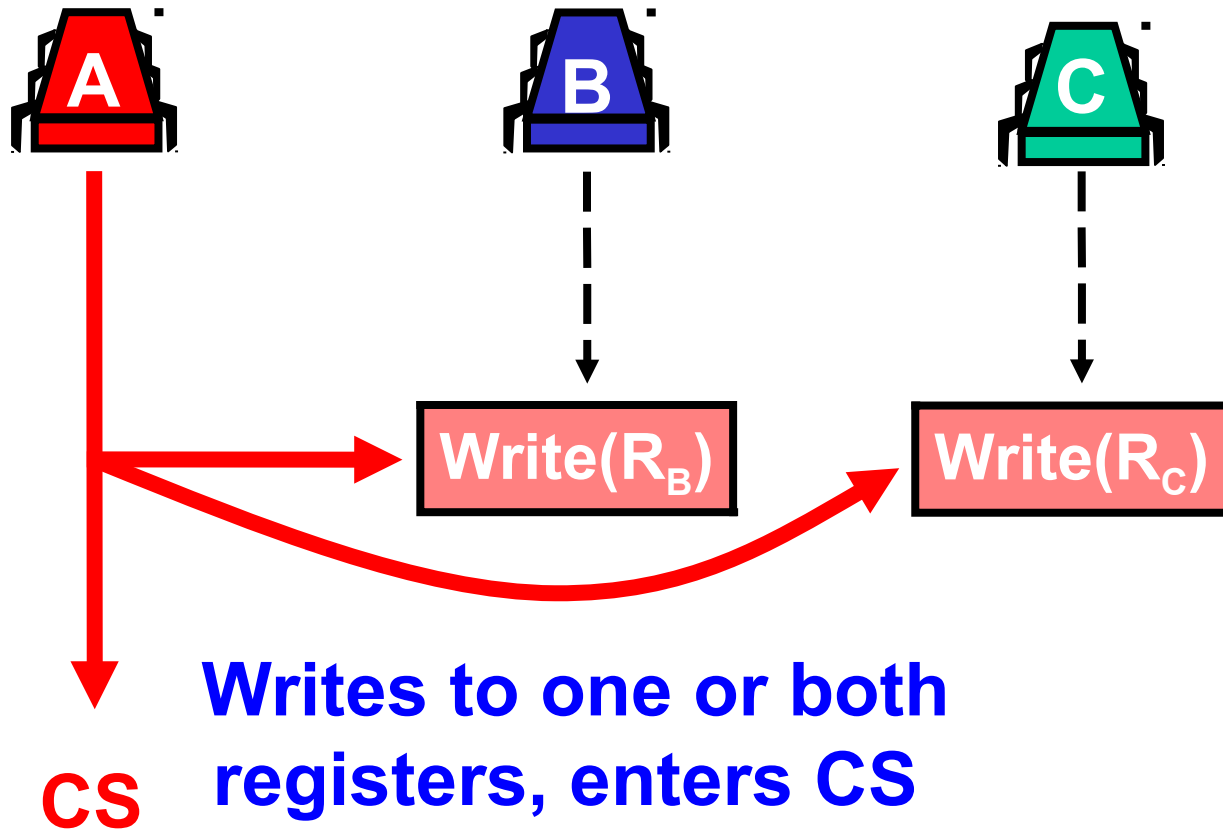
Deadlock-free mutual exclusion for 3 threads requires at least 3 multi-reader multi-writer registers



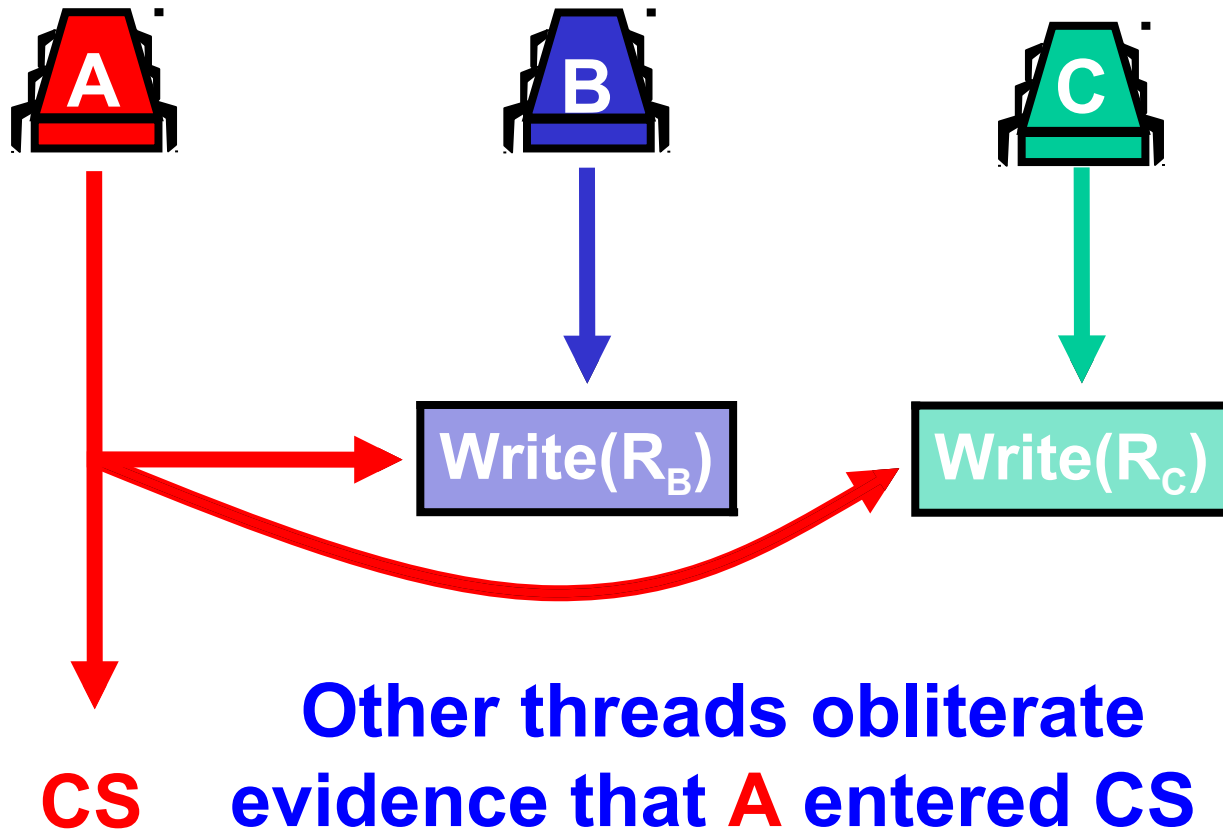
Proof: Assume Cover of 2



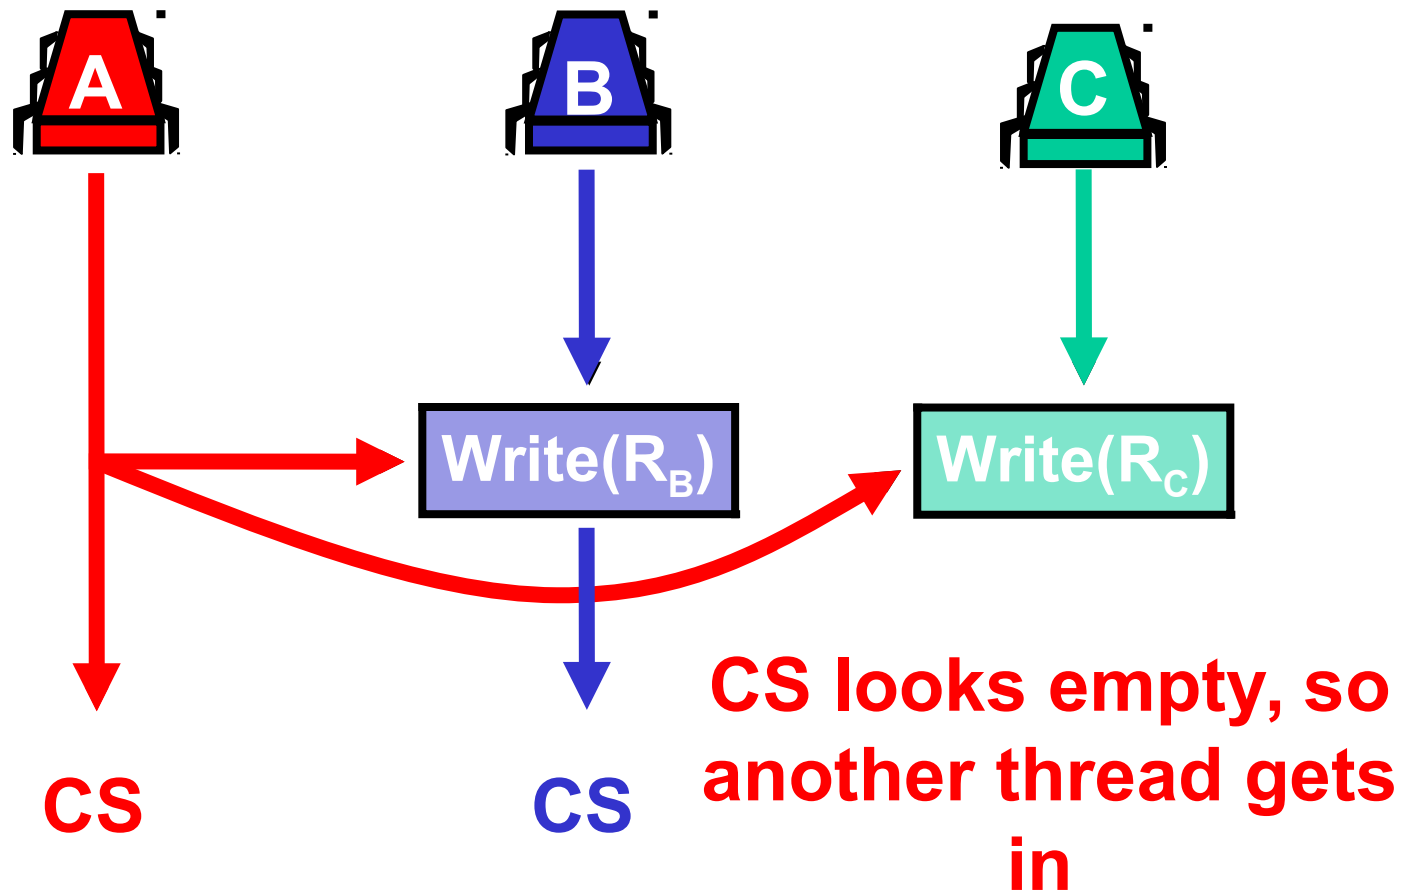
Run A Solo



Obliterate Traces of A

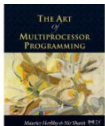


Mutual Exclusion Fails

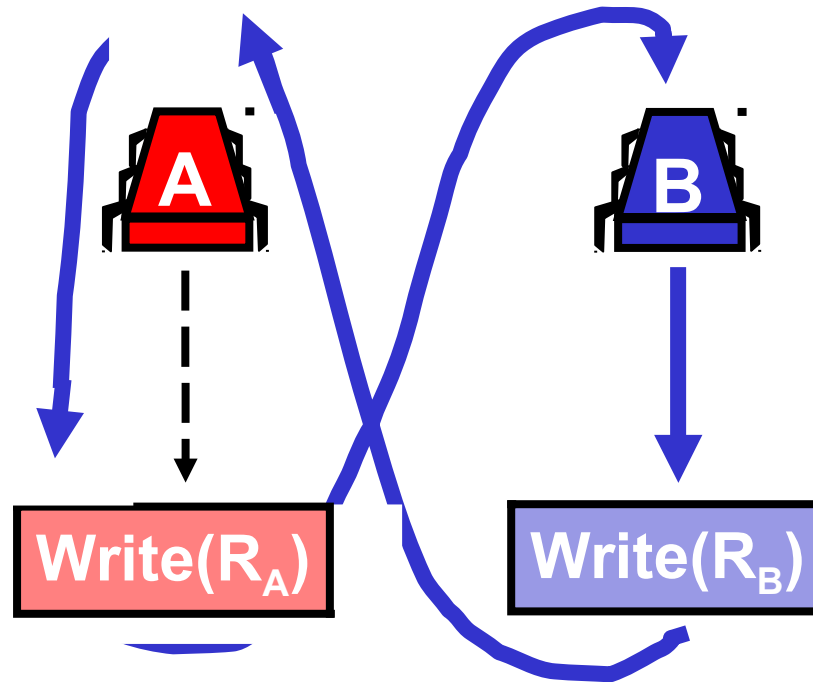


Proof Strategy

- **Proved: a contradiction starting from a covering state for 2 registers**
- **Claim: a covering state for 2 registers is reachable from any state where CS is empty**



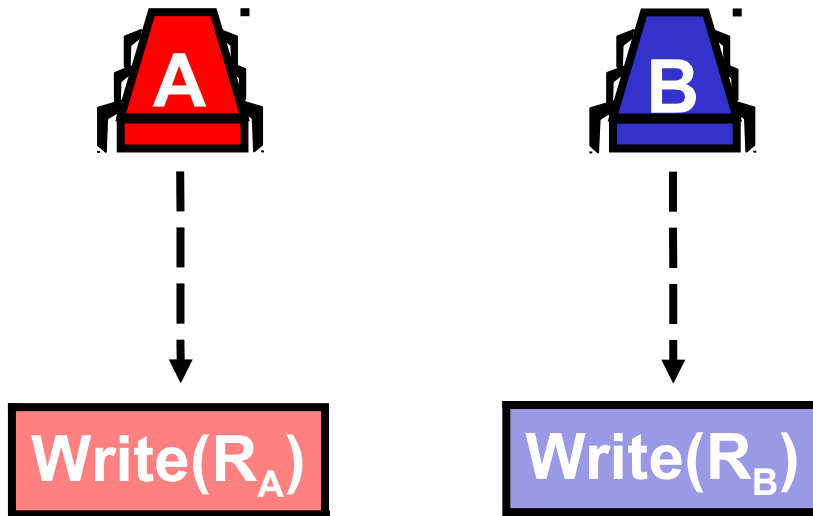
Covering State for Two




- If we run **B** through CS **3** times, **B** must return **twice** to cover some register, say R_B



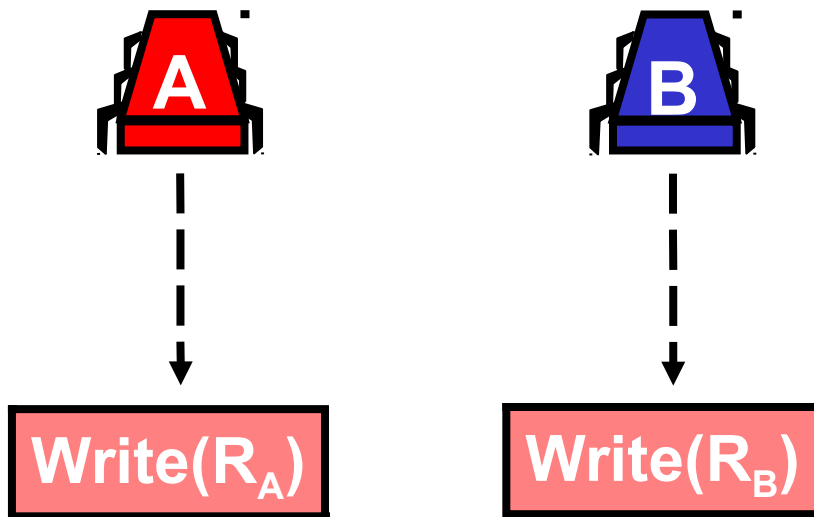
Covering State for Two



- Start with **B** covering register R_B for the 1st time
- Run **A** until it is about to write to uncovered R_A
- Are we done?

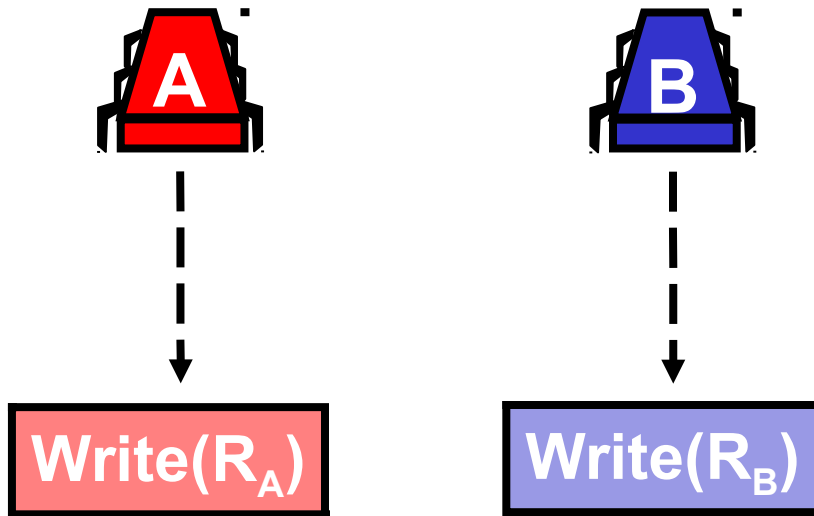


Covering State for Two

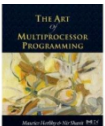


- **NO!** A could have written to R_B
- So CS no longer looks empty to thread C


Covering State for Two



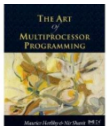
- Run **B** obliterating traces of **A** in R_B
- Run **B** again until it is about to write to R_B
- Now we are done



Inductively We Can Show

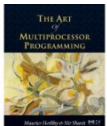


- **There is a covering state**
 - **Where k threads not in CS cover k distinct registers**
 - **Proof follows when $k = N-1$**




Summary of Lecture

- In the 1960's several **incorrect** solutions to starvation-free mutual exclusion using RW-registers were published...
- Today we know how to solve FIFO N thread mutual exclusion using $2N$ RW-Registers



Summary of Lecture

- **N RW-Registers inefficient**
 - Because writes “**cover**” older writes
- **Need stronger hardware operations**
 - that do not have the “**covering problem**”



This work is licensed under a Creative Commons Attribution-ShareAlike 2.5 License.

- **You are free:**
 - to Share — to copy, distribute and transmit the work
 - to Remix — to adapt the work
- **Under the following conditions:**
 - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
 - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- **For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to**
 - <http://creativecommons.org/licenses/by-sa/3.0/>.
- **Any of the above conditions can be waived if you get permission from the copyright holder.**
- **Nothing in this license impairs or restricts the author's moral rights.**

