Exercise 1: Bernoulli Distribution

Assume that N data points are independent, which means that the joint conditional density can be factorised into N separate terms, one for each data point:

$$\mathcal{L} = p\left(D \mid q\right) = \prod_{i=1}^{N} p\left(x_i \mid q\right) = \prod_{i=1}^{N} q^{x_i}\ (1-q)^{1-x_i}.$$

This equation tells us how likely our dataset D is, given the current model parametrised by success probability q. Since the dataset is fixed, changing the model will result in different likelihood values. Maximum Likelihood method tries to find the model that maximises the likelihood $\mathcal{L}$. Normally, we will maximise the *natural logarithm* (i.e. logarithm with base e) of the likelihood because the estimated argument $\hat{q}$ that maximises the log-likelihood will also maximise the likelihood.

$$\begin{aligned}
\log \mathcal{L} &= \sum_{i=1}^{N} \log q^{x_i}\ (1-q)^{1-x_i} = \sum_{i=1}^{N} \left(\log q^{x_i} + \log(1-q)^{1-x_i}\right) \\
&= \sum_{i=1}^{N} \left(x_i\ \log q + (1-x_i)\ \log(1-q)\right) \\
&= \sum_{i=1}^{N} (x_i)\ \log q + \sum_{i=1}^{N} (1-x_i)\ \log(1-q).
\end{aligned}$$

We now can find the optimal parameter by taking derivative, equating them to zero and solving for turning points. For q,

$$\begin{aligned}
\frac{\partial \log \mathcal{L}}{\partial q} &= \sum_{i=1}^{N} (x_i)\ \frac{\partial \log q}{\partial q} + \sum_{i=1}^{N} (1-x_i)\ \frac{\partial \log(1-q)}{\partial q} \\
&= \frac{1}{q}\ \sum_{i=1}^{N} (x_i) - \frac{1}{1-q}\ \sum_{i=1}^{N} (1-x_i)
\end{aligned}$$

Equating $\partial \log \mathcal{L}/\partial q$ to zero yields

$$\frac{1}{\hat{q}}\ \sum_{i=1}^{N} (x_i) - \frac{1}{1-\hat{q}}\ \sum_{i=1}^{N} (1-x_i) = 0,$$

equivalently

$$(1-\hat{q}) \sum_{i=1}^{N} (x_i) = \hat{q} \left(N - \sum_{i=1}^{N} (x_i)\right).$$

Simplifying it yields the maximum likelihood estimate of the Bernoulli distribution

$$\hat{q}_{\mathrm{ML}} = \frac{1}{N} \sum_{i=1}^{N} x_i$$

Exercise 2: Univariate Gaussian Distribution

Assume that a dataset D of N values (i.e. $D = \{x_1, x_2, \ldots, x_N\}$), was sampled from a Gaussian distribution $\mathcal{N}(\mu, \sigma^2)$. Assuming that the data points are independently and identically distributed. Let's get started by writing down the joint density function:

$$\mathcal{L} = p(D \mid \mu, \sigma^2) = \prod_{i=1}^{N} \mathcal{N}(x_i; \mu, \sigma^2) = \prod_{i=1}^{N} \frac{1}{\sigma\sqrt{2\pi}} \exp\left\{ -\frac{(x_i - \mu)^2}{2\sigma^2} \right\}$$

Follow the same approach as in Exercise 1 by taking the *natural logarithm* of $\mathcal{L}$ as

$$\begin{aligned} \log \mathcal{L} &= \sum_{i=1}^{N} \log \frac{1}{\sigma\sqrt{2\pi}} \exp\left\{ -\frac{(x_i - \mu)^2}{2\sigma^2} \right\} \\ &= \sum_{i=1}^{N} \left(-\log \sigma - \frac{1}{2}\log(2\pi) - \frac{(x_i - \mu)^2}{2\sigma^2} \right) \end{aligned}$$

Unlike the Bernoulli distribution which has only one parameter q, the Gaussian distribution is characterised by two parameters: mean (μ) and variance (σ). Therefore, we have to calculate two partial derivatives with respect to μ and σ, i.e. $\partial \log \mathcal{L}/\partial\mu$ and $\partial \log \mathcal{L}/\partial\sigma$. By equating both derivatives to zero, we can find the maximum likelihood estimates $\hat{\mu}_{\mathrm{ML}}$ and $\hat{\sigma}^2_{\mathrm{ML}}$ for the Gaussian model. Calculating the two derivatives as follows:

$$\frac{\partial \log \mathcal{L}}{\partial \sigma} = \sum_{i=1}^{N} \left(-\frac{1}{\sigma} + \frac{(x_i - \mu)^2}{\sigma^3} \right) \quad \text{and} \quad \frac{\partial \log \mathcal{L}}{\partial \mu} = \sum_{i=1}^{N} \frac{x_i - \mu}{\sigma^2}$$

By equating both derivatives to zero, we obtain

$$\begin{aligned} \frac{\partial \log \mathcal{L}}{\partial \mu} = 0 \Rightarrow \sum_{i=1}^{N} (x_i - \mu) &= 0 \\ \Rightarrow \mu &= \frac{1}{N} \sum_{i=1}^{N} x_i, \end{aligned}$$

and

$$\begin{aligned} \frac{\partial \log \mathcal{L}}{\partial \sigma} = 0 \Rightarrow N - \frac{1}{\sigma^2} \sum_{i=1}^{N} (x_i - \mu)^2 &= 0 \\ \Rightarrow \sigma^2 &= \frac{1}{N} \sum_{i=1}^{N} (x_i - \mu)^2 \end{aligned}$$

Hence, the ML estimate of the mean and variance of the Gaussian distribution is

$$\hat{\mu}_{\mathrm{ML}} = \frac{1}{N} \sum_{i=1}^{N} x_i \quad \text{and} \quad \hat{\sigma}^2_{\mathrm{ML}} = \frac{1}{N} \sum_{i=1}^{N} (x_i - \hat{\mu}_{\mathrm{ML}})^2$$

Exercise 3a

Since the sensors are independent, the likelihood is

$$\mathcal{L}(x) = p(z_1, z_2 \mid x) = p(z_1|x)p(z_2|x)$$

and since the sensors are gaussian

$$\mathcal{L}(x) \propto e^{-\frac{(z_1-x)^2}{2\sigma^2}} \times e^{-\frac{(z_2-x)^2}{2\sigma^2}} = e^{-\frac{(z_1-x)^2+(z_2-x)^2}{2\sigma^2}}$$

Here we ignored the irrelevant normalisation constants. Now the log-likelihood is given by

$$\log \mathcal{L}(x) = \frac{(z_1-x)^2 + (z_2-x)^2}{2\sigma^2} = \frac{(x-\bar{x})^2}{\sigma^2} + c(z_1, z_2),$$

where $\bar{x} = \frac{z_1+z_2}{2}$, and $c(z_1, z_2)$ is a constant independent of x. The maximum likelihood estimate of x is defined as

$$\hat{x} = \arg\max_x \mathcal{L}(x) = \arg\min_x \left(-\log \mathcal{L}(x)\right)$$

Now let's compute the min by differenating $-\log \mathcal{L}(x)$ with repect to x

$$\frac{\partial\{-\log \mathcal{L}(x)\}}{\partial x} = \frac{2(x-\bar{x})}{\sigma^2} = 0$$

Therefore, $\hat{x}_{\mathrm{ML}} = \bar{x} = (z_1 + z_2)/2$.

Exercise 3b

The sensors are independent

$$\mathcal{L}(x) = p(z_1, z_2|x) = p(z_1|x)p(z_2|x) \propto e^{-\frac{(z_1-x)^2}{2\sigma_1^2}} \times e^{-\frac{(z_2-x)^2}{2\sigma_2^2}}$$

The negative log-likelihood is then

$$\begin{aligned}
-\log \mathcal{L}(x) &= \frac{1}{2}\left[\frac{(z_1-x)^2}{\sigma_1^2} + \frac{(z_2-x)^2}{\sigma_2^2}\right] + \text{const} \\
&= \frac{1}{2}\left[\left(\frac{1}{\sigma_1^2} + \frac{1}{\sigma_2^2}\right)x^2 - 2\left(\frac{z_1}{\sigma_1^2} + \frac{z_2}{\sigma_2^2}\right)x\right] + \text{const} \\
&= \frac{1}{2}\left(\frac{1}{\sigma_1^2} + \frac{1}{\sigma_2^2}\right)\left[x - \frac{\sigma_1^{-2}z_1 + \sigma_2^{-2}z_2}{\sigma_1^{-2} + \sigma_2^{-2}}\right]^2 + \text{const}
\end{aligned}$$

which is maximised with respect to x when

$$\hat{x}_{\mathrm{ML}} = \frac{\sigma_1^{-2}z_1 + \sigma_2^{-2}z_2}{\sigma_1^{-2} + \sigma_2^{-2}}.$$

For example, if the sensors are $p(z_1|x) \sim \mathcal{N}(x, 10^2)$ and $p(z_2|x) \sim \mathcal{N}(x, 20^2)$. Suppose we obtain sensor readings of $z_1 = 130$ and $z_2 = 170$, then

$$\hat{x}_{\mathrm{ML}} = \frac{130/10^2 + 170/20^2}{1/10^2 + 1/20^2} = 138.0$$

It shows that the ML estimate is closer to the more confident measurement (i.e. smaller variance).