

06-20416 and 06-12412 (Intro to) Neural Computation

07 – Optimisation Algorithms

Per Kristian Lehre

Last lecture

- A **softmax** output layer allows output nodes to be interpreted as probabilities
- The probabilities indicate the likelihood of a class, given the input and the network
- A naive implementation of the softmax function can be numerically unstable

Outline

- Learning rate in stochastic gradient descent
- Alternatives to standard SGD
 - SGD with momentum
 - SGD with Nesterov momentum
 - AdaGrad
 - Adam


Repetition: Gradient Descent

Input: cost function $J: \mathbb{R}^n \rightarrow \mathbb{R}$
learning rate $\epsilon \in \mathbb{R}, \epsilon > 0$

$x \leftarrow$ some initial point in \mathbb{R}^n
while termination condition not met {

$$x \leftarrow x - \epsilon \cdot \nabla J(x)$$


}



Impact of learning rate on SGD

- The learning rate in SGD often strongly impacts the optimisation time
- Often necessary to adjust the learning rate according to the specific setting.


“Cartoon Picture”




Source: Karpathy, CS231n

Typical Behavior of Standard SGD

Case 1: Too high learning rate ϵ



Case 2: Too low learning rate ϵ



Gradient Descent with Momentum

choose an initial parameter Θ


$$v = 0$$

while termination condition not satisfied {

$$v = \alpha v - \epsilon \nabla_{\Theta} C(\Theta)$$

$$\Theta = \Theta + v$$

}



Physical interpretation:

A ball with position Θ and velocity v influenced by two forces, one which pushes the ball opposite of the current gradient, and a viscous drag determined by parameter $\alpha < 1$.

The momentum "smooths out" update steps.

The size of updates depends on how aligned the previous gradients are.

Two hyperparameters

- ϵ learning rate

- α factor which determines the influence of past gradients on the current update of the parameter
(often 0.5, 0.9, or 0.99)

Nesterov Momentum

choose an initial parameter Θ

$$v = 0$$

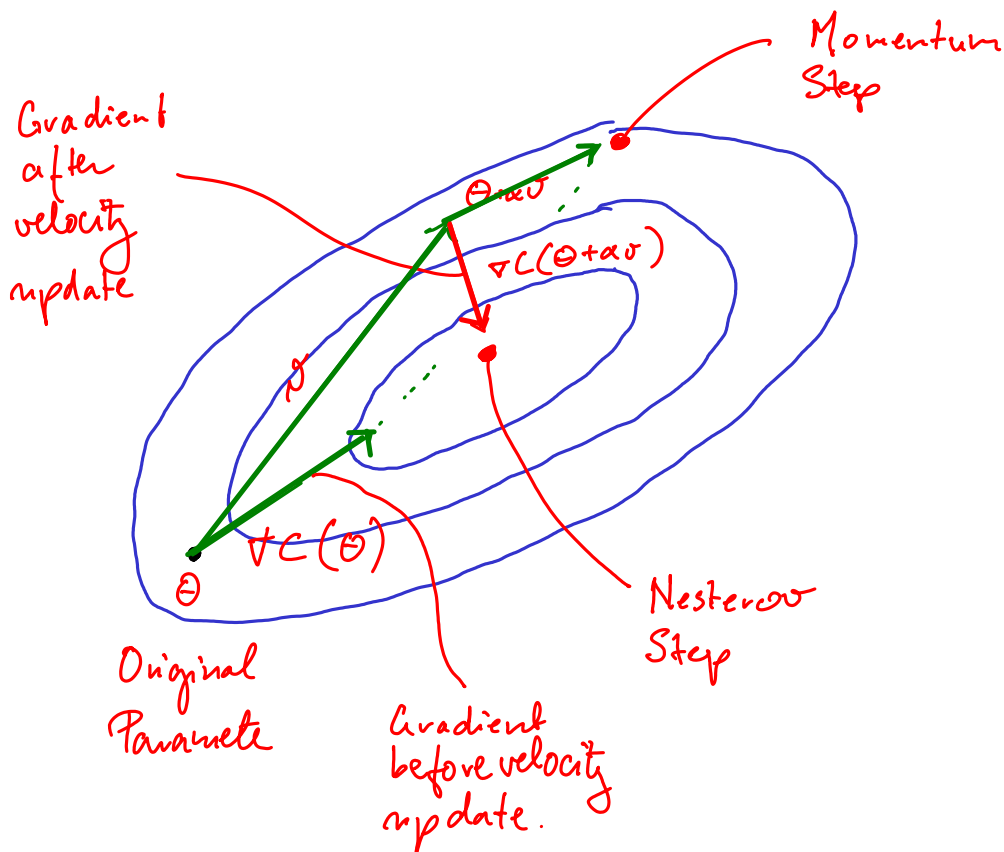
while termination condition not satisfied {

$$v = \alpha v - \epsilon \nabla_{\Theta} C(\Theta + \alpha v)$$

$$\Theta = \Theta + v$$

}

Nesterov momentum is a variant of standard momentum where the gradient is computed after the velocity is applied.



Adagrad (Duchi et al, 2011)

choose an initial parameter Θ

$$r = 0$$

while termination condition not satisfied {

$$g = \nabla_{\Theta} C(\Theta)$$

$$r = r + g \odot g$$

$$\eta = -\frac{\epsilon}{\delta + \sqrt{r}} \odot g$$

(division and square root applied componentwise)

$$\Theta = \Theta + \eta$$

}

squared
gradient

δ is a hyperparameter, typically $\delta = 10^{-6}$.

Adagrad adapts a (possibly different) "learning rate" $\frac{\epsilon}{\delta + \sqrt{r}}$ for each dimension according to accumulated square gradient r

— large r implies small $\frac{\epsilon}{\delta + \sqrt{r}}$

— small r implies large $\frac{\epsilon}{\delta + \sqrt{r}}$

Adagrad works well for problems with sparse gradients.

All gradients (new and old) weighted equally by r .

Adam (Kingma and Ba, 2014)

choose an initial parameter Θ

$$r = 0, s = 0, t = 0$$

while termination condition not satisfied {

$$t = t + 1$$

$$g = \nabla_{\Theta} C(\Theta)$$

$$s = \rho_1 s + (1 - \rho_1) g$$

$$r = \rho_2 r + (1 - \rho_2) g \odot g$$

$$\hat{s} = \frac{s}{1 - \rho_1^t}, \quad \hat{r} = \frac{r}{1 - \rho_2^t}$$

$$N = -\epsilon \cdot \frac{\hat{s}}{\sqrt{\hat{r}} + \delta}$$

$$\Theta = \Theta + N$$

}

Hyperparameters typically chosen as

$$\epsilon = 0.001$$

$$\rho_1 = 0.9$$

$$\rho_2 = 0.999$$

$$\delta = 10^{-8}$$

avg. gradient

 $\sqrt{\text{avg. squared gradient}}$

Widely used method in deep learning.

Summary

- Learning rate has strong impact on SGD
- Alternatives to SGD
 - SGD with momentum
 - SGD with Nesterov momentum
 - AdaGrad
 - Adam
- Open research problem how to choose appropriate algorithms