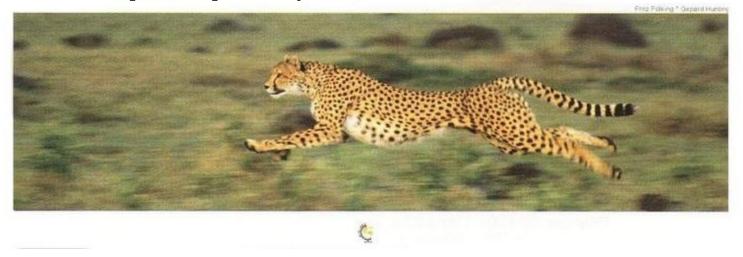
Co-Evolution

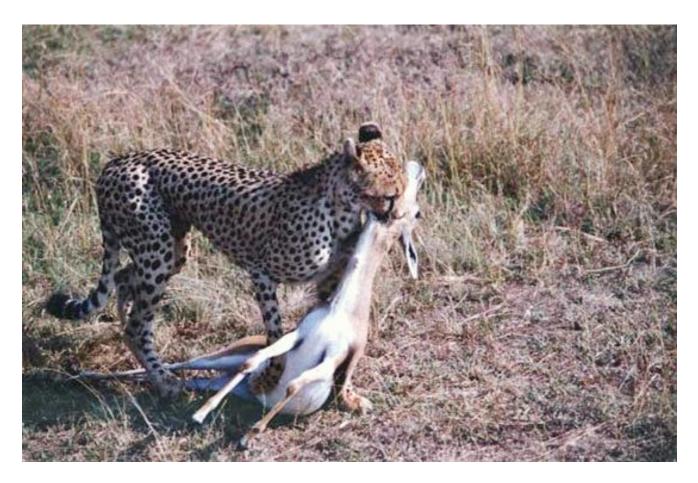
Ata Kaban
University of Birmingham

Cheetah: 60-70 mph, for up to 100 yards





Thompson's Gazelle: 50 mph...



The outcome (but about 50% of attempts are failures)

What is Co-Evolution?

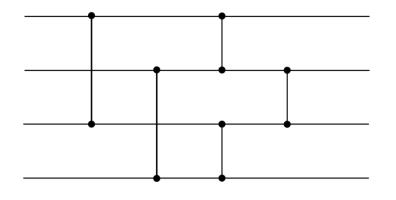
- Fitness of an individual depends on other individuals
 - Fitness landscape changes
 - Fitness of an individual may be different in different runs
- Change in one individual will change the fitness landscape in others
- Remember the Iterated Prisoner's Dilemma experiments? – the ones in which evolving strategies played against each other

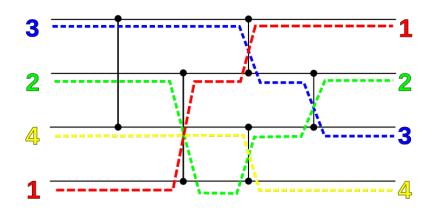
Types of co-evolution

- By evaluation
 - Competitive
 - Cooperative
- By population-organisation
 - Inter-population
 - Intra-population

Example 1: Sorting algorithm

- Goal: place the elements in a data structure (e.g. list or tree) in some specified order (e.g. numerical or alphanumeric)
- One approach in Knuth's book is the so called sorting network (for fixed nos of elements)
 - Horizontal lines=elements in the list
 - Vertical arrows=comparisons to be made (in parallel)
 - If compared elements are in wrong order than swap
- Will look at the simple case of n=16 elements (e0 —e15)





Designing sorting algorithms

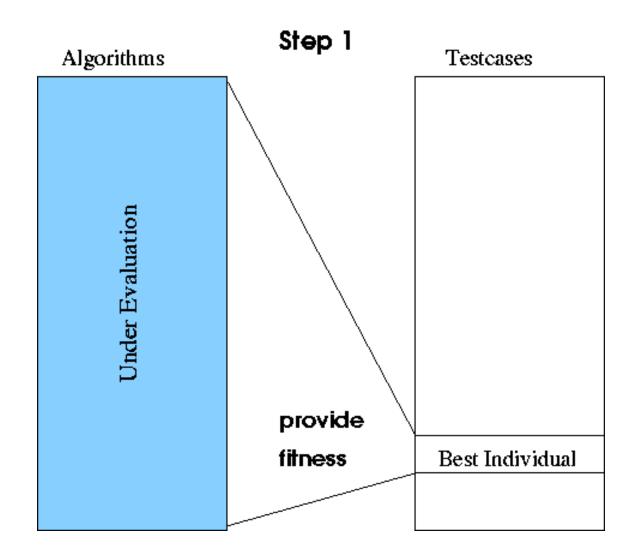
- Goal: make them
 - correct
 - efficient (reduce no of comparisons)
- What is the minimum number of comparisons necessary for correct sorting?
 - hot Q around the '60s (n=16 case)
 - '62: Bose & Nelson developed a net that needs 65 comparisons
 - '64: Batcher, Flyod, Knuth 63 comparisons
 - '69: Shapiro 62 comparisons
 - '69: Green 60 comparisons
 - '80s: W.D. Hillis can GA find an answer to this problem?


- Encoding the sorting net
 - ordered list of pairs to be compared = phenotypes
 - considered nets of 60—120 comparisons
 - genotype: diploid chromosomes
 - 1 individual = 15 pairs of 32 bit chromosomes, each encoding 4 comparisons
 - See more details on M. Mitchell, pp. 21—27.
- Fitness measure: the percentage of cases sorted correctly
 - Problem: how to compute this?
 - Test all possible inputs slow
 - Test fixed set of inputs? which?

First trials

- To foster speciation, individuals placed on a 2D grid (i.e. spatial distance between them)
- Fitness computed from random subsamples
- Half of population with lower fitness deleted replaced with a copy of a surviving neighbor
- Pairing in the local neighborhoods
- Special crossover for diploids, followed by mutation with p_m = 0.001.
- Population size between 512—1million
- 5000 generation
- → result: GA found sorting net of 65 comparisons

- Why didn't the GA do better?
 - After early generations, with randomly generated test cases used to compute fitness, the difficulty of test cases stayed roughly the same!
- Solution: Co-Evolution
 - = evolve both algorithms and test cases!
 - Algorithms try to sort
 - Test cases try to 'trip up' algorithms
 - → predator/prey (or host/parasite) relationship inspired from nature
 - → as the algorithms got better, the test cases got harder (specifically targeting weaknesses in the networks)
- Result: 61 comparisons





- Importance of the work:
 - introduces new technique inspired by coevolution. The results are convincing in that this new technique is potentially powerful.

Inter-population competitive co-evolution

Example2: Game Playing

- Intra-population competitive co-evolution
- Task: evolve a backgammon player
- Problem: evaluation
 - Against human player
 - Against 'conventional' program
 - Against internet players
- Solution: co-evolution
 - Play against other evolving programs
- Intra-population
 - All genotypes are of the same type
 - Only one population

Co-Evolving Backgammon Players

- TD-Gammon
 - Grand Master level player
 - Learns by self-playing
- What makes it successful?
 - Temporal-Difference Learning ?
 - Or simply self-playing ?
- Simple use of NN in game play:
 - Generate all legal moves
 - Feed them to the NN
 - Best output is new move
 - NN used to evaluate positions

Simple Backgammon Learner

Evolve a NN that plays backgammon

- 1. Initial NN is NN(k); k=0;
- 2. Generate a mutant challenger of NNk:

$$w'(i,j) = w(i,j) + Gaussian (0, s);$$

3. If NN'(k) is beaten by NN(k)

$$NN(k+1) = NN(k)$$

else

$$NN(k+1) = 0.95 * NN(k) + 0.05 * NN'(k)$$

4. k = k+1; go to step 2 unless finished

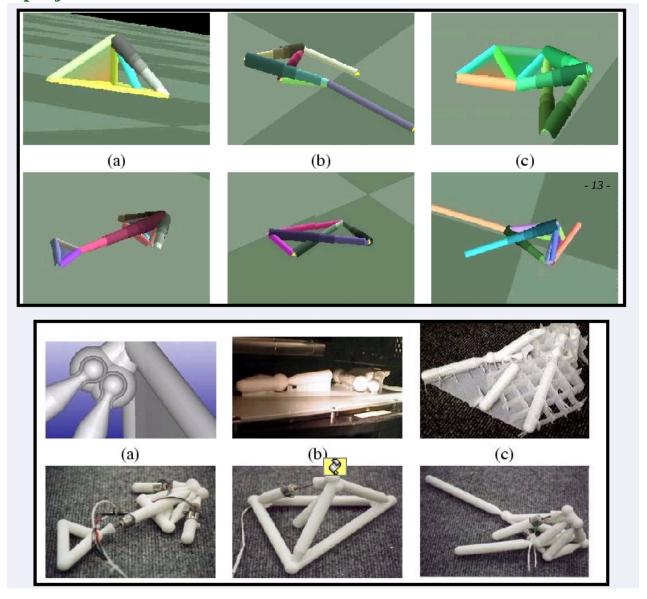
Setting

- Uses 197-20-1 fully connected feed-forward NN; initial weights 0s
- No training for NNs!
- EA with population size 1
- Only simple mutation: add Gaussian noise to weights
- No recombination
- Performance
- 40% winning against a strong expert-trained program (PUBEVAL) after 100.000 generations

Example3: Evolution of a complete robot

- Inter-population Cooperative Co-Evolution
- Task: Evolve Morphology and Behaviour of a robot
 - Structure
 - Control
- Evaluation: no problem
 - Simulation
 - But: slow
- Problem: Evolution
 - All-in-one is possible, but:
 - Very large search space
- Solution: Co-Evolution
 - Divide Problems
 - Evolve Structure and Control separately

GOLEM project, Pollack et. al



Coupling cooperative Co-Evolution

Tight coupling (1)

- Body-Individual and Controller-Individuals evolve into 'matched pair'
- Body 'a' only works well with controller 'b'

Option: Two-Part Genotype

- Keep both genotypes in one Individual
- Evolve both together

Tight coupling (2)

- Change in one might depend on change in the other
- Can slow down evolution

Option: delayed evaluation

- Mutate morphology individuals
- Evolve controller individuals (possibly more than one generation)
- Assign fitness to morphology individuals

Example4: Pattern recognition

- Intra-population Cooperative Co-Evolution
- Taks: Evolve a NN to recognize a set of letters
 - Grey values fed into NN
 - Output classifies letters
 - Classification problem
- Co-Evolution: Divide and Conquer
 - Different Individuals specialize on different letters
 - Group output is combined into one by external mechanism (e.g. voting)
- Speciation

When and Why Co-Evolution?

- No fitness-function known
 - Bootstrapping by co-evolution
- Too many fitness cases
 - Co-Evolve fitness and cases
- Modularizable Problem
 - Divide and Conquer

Other examples

- Creative Design Systems
 - Evolve designs and design specifications
- Other Games
 - Tic-Tac-Toe
 - Prisoner's Dilema
 - Checkers
- Artificial Life
 - Complex simulated Ecosystems

Summary

- Co-Evolution is Everywhere
- Can be cooperative and competitive
- Can be in one population, or more than one
- An individual's fitness is not fixed in coevolution
- Co-evolution is not well explored in Evolutionary Computation

References

- Mitchell, M, Introduction to Genetic Algorithms (book), MIT Press, 1996, pp.21—27.
- Pollack, J, Blair, A. and Land, M. (1996). Coevolution of A Backgammon Player. *Proceedings Artificial Life V*, C. Langton, (Ed), MIT Press. Available at: http://www.demo.cs.brandeis.edu/papers/alife5.pdf
- Pollack, Jordan B., Lipson, Hod, Hornby, Gregory S., and Funes, Pablo. Three Generations of Automatically Designed Robots. *Artificial Life*, 7:3, pg 215-223. 2001. Available at: http://www.demo.cs.brandeis.edu/papers/long.html#pollack_alife01
- Darwen, P. J. and Yao, X. (1996). Automatic Modularization by Speciation. *Third IEEE International Conference on Evolutionary Computation*, available

at: http://www.demo.cs.brandeis.edu/papers/icec96darwen.ps.gz