

Problem solving exercise class
N-player Iterated Prisoner's Dilemma
(co-evolution case study)

Ata Kaban

University of Birmingham

Iterated Prisoner's Dilemma

- Invented by Merrill Flood & Melvin Dresher in 1950s
- Studied in game theory, economics, political science
- The story
 - Alice and Bob arrested, no communication between them
 - They are offered a deal:
 - If any of them confesses & testifies against the other then gets suspended sentence while the other gets 5 years in prison
 - If both confess & testify against the other, they both get 4 years
 - If none of them confesses then they both get 2 years
 - What is the best strategy for maximising one's own payoff?

N-Player Version

The payoff matrix of the N-player iterated prisoner's dilemma game, for Player A is:

	0	1	2	...	N-1
Coop	0	2	4		$2(N-1)$
Defect	1	3	5		$2(N-1)+1$

All players are treated equally.

Design a co-evolutionary algorithm for learning to play the iterated 4-player prisoner's dilemma game.

- Chromosome representation for strategies (players)
- Fitness evaluation function
- Evolutionary operators (crossover, mutations)
- Selection scheme
- Comment on parameters of your design.
- Comment on strengths and weaknesses of your design

Recap: 2-player version of the game
(To skip the recap, jump to slide 13)

- Abstract formulation through a payoff matrix

		Player A	
		Cooperate	Defect
Player B	Cooperate	3,3	0,5
	Defect	5,0	1,1

- 2 tournaments – participants have sent strategies
- Human strategies played against each other
- Winner: TIT FOR TAT
 - Cooperates as long the other player does not defect
 - Defects on defection until the other player begins to cooperate again
- Can GA evolve a better strategy?

- Individuals = strategies
- How to encode a strategy by a string?
- Let memory depth of previous moves=1

Fix a canonical order of cases:

	A	B
– Case 1:	C	C
– Case 2:	C	D
– Case 3:	D	C
– Case 4:	D	D

e.g. strategy encoding (for A): ‘CDCD’

- Now let memory depth of previous moves=3
 - How many cases?
 - Case 1:
 - Case 2:
 - ...
 - How many letters are needed to encode a strategy as a string?
 - How many strategies there are?
 - Is that a large number?

- Experiment 1
 - 40 runs with different random initialisations
 - 50 generations each
 - Population of 20
 - Fitness=avg score over all games played
 - A fixed environment of 8 human-designed strategies
- Results
 - Found better strategy than those 8 strategies in the environment!
 - Even though – how many strategies were only tested in a run out of all possible strategies?
 - What does this result mean?

- Experiment 2
 - changing environment: the evolving strategies played against each other.
- Results
 - Found strategies similar in essence with the winner human-designed strategy
- Idealised model of evolution & co-evolution

N-Player Version

The payoff matrix of the N-player iterated prisoner's dilemma game, for Player A is:

	0	1	2	...	N-1
Coop	0	2	4		$2(N-1)$
Defect	1	3	5		$2(N-1)+1$

All players are treated equally.

Representation

- Strategy = lookup table
 - Situation (history) \rightarrow action, for each situation
- How to represent history of the game?
 - Let l denote the length of the history considered
- How many histories are possible in this game?

...representing history

- The player's own previous l moves
 - Requires bits
 - The number of co-operators in the last l moves
 - Requires bits
- That is bits in total

- An example of encoded history, if $l=3$:

001111001

What does it mean?

- need a convention as of which bit means what
 - o Let the first l bits indicate the player's own actions
 - o Let the leftmost bit refer to the most recent move
 - o Let the next groups of 2 bits indicate the nos of collaborators
 - o Let the leftmost group refer to the most recent move

001 11 10 01

Now the bit-string 'makes sense'!

Can you read the story from the bit-string now?

- How many histories are there in total?

If 9 bits are needed to represent a history

Then there are 2^9 histories possible.

Remember, we agreed that strategies will be stored as
‘lookup tables’:

One strategy is a binary string (0=coop, 1=defect) that gives an action for all possible histories. How long this string is?.....

So 2^9 bits are needed to represent a strategy.

e.g. for history ‘001 11 10 01’=121, the action is whatever is listed in entry (bit) 121.

- Sure? Anything missing?

- What is missing?
 - Actions are taken as function of the history
 - What about the very first action?
- Need some extra bits to represent the “pre-history” of $l=3$ virtual previous rounds at the beginning of the game!
 - That is 9 extra bits. (But there several possibilities here.)
- Length of bit string that represents one strategy:
 - It's not 2^9 but 2^9+9
- *Q: Would you be able to write this quantity more generally, with history length l and nos of players N ?*

Fitness evaluation function

- Fitness of an individual player is evaluated by playing a number K of (N -player) games with adversaries randomly drawn from the population.
- Adding the payoffs obtained by each individual is a measure of its fitness.

Evolutionary operators

- Since binary strings are used, standard operators are available, e.g.
 - Uniform crossover
 - Bit-wise flipping can be used

Selection scheme

- Fitness ranking or tournament selection
- Important is to maintain the *selection pressure* constantly!

Discussion. Strengths, weaknesses

- Strengths:
 - Generic, the same design can be applied for more general number of players N
 - Simplicity in evolution and game playing due to bit string representation
- Weaknesses:
 - In N is large, computation time is long
 - Inability to capture multiple cooperation levels
- Parameters that influence the results:
 - History length
 - Nos of generations