

Probabilistic Robotics*

Non-parametric Bayes Filter Implementations

Particle filters


Mohan Sridharan

University of Birmingham, UK

m.sridharan@bham.ac.uk

*Revised original slides that accompany the book by Thrun, Burgard and Fox.

Sample-based Localization (sonar)



Particle Filters


- Represent belief by random **samples**.
- Estimation of **non-Gaussian**, **nonlinear** processes.
- Monte Carlo filter, survival of the fittest, I-condensation, bootstrap filter, particle filter.
- Filtering: [Rubin, 88], [Gordon et al., 93], [Kitagawa 96].
- Computer vision: [Isard and Blake 96, 98].
- Dynamic Bayesian Networks: [Kanazawa et al., 95].

Particle Filter Algorithm (basic)

Algorithm **particle_filter**(χ_{t-1}, u_t, z_t):

1. $\bar{\chi}_t = \chi_t = \emptyset$
2. **For** $m = 1 \dots M$
3. Sample $x_t^{[m]} \sim p(x_t \mid x_{t-1}^{[m]}, u_t)$
4. $w_t^{[m]} = p(z_t \mid x_t^{[m]})$
5. $\bar{\chi}_t = \bar{\chi}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$
6. **For** $m=1 \dots M$
7. Draw i with probability $\propto w_t^{[i]}$
8. Add $x_t^{[i]}$ to χ_t
9. **Return**

Importance Sampling



Weight samples: $w = f/g$

Importance Sampling

$$f(.) = \text{bel}(x_t) = \eta p(z | x) \overline{\text{bel}}(x_t)$$

$$g(.) = \overline{\text{bel}}(x_t) = \sum p(x_t | u_t, x_{t-1}) \text{bel}(x_{t-1})$$

Function $f(.)$: target distribution.

Function $g(.)$: proposal distribution.

Weights: $w(x) = f(x) / g(x)$

Need: $f(x) > 0 \rightarrow g(x) > 0$






Converges to desired distribution iteratively.

PF derivation ([Section 4.3.3, PR](#))


Importance Sampling with Resampling: Landmark Detection Example






Distributions



Distributions






Wanted: samples distributed according to $p(x | z_1, z_2, z_3)$



This is Easy!

We can draw samples from $p(x|z_i)$ by adding noise to the detection parameters.




Importance Sampling with Resampling

$$\text{Target distribution } f : p(x | z_1, z_2, \dots, z_n) = \frac{\prod_k p(z_k | x) p(x)}{p(z_1, z_2, \dots, z_n)}$$


$$\text{Sampling distribution } g : p(x | z_l) = \frac{p(z_l | x) p(x)}{p(z_l)}$$

$$\text{Importance weights } w : \frac{f}{g} = \frac{p(x | z_1, z_2, \dots, z_n)}{p(x | z_l)} = \frac{p(z_l) \prod_{k \neq l} p(z_k | x)}{p(z_1, z_2, \dots, z_n)}$$

Importance Sampling with Resampling




Weighted
samples



After
resampling

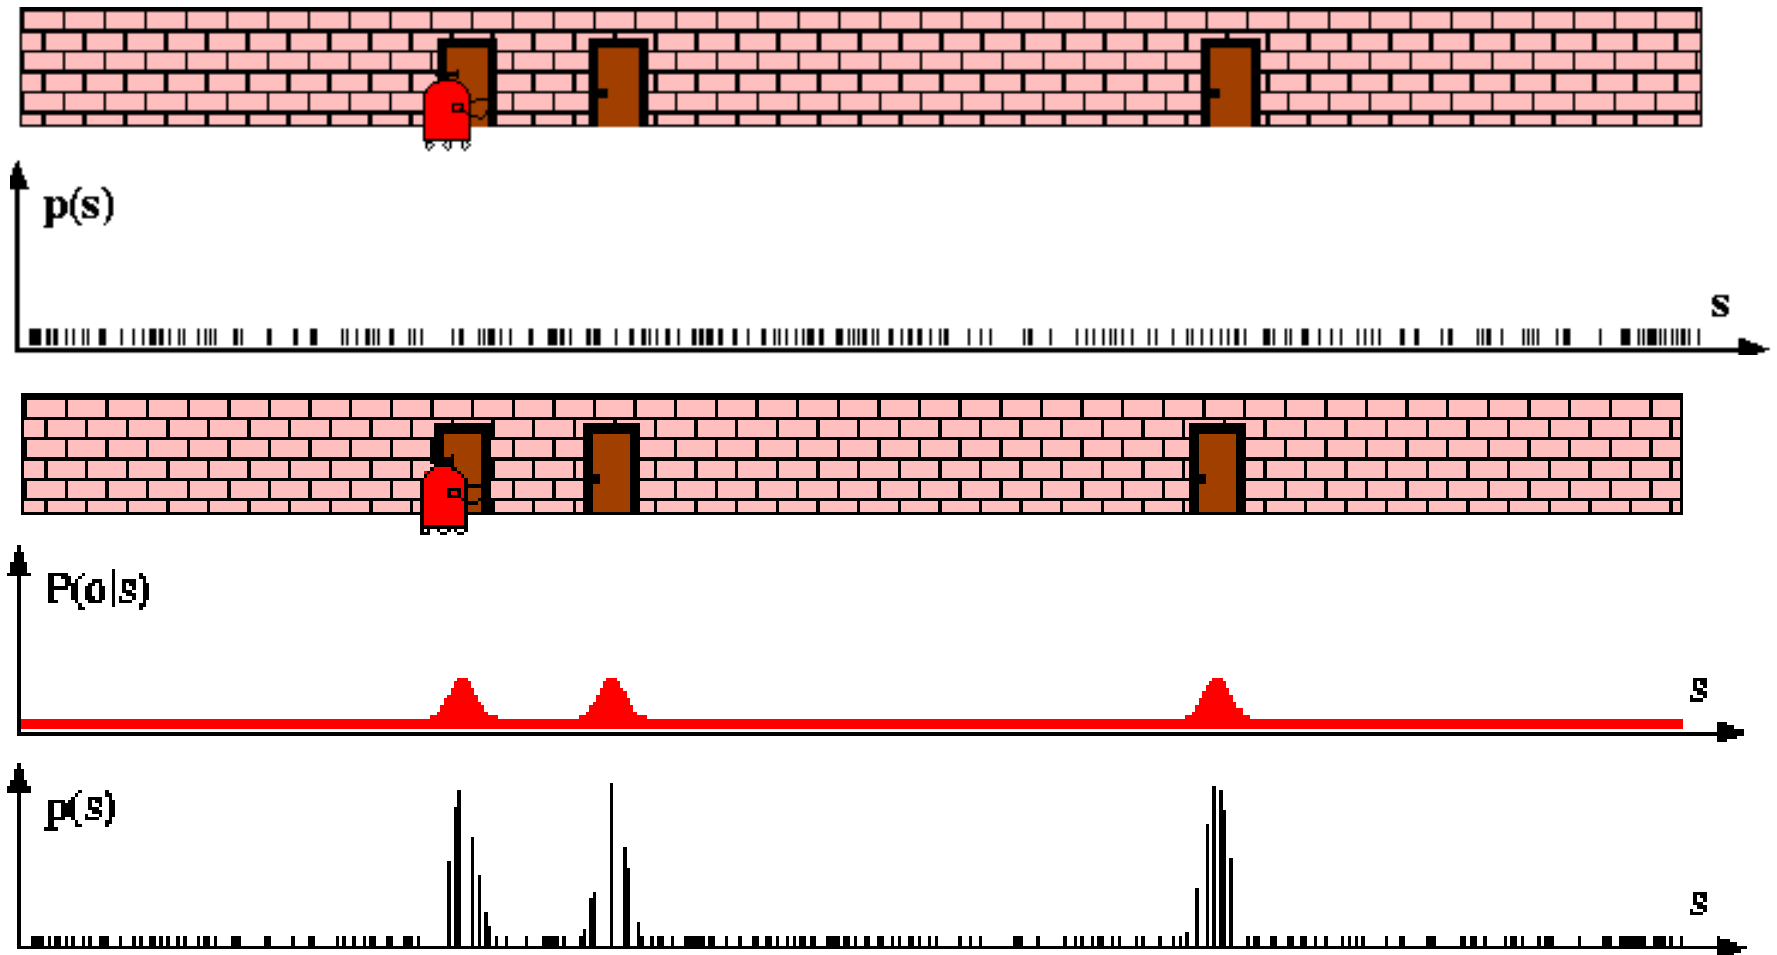
Particle Filters



Sensor Information: Importance Sampling



$$Bel(x) \leftarrow \alpha p(z | x) Bel^-(x)$$

$$w \leftarrow \frac{\alpha p(z | x) Bel^-(x)}{Bel^-(x)} = \alpha p(z | x)$$



Robot Motion

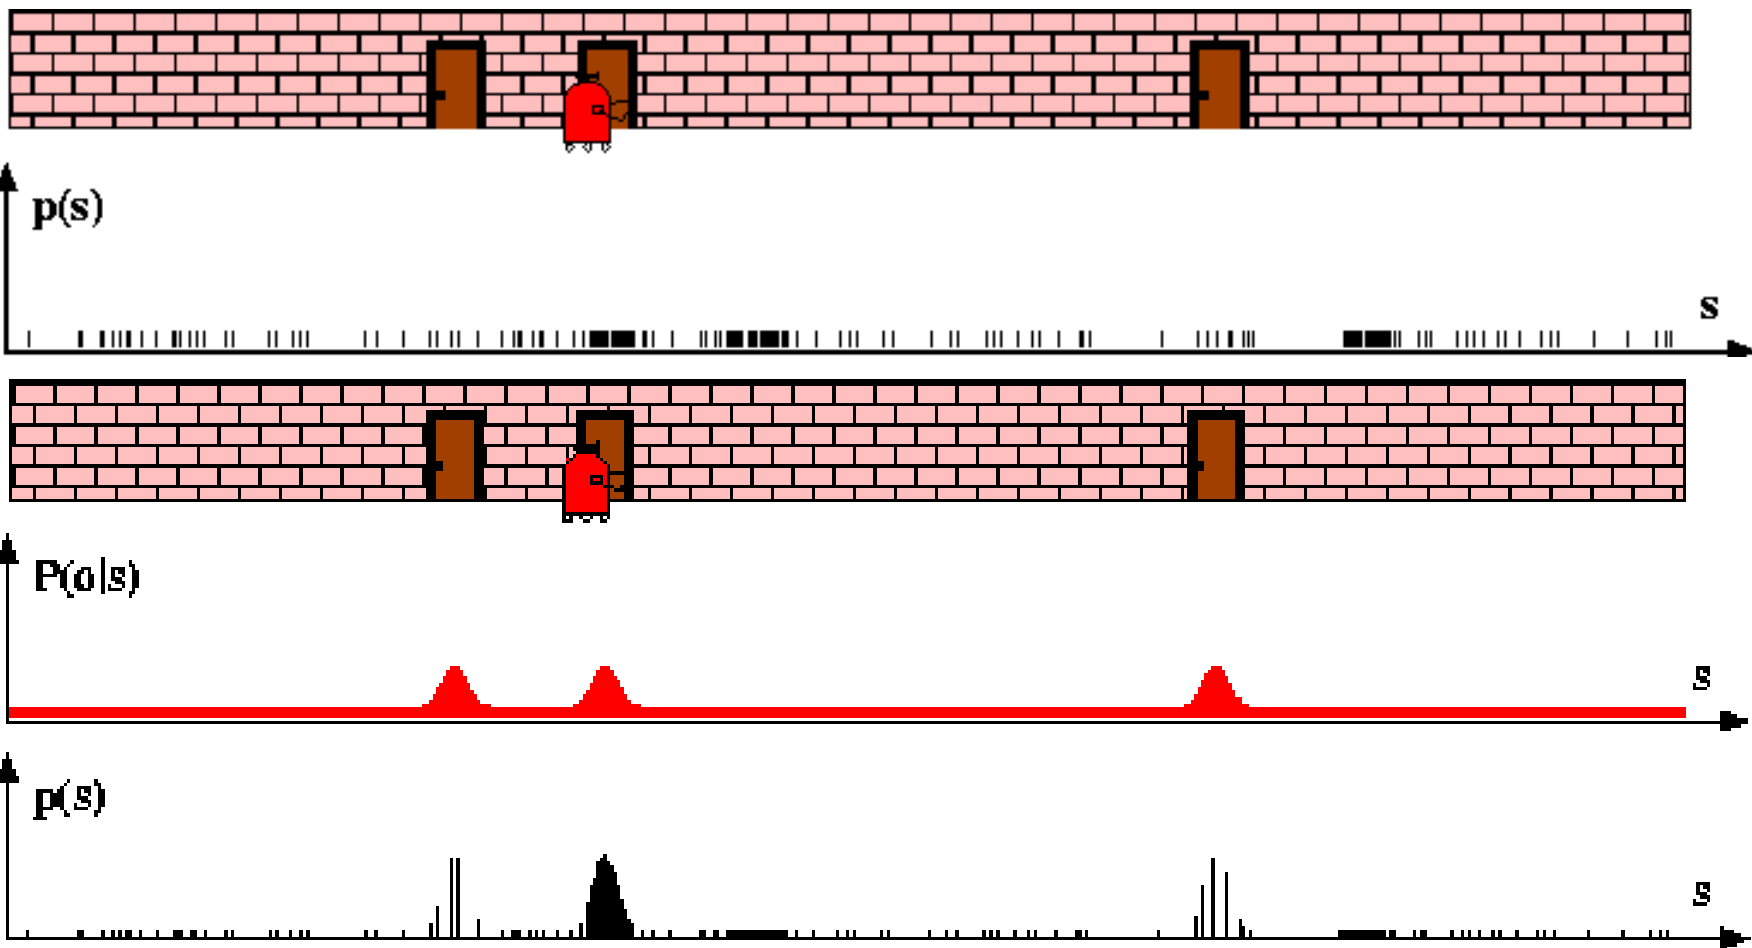
$$Bel^-(x) \leftarrow \int p(x|u, x') Bel(x') dx'$$



Sensor Information: Importance Sampling



$$Bel(x) \leftarrow \alpha p(z | x) Bel^-(x)$$

$$w \leftarrow \frac{\alpha p(z | x) Bel^-(x)}{Bel^-(x)} = \alpha p(z | x)$$




Robot Motion

$$Bel^-(x) \leftarrow \int p(x|u, x') Bel(x') \, dx'$$



Particle Filter Algorithm

$$Bel(x_t) = \eta p(z_t | x_t) \int p(x_t | x_{t-1}, u_t) Bel(x_{t-1}) dx_{t-1}$$




$$\begin{aligned} w_t^i &= \frac{\text{target distribution}}{\text{proposal distribution}} \\ &= \frac{\eta p(z_t | x_t) p(x_t | x_{t-1}, u_t) Bel(x_{t-1})}{p(x_t | x_{t-1}, u_t) Bel(x_{t-1})} \\ &\propto p(z_t | x_t) \end{aligned}$$

Resampling

- **Given**: Set S of weighted samples.
- **Wanted** : Random sample, where the probability of drawing x_i is given by w_i .
- Typically done M times with replacement to generate new sample set S' .

Resampling



- Roulette wheel
- Binary search, $n \log n$
- Stochastic universal sampling
- Systematic resampling
- Linear time complexity
- Easy to implement, low variance

Resampling Algorithm

1. Algorithm **systematic_resampling**(S, M):

1. $S' = \emptyset, c_1 = w^1$

2. **For** $i = 2 \dots M$

Generate cdf

3. $c_i = c_{i-1} + w^i$

4. $u_1 \sim U(0, M^{-1}], i = 1$

Initialize threshold

1. **For** $j = 1 \dots M$

Draw samples ...

2. **While** ($u_j > c_i$)

Skip until next threshold reached

3. $i = i + 1$

4. $S' = S' \cup \{x^i, M^{-1} >\}$


Insert


5. $u_{j+1} = u_j + M^{-1}$


Increment threshold


6. **Return** S'


Also called **stochastic universal sampling**


















































Initial Distribution




After Incorporating Ten Ultrasound Scans




After Incorporating 65 Ultrasound Scans



Estimated Path




Using Ceiling Maps for Localization



Under a Light

Measurement

z:




$P(z/x):$




Next to a Light

Measurement

z:



$P(z/x)$:




Elsewhere

Measurement


z:




$P(z/x)$:



Global Localization Using Vision



Localization for AIBO robots



Limitations

- The approach described so far is able to:
 - Track the pose of a mobile robot.
 - Globally localize the robot.
- Can amplify sampling variance, i.e., variability from original distribution due to random sampling.
- *Sampling bias* and *particle deprivation*.
- How can we deal with localization errors, e.g., the *kidnapped robot problem*?

Approaches

- Randomly insert samples;
 - Robot can be “*teleported*” at any point in time 😊
- Insert random samples proportional to the average likelihood of the particles:
 - Robot has been teleported with higher probability when the likelihood of its observations drops.

Summary

- Particle filters instance of recursive Bayesian filtering.
- Represent the posterior by a set of weighted samples.
- In the context of localization, particles are propagated according to the motion model.
- Particles are then weighted according to the likelihood of the observations.
- During re-sampling, new particles are drawn with probability proportional to the weights.

What Next?

- SLAM!
- EKF-SLAM and Fast-SLAM.
- Probabilistic sequential decision making.