

Probabilistic Robotics*

Partially Observable Markov Decision Process

Mohan Sridharan

University of Birmingham, UK

m.sridharan@bham.ac.uk

*Revised original slides that accompany the book by Thrun, Burgard and Fox.

POMDPs

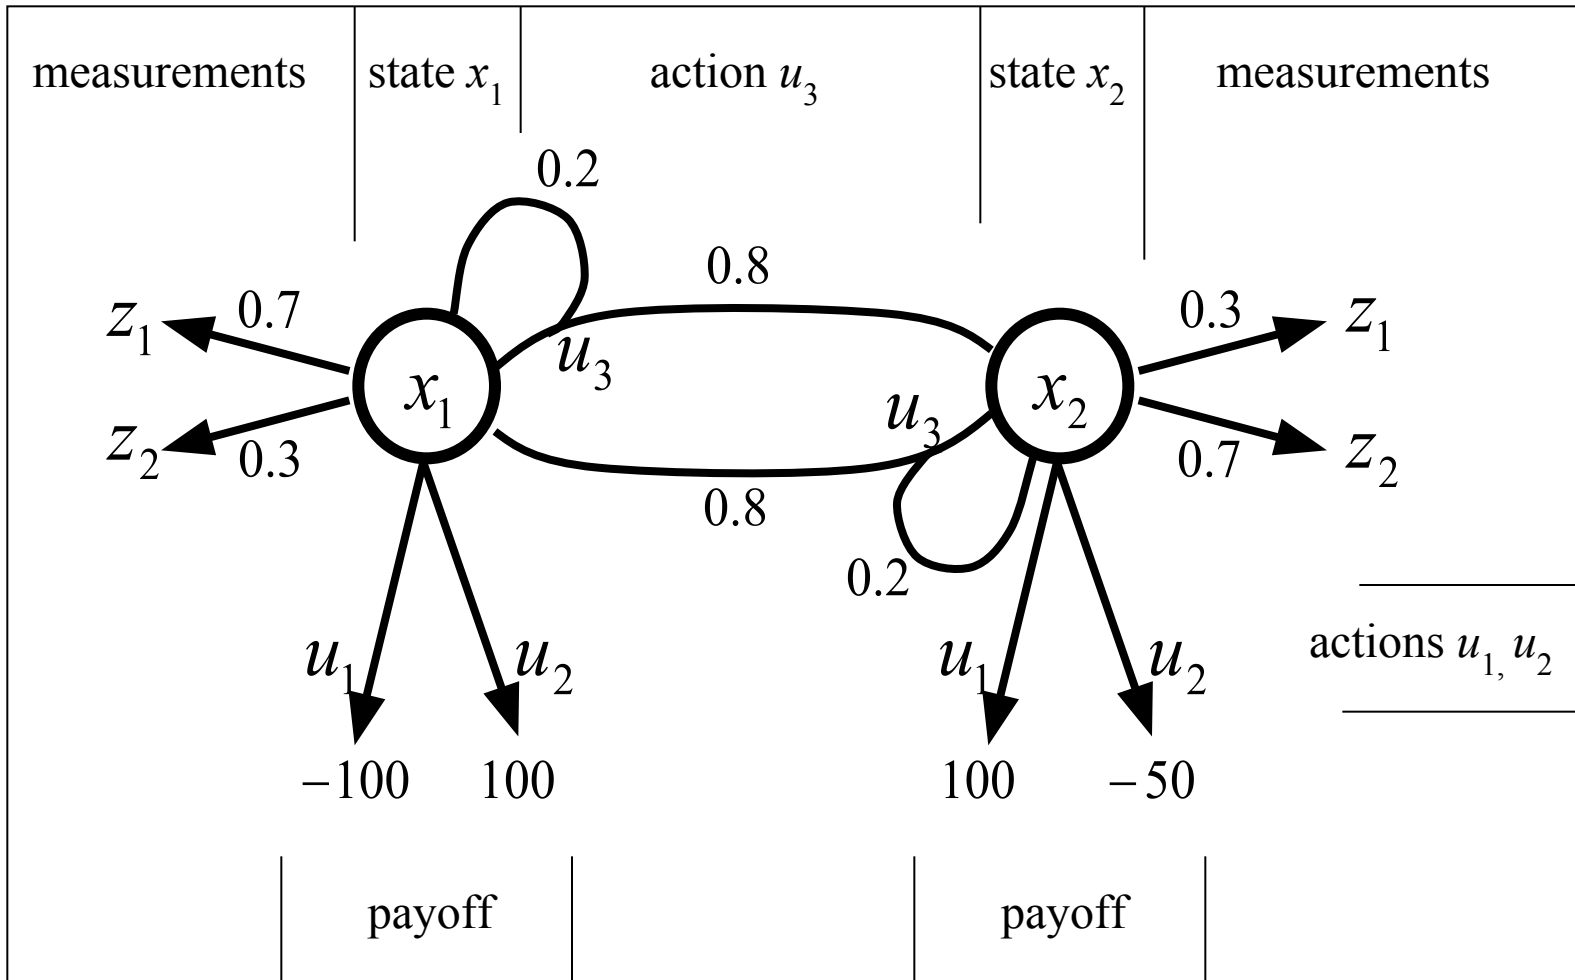
- **State is not observable** – agent has to make decisions based on belief state which is a **posterior distribution over states**.
- Let b be the belief of the agent about the state under consideration.
- POMDPs compute a **value function over belief space**:

$$V_T(b) = \gamma \max_u \left\{ r(b, u) + \int V_{T-1}(b') p(b' | u, b) db' \right\}$$
$$\pi_T(b) = \arg \max_u \left\{ r(b, u) + \int V_{T-1}(b') p(b' | u, b) db' \right\}$$
$$V_T(x) = \gamma \max_u \left\{ r(x, u) + \int V_{T-1}(x') p(x' | u, x) dx' \right\}$$

Problems

- Belief is a probability distribution – each value in a POMDP is a function of an entire probability distribution!
- Probability distributions are continuous.
- Huge complexity of belief spaces.
- For **finite worlds** with finite state, action, and observation spaces and finite horizons, we can **effectively represent the value functions by piecewise linear functions**.

An Illustrative Example



The Parameters of the Example

- The actions u_1 and u_2 are terminal actions.
- The action u_3 is a sensing action that potentially leads to a state transition.
- The horizon is finite and $\gamma=1$.

$$r(x_1, u_1) = -100, \quad r(x_2, u_1) = 100$$

$$r(x_1, u_2) = 100, \quad r(x_2, u_2) = -50$$

$$r(x_1, u_3) = -1, \quad r(x_2, u_3) = -1$$

$$p(x'_1 | x_1, u_3) = 0.2 \quad p(x'_2 | x_1, u_3) = 0.8$$

$$p(x'_1 | x_2, u_3) = 0.8 \quad p(x'_2 | x_2, u_3) = 0.2$$

$$p(z_1 | x_1) = 0.7 \quad p(z_2 | x_1) = 0.3$$

$$p(z_1 | x_2) = 0.3 \quad p(z_2 | x_2) = 0.7$$

$$p_1 = b(x_1), \quad p_2 = b(x_2), \quad p_2 = 1 - p_1$$

$$\text{Policy } \pi : [0, 1] \rightarrow u$$

Payoff in POMDPs

- In MDPs, the payoff (or return) depends on the state of the system.
- In POMDPs, the true state is not known.
- Therefore, we compute the **expected payoff** by **integrating over all states**:

$$\begin{aligned} r(b, u) &= E_x[r(x, u)] \\ &= \int r(x, u) p(x) dx \\ &= p_1 r(x_1, u) + p_2 r(x_2, u) \end{aligned}$$

Payoffs in Our Example (1)





- If we are certain that we are in state x_1 and execute action u_1 we receive reward of -100.
- If we definitely know that we are in x_2 and execute u_1 the reward is +100.
- In between it is the linear combination of the extreme values weighted by the probabilities:

$$\begin{aligned} r(b, u_1) &= -100 p_1 + 100 p_2 \\ &= -100 p_1 + 100 (1 - p_1) \end{aligned}$$

$$r(b, u_2) = 100 p_1 - 50 (1 - p_1)$$

$$r(b, u_3) = -1$$

Payoffs in Our Example (2)



The Resulting Policy for T=1

- Given we have a finite POMDP with T=1, we would use $V_I(b)$ to determine the optimal policy.
- In our example, the optimal policy for T=1 is:

$$\pi_1(b) = \begin{cases} u_1 & \text{if } p_1 \leq \frac{3}{7} \\ u_2 & \text{if } p_1 > \frac{3}{7} \end{cases}$$

- This is the upper thick graph in the diagram.

Piecewise Linearity, Convexity

- The resulting value function $V_I(b)$ is the maximum of the three functions at each point:

$$\begin{aligned} V_1(b) &= \max_u r(b, u) \\ &= \max \left\{ \begin{array}{cc} -100 p_1 & +100 (1 - p_1) \\ 100 p_1 & -50 (1 - p_1) \\ & -1 \end{array} \right\} \end{aligned}$$

- It is piecewise linear and convex.


Pruning

- Carefully consider $V_I(b)$ – only the first two components contribute.
- The third component can be pruned away from $V_I(b)$:

$$V_1(b) = \max \left\{ \begin{array}{cc} -100 p_1 & +100 (1 - p_1) \\ 100 p_1 & -50 (1 - p_1) \end{array} \right\}$$

Increasing the Time Horizon

- Assume the robot can make an observation before deciding on an action.



Increasing the Time Horizon

- Assume the robot can make an observation before deciding on an action.
- Suppose the robot perceives z_1 for which:

$$p(z_1 | x_1) = 0.7 \text{ and } p(z_1 | x_2) = 0.3.$$

- Given the observation z_1 we update the belief using Bayes rule:




$$p'_1 = p(x_1 | z) = \frac{p(z_1 | x_1)p(x_1)}{p(z_1)} = \frac{0.7 p_1}{p(z_1)}$$

$$p'_2 = \frac{0.3(1 - p_1)}{p(z_1)}$$

$$p(z_1) = 0.7 p_1 + 0.3(1 - p_1) = 0.4 p_1 + 0.3$$

Value Function

$p_1' = p_1$ after sensing z_1



Increasing the Time Horizon

- Coming back to our assumption that robot can make an observation before deciding on an action.
- If the robot perceives z_1 :
 $p(z_1 | x_1)=0.7$ and $p(z_1 | x_2)=0.3$.
- We update the belief $V_1(b | z_1)$ using Bayes rule to obtain:

$$\begin{aligned} V_1(b | z_1) &= \max \left\{ \begin{array}{cc} -100 \cdot \frac{0.7 p_1}{p(z_1)} & +100 \cdot \frac{0.3 (1-p_1)}{p(z_1)} \\ 100 \cdot \frac{0.7 p_1}{p(z_1)} & -50 \cdot \frac{0.3 (1-p_1)}{p(z_1)} \end{array} \right\} \\ &= \frac{1}{p(z_1)} \max \left\{ \begin{array}{cc} -70 p_1 & +30 (1 - p_1) \\ 70 p_1 & -15 (1 - p_1) \end{array} \right\} \end{aligned}$$

Expected Value after Measuring

- Since we do not know what the next measurement will be, we have to compute the expected belief:

$$\begin{aligned}\bar{V}_1(b) &= E_z[V_1(b | z)] = \sum_{i=1}^2 p(z_i) V_1(b | z_i) \\ &= \sum_{i=1}^2 p(z_i) V_1\left(\frac{p(z_i | x_1) p_1}{p(z_i)}\right) \\ &= \sum_{i=1}^2 V_1(p(z_i | x_1) p_1)\end{aligned}$$

Expected Value after Measuring

- Since we do not know what the next measurement will be, we have to compute the expected belief:

$$\begin{aligned}\bar{V}_1(b) &= E_z[V_1(b \mid z)] \\ &= \sum_{i=1}^2 p(z_i) V_1(b \mid z_i) \\ &= \max \left\{ \begin{array}{cc} -70 p_1 & +30 (1 - p_1) \\ 70 p_1 & -15 (1 - p_1) \end{array} \right\} \\ &\quad + \max \left\{ \begin{array}{cc} -30 p_1 & +70 (1 - p_1) \\ 30 p_1 & -35 (1 - p_1) \end{array} \right\}\end{aligned}$$




Resulting Value Function

- The four possible combinations yield the following function which then can be simplified and pruned:




$$\begin{aligned}\bar{V}_1(b) &= \max \left\{ \begin{array}{cccc} -70 p_1 & +30 (1 - p_1) & -30 p_1 & +70 (1 - p_1) \\ -70 p_1 & +30 (1 - p_1) & +30 p_1 & -35 (1 - p_1) \\ +70 p_1 & -15 (1 - p_1) & -30 p_1 & +70 (1 - p_1) \\ +70 p_1 & -15 (1 - p_1) & +30 p_1 & -35 (1 - p_1) \end{array} \right\} \\ &= \max \left\{ \begin{array}{cc} -100 p_1 & +100 (1 - p_1) \\ +40 p_1 & +55 (1 - p_1) \\ +100 p_1 & -50 (1 - p_1) \end{array} \right\}\end{aligned}$$

Value Function

$p_1' = p_1$ after sensing z_1



Value Function




$$\bar{V}_1(b) = p(z_1) V_1(b|z_1) + p(z_2) V_1(b|z_2)$$

State Transitions (Prediction)

- When the agent selects u_3 its state potentially changes.
- When computing the value function, we have to take these potential state changes into account.

$$\begin{aligned} p_1' &= E_x[p(x_1' \mid x, u_3)] \\ &= \sum_{i=1}^2 p(x_1' \mid x_i, u_3) p_i \\ &= 0.2 p_1 + 0.8(1 - p_1) \\ &= 0.8 - 0.6 p_1 \end{aligned}$$

State Transitions (Prediction)



$$p_1' = 0.8 - 0.6p_1$$

Value Function after executing u_3




- Taking the state transitions into account:

$$\bar{V}_1(b) = \max \left\{ \begin{array}{cccc} -70 p_1 & +30 (1 - p_1) & -30 p_1 & +70 (1 - p_1) \\ -70 p_1 & +30 (1 - p_1) & +30 p_1 & -35 (1 - p_1) \\ +70 p_1 & -15 (1 - p_1) & -30 p_1 & +70 (1 - p_1) \\ +70 p_1 & -15 (1 - p_1) & +30 p_1 & -35 (1 - p_1) \end{array} \right\}$$

$$= \max \left\{ \begin{array}{cc} -100 p_1 & +100 (1 - p_1) \\ +40 p_1 & +55 (1 - p_1) \\ +100 p_1 & -50 (1 - p_1) \end{array} \right\}$$

$$\bar{V}_1(b \mid u_3) = \max \left\{ \begin{array}{cc} 60 p_1 & -60 (1 - p_1) \\ 52 p_1 & +43 (1 - p_1) \\ -20 p_1 & +70 (1 - p_1) \end{array} \right\}$$

Value Function after executing u_3




Value Function for T=2

- Taking into account that the agent can either directly perform u_1 or u_2 or first u_3 and then u_1 or u_2 , we obtain (after pruning).



$$\bar{V}_2(b) = \max \left\{ \begin{array}{cc} -100 p_1 & +100 (1 - p_1) \\ 100 p_1 & -50 (1 - p_1) \\ 51 p_1 & +42 (1 - p_1) \end{array} \right\}$$

Graphical Representation of $V_2(b)$









Deep Horizons and Pruning

- We have now completed a full backup in belief space.
- This process can be applied recursively.
- The value functions for $T=10$ and $T=20$:



Deep Horizons and Pruning



Why Pruning is Essential

- Each update adds additional linear components to V .
- Each measurement squares number of linear components.
- Unpruned value function for $T=20$ has more than $10^{547,864}$ linear functions.
- At $T=30$ we have $10^{561,012,337}$ linear functions.
- The pruned value functions at $T=20$, in comparison, contains only 12 linear components.
- The combinatorial explosion of linear components in the value function is the major reason why POMDPs are impractical for most applications.

POMDP Summary

- POMDPs compute the optimal action in partially observable, stochastic domains.
- For finite horizon problems, the resulting value functions are piecewise linear and convex.
- In each iteration the number of linear constraints grows exponentially.
- POMDPs so far have only been applied successfully to very small state spaces with small numbers of possible observations and actions.

POMDP Approximations


- Point-based value iteration.
- QMDPs.
- AMDPs.
- MC-POMDP.

Point-based Value Iteration


- Maintains a set of example beliefs.
- Only considers constraints that maximize value function for at least one of the examples.

Point-based Value Iteration

Value functions for $T=30$



Exact value function



PBVI

QMDPs

- QMDPs only consider state uncertainty in the first step.
- After that, the world becomes fully observable!
- Planning only marginally less efficient than MDPs, but performance significantly better!

Monte Carlo POMDPs

- Represent beliefs by samples.
- Estimate value function on sample sets.
- Simulate control and observation transitions between beliefs.

Summary

- POMDPs ideal for modeling systems with partially observable state and non-deterministic actions.
- Exponential state space explosion is a problem!
- Methods exist to make the problem more tractable – not good enough for real-world robot problems.
- Possible solutions:
 - Impose hierarchy by exploiting inherent structure.
 - Speed up POMDP solution techniques.